



International Business School



PANAMERICAN
UNIVERSITY

Master in AI for Business and Data Science

Spain | 2026

DEMAND FORECASTING AND INVENTORY OPTIMISATION USING MACHINE LEARNING

A simulation-based evaluation on the M5-Forecasting (Walmart) dataset, with an SME-accessible Power BI dashboard.

Authors

Desmond Korbla Aziega

Theresa Korlekuor Apla-Kweku

Abdul-Razak Seidu

Supervisor

Prof. Rocío González Martínez

Supervisor Approval



Mr./Ms. Rocío González Martínez, as tutor of the project titled **Demand Forecasting and Inventory Optimisation Using Machine Learning, A Simulation-Based Evaluation for Retail Decision-Makers, with an SME-Accessible Power BI Dashboard**, pertaining to the Master Master in AI for Business and Data Science,

I declare that I have supervised and assessed, on an overall basis, the work carried out by the students: Desmond Korbla Aziega, Theresa Korlekuor Apla-Kweku and Abdul-Razak Seidu

CIF / G30164099

GONZALEZ Firmado
MARTINEZ digitalmente por
MARIA DEL GONZALEZ
ROCIO - MARTINEZ MARIA
03105420Y DEL ROCIO -
03105420Y 03105420Y
Fecha: 2026.07.03
09:45:27 +02'00'

Sing TUTOR/A

Murcia, 2 July 2026

Centro adscrito a:



Universidad
Politécnica
de Cartagena



Edificio ENAE nº13. Campus Universitario de Espinardo, CP 30100 Murcia. Telf: 968 899 899 Fax: 868 88 41 33 Email: info@enae.es www.enae.es

Acknowledgements

We would like to express our deepest gratitude to our supervisor, Prof. Rocío González Martínez, for her guidance, patience, and rigorous feedback throughout this research. Her insistence that every claim in this thesis be measured rather than assumed; most visibly in her critique of the original inventory-cost finding fundamentally shaped this work for the better, and we are profoundly grateful for the standard of honesty and rigour she held us to.

We are equally grateful to our families back home, whose patience, encouragement, and sacrifice made it possible for us to devote the time and energy this thesis demanded, often from a distance and across many long nights. Finally, we thank everyone else; friends, colleagues, and mentors, who supported us in ways large and small throughout this journey, whether through a timely word of encouragement, a technical suggestion, or simply their belief that this work was worth finishing. This thesis would not have been possible without you. And finally, we are grateful to God for his grace throughout this masters program.

Table of Contents

Supervisor Approval	3
Acknowledgements.....	4
Table of Contents.....	5
List of Tables	11
List of Figures	12
List of Abbreviations	14
Abstract	16
Resumen.....	17
Executive Summary	18
1.1 General Objectives of the Study	33
1.2 Justification and Relevance of the Research Topic	34
1.2.1 Practical Industry Context.....	34
1.2.2 Academic Gaps and Research Motivation.....	35
1.2.3 SME Transferability: Data, Tools, and Technical Requirements.....	36
1.3 Methodology Overview.....	37
1.4 Structure of the Thesis.....	40
1.5 Theoretical and Conceptual Foundations.....	41
1.5.1 Classical Inventory Theory	41
1.5.2 Machine Learning and Ensemble Learning Theory	41
1.5.3 Supply Chain and Demand Uncertainty Theory	42
1.5.4 CRISP-DM as a Research Methodology Framework.....	42
1.6 Contributions of the Study.....	43
1.6.1 Methodological Contribution.....	43
1.6.2 Practical Contribution.....	43
1.6.3 Research Contribution	44
2.1 Introduction.....	46
2.2 Problem Framing: What the Literature Collectively Addresses.....	47
2.2.1 The Inadequacy of Traditional Forecasting Methods	47
2.2.2 The Challenge of Demand Volatility and Uncertainty	47

2.2.3 Integration of Forecasting with Inventory Policy	48
2.2.4 Scalability, Accessibility, and Context-Specificity	49
2.3 Theoretical Frameworks and Conceptual Foundations.....	50
2.3.1 Classical Inventory Theory	50
2.3.2 Ensemble Learning Theory.....	50
2.3.3 Machine Learning Methodology Frameworks.....	51
2.3.4 Supply Chain Theory	51
2.4 Data Sources, Characteristics, and Preprocessing	52
2.4.1 Benchmark Retail Datasets	52
2.4.2 Industry and Proprietary Datasets.....	52
2.4.3 Data Characteristics and Preprocessing	53
2.5 Forecasting Methodologies: A Layered Overview	55
2.5.1 Classical Statistical Methods (Baseline Tier)	55
2.5.2 Standard Machine Learning Algorithms.....	55
2.5.2.1 Tree-Based Methods	55
2.5.2.2 Support Vector Methods	56
2.5.2.3 Neural Networks and Basic Deep Learning.....	56
2.5.3 Deep Learning Architectures	57
2.5.3.1 LSTM and Recurrent Neural Networks.....	57
2.5.3.2 Advanced Deep Learning: Transformers, GRUs, and Probabilistic Networks	57
2.5.4 Ensemble and Hybrid Architectures	58
2.5.4.1 Gradient Boosting Sequential Ensembles.....	58
2.5.4.2 Heterogeneous Deep Learning Ensembles with Meta-Learning	59
2.5.4.3 Classical-ML Hybrid Frameworks	59
2.5.4.4 Reinforcement Learning for Inventory Policy	59
2.5.5 Specialised Applications and Extensions.....	60
2.6 Inventory Optimisation Frameworks and Operational Integration.....	60
2.6.1 The Order-Up-To-Level Policy	60
2.6.2 Safety Stock and Service Level Calibration	61
2.6.3 Simulation and Operational Validation	61
2.7 Evaluation Metrics and Comparative Results.....	62

2.7.1 Standard Accuracy Metrics.....	62
2.7.2 Comparative Performance Across Studies.....	62
2.7.3 Feature Importance and Key Demand Drivers	63
2.8 Research Gaps and Limitations.....	64
2.8.1 Single-Echelon Focus and Multi-Echelon Neglect.....	64
2.8.2 Point Forecasting Dominance and Probabilistic Neglect.....	64
2.8.3 Absence of Real-Time Deployment Validation.....	65
2.8.4 Exclusion of Shortage and Perishability Costs	65
2.8.5 SME and Mid-Market Accessibility	65
2.8.6 Ethical, Governance, and Sustainability Dimensions	66
2.8.7 Concept Drift and Model Maintenance.....	66
2.9 Directions for Future Research	66
2.10 Synthesis: Convergent Findings and Key Contributions	68
2.10.1 What the Literature Has Established.....	68
2.10.2 Where the Literature Converges and Diverges	68
2.10.3 Contribution of This Review	69
2.11 Conclusion	69
3.1 Introduction.....	71
3.2 Research Design and Methodological Framework	72
3.2.1 Research Philosophy	72
3.2.2 CRISP-DM as the Structuring Framework	72
3.3 Data Source: The M5-Forecasting (Walmart) Dataset	72
3.3.1 Dataset Description and Provenance	72
3.3.2 Justification of Dataset Choice	73
3.3.3 Stratified Subsample for Computational Tractability	73
3.4 Data Preprocessing Pipeline	74
3.4.1 Data Integration	74
3.4.2 Handling Missing Values.....	74
3.4.3 Handling Zero-Sale Records.....	75
3.4.4 Outlier Detection and Treatment	75
3.4.5 Categorical Variable Encoding.....	75

3.4.6 Target-Encoded Mean Features	75
3.5 Feature Engineering	76
3.5.1 Lag Features	76
3.5.2 Rolling Average Features	76
3.5.3 Calendar and Temporal Features	76
3.5.4 Promotional and Event Features	77
3.5.5 Price Features	77
3.5.6 Target-Encoded Mean Features	77
3.6 Model Development and Evaluation Framework	78
3.6.1 Model Selection Rationale	78
3.6.2 Training, Validation, and Test Split	80
3.6.3 Hyperparameter Optimisation	80
3.6.4 Evaluation Metrics	81
3.7 Inventory Policy Integration: The Order-Up-To-Level Framework	82
3.7.1 Policy Selection and Justification	82
3.7.2 Gaussian Policy Parameter Formulations (Primary)	83
3.7.3 Lead Time Scenarios	84
3.7.4 Robustness Check: Empirical Quantile Safety Stock	84
3.7.5 Stockout Cost Modelling	84
3.8 Power BI Dashboard Development	85
3.9 Sensitivity Analysis Design	86
3.10 Ethical Considerations and Model Transparency	86
3.11 SME Transferability Specification	87
3.12 Chapter Summary	88
4.1 Introduction	90
4.2 Implementation Environment	90
4.3 Data Preparation Summary	91
4.4 Tier 1: Classical Baseline Results	91
4.5 Tier 2: Machine Learning Model Results	92
4.6 Tier 3: LSTM Results	93
4.7 Tier 4: Stacking Ensemble Results	94

4.8 Tier 5: Quantile Forecasting Results.....	94
4.9 Consolidated Model Comparison	95
4.10 Performance Against Chapter 1 Targets	96
4.11 Model Interpretability — SHAP Analysis.....	98
4.12 Performance Fairness Across Categories and States	102
4.13 Discussion	103
4.14 Chapter Summary	104
5.1 Introduction.....	105
5.2 Inventory Policy and Simulation Design	106
5.3 Validation of the Simulation Engine	106
5.4 Central-Case Results	107
5.5 The Mechanism: Same Spread, Opposite Conclusion.....	108
5.6 The Boundary Condition: When the Quantile Policy Pays Off.....	109
5.7 Sensitivity Analysis	110
5.7.1 Stockout-cost multiplier	110
5.7.2 Lead time	111
5.8 Power BI Dashboard	112
5.9 Discussion	115
5.10 Limitations.....	115
5.11 Chapter Summary	116
6.1 Introduction.....	118
6.2 Synthesis of Empirical Findings	118
6.3 Revisiting the Four Research Gaps.....	119
6.3.1 Gap 1 — Disconnect between forecast accuracy and inventory decision-making.....	119
6.3.2 Gap 2 — Real-world deployment and operational validation	120
6.3.3 Gap 3 — Underrepresentation of SME retailers and practical constraints.....	120
6.3.4 Gap 4 — Integration of uncertainty in inventory calibration	121
6.4 Comparison with the Literature	122
6.5 Theoretical Contributions	123
6.6 Practical Contributions.....	124
6.7 Validity Threats	125

6.8 Extended Limitations	126
6.9 Implications for Future Research.....	127
6.10 Chapter Summary	127
7.1 Introduction.....	129
7.2 Restatement of the Research Objectives.....	129
7.3 Achievement Against Research Objectives	130
7.3.1 Objective 1 — Comparative evaluation of forecasting approaches.....	130
7.3.2 Objective 2 — Systematic feature engineering	130
7.3.3 Objective 3 — OUTL inventory integration.....	130
7.3.4 Objective 4 — Quantification of operational impact.....	131
7.3.5 Objective 5 — Power BI dashboard for SME deployment.....	132
7.4 Principal Contributions	132
7.5 Closing Remarks.....	133
References	123
A.1 Library Versions	128
A.2 Final Chosen Hyperparameters.....	129
A.3 OUTL and Economic Constants	130
A.4 Deliverables Inventory	130
A.5 How to Reproduce	131
A.6 Full Simulation Grid	132
A.7 Falsification Test — Per-Policy Detail	133

List of Tables

<i>Table 1: Summary of Key Datasets Used Across the Literature</i>	53
<i>Table 2: Key Forecasting Performance Results Across Selected Studies</i>	62
Table 3.1: Subsample composition by category and state	74
Table 3.2: Complete feature engineering summary	77
Table 3.3: Hyperparameter search spaces	81
Table 4.1: Implementation environment	90
Table 4.2: Tier 1 — Classical baseline results (test set)	92
Table 4.3: Tier 2 — Machine learning model results (test set)	93
Table 4.4: Tier 3 — LSTM results (test set)	93
Table 4.5: Tier 4 — Stacking ensemble results (test set)	94
Table 4.6: Tier 5 — Quantile forecast empirical coverage (test set)	95
Table 4.7: Consolidated cross-tier model comparison (sorted by MAE)	95
Table 4.8: Stacking ensemble (Tier 4) vs best classical baseline	97
Table 4.9: Top 15 features by mean absolute SHAP value	100
Table 4.10: Ensemble metrics by product category	102
Table 4.11: Ensemble metrics by state	102
<i>Table 5.1: Falsification test of the simulation engine (central case, $L = 14$)</i>	107
<i>Table 5.2: Simulated inventory outcome ($L = 14$, $SL = 95\%$, $m = 1.0$)</i>	107
<i>Table 5.3: Sensitivity to the stockout multiplier ($L = 14$, $SL = 95\%$)</i>	110
<i>Table 5.4: Sensitivity to lead time ($SL = 95\%$, $m = 1.0$)</i>	111
<i>Table A.1: Pinned library versions used in the experiments</i>	141
<i>Table A.2: Final hyperparameter values per model</i>	142
<i>Table A.3: OUTL and economic constants</i>	143
<i>Table A.4: Full forward-simulation grid (SL target = 95%)</i>	146
<i>Table A.5: Per-policy outcome under the demand $\times 3$ stress ($L = 14$, $SL = 95\%$, $m = 1.0$)</i>	149

List of Figures

Figure 4.1. Consolidated model comparison across all twelve forecasting approaches evaluated. Bars show MAE on the test set (lower is better), ordered from best to worst. The stacking ensemble (STACK_Ridge_trees) achieves the lowest MAE at 0.952, narrowly beating the three individual tree-based ML models. The LSTM and Seasonal Naive baselines are the weakest performers; Croston-SBA leads the classical baselines.

96

Figure 4.2. SHAP feature importance ranking for the LightGBM model (mean absolute SHAP values, computed on a 5,000-row stratified sample of the test set). The top five features — rolling_mean_7, rolling_mean_28, item_month_mean, rolling_std_7, and item_dow_mean — together account for 78.6% of model behaviour, confirming that engineered temporal statistics and target-encoded item-level seasonality are the dominant predictive drivers.

99

Figure 4.3. SHAP beeswarm plot showing the directional effect of each feature on individual predictions. Each point represents one observation; horizontal position shows the SHAP value (positive = higher predicted demand) and colour encodes the feature value (red = high, blue = low). The plot reveals that high recent demand (rolling_mean_7) and high item-specific historical levels push predictions upward, while is_snap_day and weekend indicators contribute meaningful but smaller signals.

100

Figure 4.4. Performance fairness analysis. WMAPE is reported by product category (left) and by state (right). Variation across categories (26.8% coefficient of variation) and across states (19.8%) reflects underlying demand structure (intermittence levels, demand volumes) rather than systematic algorithmic bias: no category or state performs more than $2\times$ worse than another on any metric. Foods, the highest-volume category, has the highest absolute MAE but the best WMAPE and Pred(10%).

101

Figure 5.1. Power BI dashboard — Forecast Overview page. KPI tiles show forecast-accuracy and realised-service-level summaries; the time series shows daily actual, ML-ensemble (Tier 4), and Croston-SBA classical forecasts over the test window for a selected SKU-store combination.

113

Figure 5.2. Power BI dashboard — Inventory Status page. Summary cards report mean safety stock, reorder point, order-up-to-level, and total annual cost at the user-selected lead time, service level, and policy. The per-SKU table is sorted by total cost descending with conditional formatting flagging the highest safety-stock items. The bar chart compares total annual cost across the three policies.

113

Figure 5.3. Power BI dashboard — Cost Analysis page. KPI tiles report the central-case simulated cost decomposition (ML-Empirical-Quantile policy, $L = 14$, $m = 1.0$): \$64,573 total annual cost, \$0 stockout cost, holding cost dominating. The table reports realised service level by policy (all policies tie at 100%), the matrix breaks total cost down by lead time $\times$ policy, and the bar chart compares cost composition across the three policies at $L = 14$. 114

Figure 5.4. Power BI dashboard — Sensitivity Analysis page. The line chart shows ML-Quantile cost reduction by lead time, sitting at approximately -1.5% across all four scenarios — well below the Chapter 1 5–15% target band (shown in green). The bar chart confirms the §5.7.1 invariance: cost change is identical across the three stockout-cost multipliers because realised stockouts are zero across all policies. The full sensitivity grid reports per-condition values, matching the appendix simulation results. 114

List of Abbreviations

Abbreviation	Definition
AI	Artificial Intelligence
ANN	Artificial Neural Network
ARIMA	Autoregressive Integrated Moving Average
BI	Business Intelligence (as in Power BI)
CNN	Convolutional Neural Network
CPI	Consumer Price Index
CRISP-DM	Cross-Industry Standard Process for Data Mining
CSV	Comma-Separated Values
DRL	Deep Reinforcement Learning
EOQ	Economic Order Quantity
ERP	Enterprise Resource Planning
GNN	Graph Neural Network
GPU	Graphics Processing Unit
HC	Holding Cost
KNN	K-Nearest Neighbours
KPI	Key Performance Indicator
L1 / L2	L1/L2 regularisation (Lasso/Ridge penalty terms)
LSTM	Long Short-Term Memory (neural network architecture)
M5	M5-Forecasting (Walmart) dataset / competition
MA	Moving Average
MAE	Mean Absolute Error
MAPE	Mean Absolute Percentage Error
ML	Machine Learning
MLP	Multi-Layer Perceptron
MSE	Mean Squared Error
OC	Ordering Cost
OUTL	Order-Up-To-Level (inventory review policy)
P10 / P50 / P90 / P95 / P99	10th / 50th / 90th / 95th / 99th percentile of the forecast demand distribution
POS	Point-of-Sale
RAM	Random Access Memory
RF	Random Forest
RMDN	Recurrent Mixture Density Network
RMSE	Root Mean Squared Error
SARIMA / SARIMAX	Seasonal ARIMA / Seasonal ARIMA with eXogenous variables
SBA	Syntetos–Boylan Approximation (as in Croston-SBA)
SC	Stockout Cost
SES	Simple Exponential Smoothing
SHAP	SHapley Additive exPlanations
SKU	Stock-Keeping Unit
SL	Service Level
SME	Small and Medium-sized Enterprise
SNAP	Supplemental Nutrition Assistance Program
SS	Safety Stock
SVM / SVR	Support Vector Machine / Support Vector Regression
TC	Total Cost
TFT	Temporal Fusion Transformer

WMAPE	Weighted Mean Absolute Percentage Error
--------------	---

Abstract

Retail demand forecasting and inventory optimisation are central to operations management, yet the operational value of machine-learning forecasting is more often asserted than measured. This thesis develops and evaluates an end-to-end forecasting-to-inventory framework on the M5-Forecasting (Walmart) dataset, a stratified subsample of 502 product–store series of daily, intermittent SKU-level demand, comparing classical statistical baselines, gradient-boosting and deep-learning models, a stacking ensemble, and LightGBM quantile regression, and integrating their outputs with a periodic (R, s, S) Order-Up-To-Level inventory policy. On forecast accuracy, the stacking ensemble is the most accurate model but improves on the strongest classical baseline by only about 4.5% (MAE), well below the 10–25% reported in prior work. Critically, when inventory outcomes are measured by forward simulation over the real test window rather than assumed, all policies achieve full service with no stockouts, so neither the more accurate point forecast nor a quantile-based safety-stock reformulation lowers inventory cost: the apparent saving reported under assumed service levels does not survive measurement. The thesis interprets this as a finding rather than a failure. The 10–25% and 5–15% benchmarks of the literature were obtained on aggregated data and do not transfer to SKU-daily intermittent demand, and the operational value of ML forecasting is conditional on the demand regime, material only where the baseline policy faces genuine stockout risk. The contribution is a rigorous, reproducible, simulation-based characterisation of when ML forecasting and distributional safety stock affect inventory outcomes, delivered through an accessible Power BI dashboard designed for SME retailers.

Keywords: demand forecasting; machine learning; inventory optimisation; intermittent demand; inventory simulation; safety stock; quantile regression; Order-Up-To-Level (R, s, S) policy.

Resumen

La previsión de la demanda y la optimización de inventario son centrales en la gestión de operaciones del comercio minorista, pero el valor operativo de la previsión basada en aprendizaje automático se suele afirmar más que medir. Esta tesis desarrolla y evalúa un marco integral de previsión e inventario sobre el conjunto de datos M5-Forecasting (Walmart), una submuestra estratificada de 502 series producto–tienda de demanda diaria, intermitente y a nivel de SKU, comparando modelos estadísticos clásicos, modelos de gradient boosting y aprendizaje profundo, un ensemble por apilamiento (stacking) y regresión cuantílica con LightGBM, e integrando sus salidas con una política periódica (R, s, S) de nivel de reposición (Order-Up-To-Level). En precisión de previsión, el ensemble es el modelo más preciso, pero mejora al mejor modelo clásico en apenas un 4,5% (MAE), muy por debajo del 10–25% reportado en trabajos previos. De forma decisiva, cuando los resultados de inventario se miden mediante simulación hacia delante sobre la ventana de test real, en lugar de asumirse, todas las políticas alcanzan servicio completo sin roturas de stock, de modo que ni la mayor precisión de la previsión ni la reformulación cuantílica del stock de seguridad reducen el coste de inventario: el ahorro aparente obtenido bajo niveles de servicio asumidos no sobrevive a la medición. La tesis interpreta esto como un hallazgo, no como un fracaso. Los umbrales del 10–25% y del 5–15% de la literatura se obtuvieron sobre datos agregados y no se trasladan a la demanda intermitente a nivel SKU-día, y el valor operativo de la previsión con ML es condicional al régimen de demanda, relevante solo cuando la política base se expone a un riesgo real de rotura. La contribución es una caracterización rigurosa, reproducible y basada en simulación de cuándo la previsión con ML y el stock de seguridad distribucional afectan a los resultados de inventario, presentada mediante un cuadro de mando accesible en Power BI diseñado para pymes minoristas.

Palabras clave: previsión de la demanda; aprendizaje automático; optimización de inventario; demanda intermitente; simulación de inventario; stock de seguridad; regresión cuantílica; política Order-Up-To-Level (R, s, S).

Executive Summary

1. Purpose of the Thesis

This thesis was undertaken because the operational value of machine learning-based demand forecasting is, in retail practice and in much of the academic literature, more often asserted than measured. Published studies routinely report accuracy improvements of 10-25 percent and inventory cost reductions of 5-15 percent, but these figures are typically obtained on data aggregated to the weekly or category level, evaluated analytically rather than through direct simulation, and rarely tested against the actual conditions under which a retailer places an order: a single product, at a single store, on a single day. The purpose of this thesis is to close that gap; to build a complete, reproducible forecasting-to-inventory pipeline and to test, rigorously and honestly, whether those published benefits survive contact with real, SKU-daily, intermittent retail demand, and with a decision-support tool that resource-constrained organisations can actually use.

To do this, the thesis compares twelve forecasting models spanning classical statistics, machine learning, deep learning, and ensemble methods; integrates the best-performing forecast with a standard periodic inventory policy; and, critically, measures the resulting inventory performance by simulating that policy forward, day by day, against real held-out demand rather than assuming an outcome analytically. None of that technical depth is necessary to understand the conclusion, which can be stated in one sentence: the value of machine learning forecasting for retail inventory management is real, but it is not automatic, and this thesis provides a tool that tells an organisation, before it spends money, whether its own situation is one where the investment is likely to pay off.

This summary is written for the people who will ultimately decide whether, and how, to invest in this kind of work, retail operations managers, finance leads, and SME owners, as well as the academic committee evaluating this thesis, and is organised to mirror the questions a decision-maker actually asks: What decision am I facing? Why is this problem hard? What does the existing evidence say? How was this specific research conducted, and can I trust it? What did it find? And, most importantly, what should I do with that information?

1.1 Objectives of the Study

Five objectives guided the research: (1) comparatively evaluate classical statistical, machine learning, deep learning, ensemble, and probabilistic forecasting approaches on a hierarchical, multi-product, multi-store retail benchmark to identify the most accurate method; (2) develop a

systematic feature-engineering strategy capturing recent sales trends, seasonality, promotions, and price signals that is replicable in SME retail settings with comparable data; (3) integrate the best-performing forecast with the Order-Up-To-Level inventory policy to compute safety stock, reorder points, and order-up-to levels at a target service level; (4) directly test whether the accuracy and cost-reduction benchmarks reported in closely related published studies transfer to the SKU-daily, intermittent-demand granularity at which real inventory decisions are made, treating any shortfall as a finding rather than a target to be forced; and (5) deliver a Power BI dashboard that translates the technical outputs into an accessible, interactive decision-support tool for resource-constrained retail organisations.

Objective 4 deserves particular emphasis for a decision-making audience: unlike a typical applied study that sets out to demonstrate a benefit, this research was explicitly designed to test whether a claimed benefit actually holds at the level of granularity where a retailer places real orders. That design choice is what allows Section 6 to report an honest result rather than a predetermined one.

2. The Business Decision at Stake

Every retailer, from a single-location shop to a national chain, makes the same recurring decision: how much stock to hold, and when to reorder it. Get the answer wrong in one direction and the business runs out of product customers want to buy; a stockout; losing the sale immediately and, if it happens repeatedly, losing the customer's future business as well. Get it wrong in the other direction and the business ties up cash in inventory that sits on shelves or in warehouses, incurring storage cost, spoilage risk for perishables, and the opportunity cost of capital that could have been used elsewhere.

Large retailers such as Walmart address this problem with dedicated data science teams, enterprise forecasting software, and analysts who spend their careers tuning inventory parameters. Small and medium-sized retailers typically have none of this. They rely on the experience and intuition of a manager, on simple spreadsheet rules, or on the basic reorder logic built into off-the-shelf point-of-sale software. This is not a criticism of SME operators, it reflects a genuine resource gap. Data science talent, enterprise software licences, and cloud computing budgets are simply not accessible to a business running on thin retail margins.

In the last decade, a large and growing body of academic and industry research has argued that machine learning can close this gap: that algorithms trained on a retailer's own sales history can forecast demand more accurately than traditional statistical methods, and that these more accurate forecasts can be converted into inventory policies that reduce both stockouts and excess

stock simultaneously. Vendors selling forecasting software, and many academic papers, report double-digit percentage improvements forecast accuracy gains of 10 to 25 percent, and inventory cost reductions of 5 to 15 percent, relative to the classical methods that most SMEs currently use.

This creates a genuine decision for any retail operator: should the organisation invest time, money, and change-management effort into adopting a machine learning forecasting and inventory system? The honest answer, prior to this thesis, was difficult to give with confidence, for a reason explained in detail in Section 4 below: most of the evidence cited in favour of such an investment was generated under conditions, aggregated weekly or category-level sales data, and inventory calculations that assumed rather than measured their own performance, that do not match how a typical SME actually operates or how the resulting software would actually be used, at the day-by-day, product-by-product level.

This thesis was designed to close that evidence gap. It builds a complete, end-to-end forecasting-to-inventory pipeline using the same kind of data an SME already collects, tests it under conditions that match real day-to-day retail operation, and, critically, measures the resulting inventory performance by actually running the inventory policy forward through real demand data, rather than assuming its performance analytically. The result is a decision-support tool, delivered as an interactive dashboard, and a specific, evidence-based answer to the investment question above: not a blanket yes or no, but a diagnostic that tells an individual organisation whether its own demand pattern makes the investment worthwhile.

3. The Nature of the Problem

To understand why this is a genuinely hard problem, and why the answer turned out to be more nuanced than the marketing literature suggests, it helps to unpack what "demand forecasting and inventory optimisation" actually involves at the level of detail a retail operation experiences it.

3.1 Demand Is Volatile and, at the Product Level, Highly Intermittent

Aggregate retail demand, total sales across a whole store or category, summed over a week or a month, tends to be relatively smooth and predictable. But inventory decisions are not made at that aggregate level. They are made product by product, store by store, day by day: how many units of this specific item, at this specific location, will sell tomorrow? At that level of detail, demand for most retail products is intermittent. Many days see zero sales of a given item at a given store, punctuated by occasional spikes driven by promotions, paydays, holidays, or simple randomness. This intermittency is the single most important technical fact about retail demand, and it is the

reason forecasting accuracy figures that sound impressive in academic papers often do not describe the reality of day-to-day, item-level operations.

3.2 The Forecast Is Only Half the Problem

A more accurate demand forecast is not, by itself, worth anything to a business. Its value is realised only when it is translated into an inventory decision, how much safety stock to hold, when to trigger a reorder, and how much to order. Most published research on retail forecasting stops at the forecast: it reports how close the model's predictions were to actual sales, using statistical accuracy metrics, and stops there. Very little published research goes the extra step of actually running the resulting inventory policy and measuring what happens to costs and stockouts. This thesis identifies this disconnect as the first and most consequential gap in the existing literature, and it is the gap that this research is primarily designed to close.

3.3 The Textbook Safety-Stock Formula Was Built for a Different Kind of Demand

The standard method used across industry and academia to calculate how much buffer stock ("safety stock") to hold assumes that demand forecast errors follow a symmetric, bell-shaped statistical distribution, the same assumption underlying most classical statistics. Intermittent, spiky retail demand of the kind described in Section 3.1 does not follow this pattern; it is skewed, with a small number of large demand spikes and a large number of zero- or near-zero demand days. This mismatch between the assumption built into the textbook formula and the actual shape of real retail demand is a second identified gap in current practice, and one of the central technical questions this thesis investigates: does correcting this mismatch, replacing the textbook formula with one that reflects the true shape of the demand distribution, actually save money in practice?

3.4 SMEs Face a Structural Resource Disadvantage

Even where the technical solutions above are well understood, SME retailers face a separate and equally real barrier: access. Enterprise forecasting software, cloud computing infrastructure, and in-house data science expertise are typically priced and designed for organisations far larger than a typical SME. This is the third gap the thesis addresses directly, through an explicit design commitment to use only free, open-source software, to require no computing hardware beyond a standard laptop, and to require no data that a typical retail point-of-sale system does not already collect.

3.5 Point Forecasts Hide the Uncertainty That Inventory Decisions Actually Depend On

Finally, most forecasting systems, including many marketed commercially, produce a single "best guess" number for tomorrow's demand. But the inventory decision does not depend on the best

guess; it depends on the range of plausible outcomes and, specifically, on how bad the worst realistic case could be, because that is what safety stock is meant to protect against. A single point forecast cannot answer that question. This is the fourth gap the thesis addresses: incorporating the full range of forecast uncertainty, not just a single predicted number, into the inventory calculation.

4. What Is Already Known: Insights from the Research Literature

This thesis is grounded in a structured review of 52 academic sources spanning classical statistical forecasting, machine learning, deep learning, ensemble methods, and inventory theory. Three consistent findings emerge from that body of work, alongside the four gaps already introduced in Section 3.

First, there is strong and consistent evidence that machine learning and ensemble methods, models that combine the predictions of several algorithms, outperform traditional statistical forecasting methods (moving averages, exponential smoothing, ARIMA-family models) across a wide range of retail datasets and sectors. This finding is well replicated and is not in serious dispute in the literature.

Second, published studies report substantial operational benefits from this accuracy improvement: forecast accuracy gains in the 10-25 percent range and inventory cost reductions of 5-15 percent relative to classical baselines are commonly cited figures, drawn from studies such as Barghi (2025) and Seyedan, Mafakheri and Wang (2023). These are the benchmark figures against which this thesis explicitly tests its own results, and they are the figures most frequently used to justify commercial investment in forecasting technology.

Third, and this is the critical qualification that a decision-maker needs to understand before acting on the figures in the paragraph above, the studies reporting these benchmark figures typically evaluate their models on data aggregated to the weekly or product-category level, not at the daily, individual-product level at which actual inventory reorder decisions are made. This distinction matters enormously. Aggregation smooths out the intermittency described in Section 3.1: a category's total weekly sales are far more predictable than one specific product's sales on one specific day. A forecasting improvement measured on smoothed, aggregated data does not automatically apply at the granularity where inventory decisions actually get made.

This thesis was explicitly designed to test whether the published benchmark figures, 10-25 percent accuracy improvement, 5-15 percent cost reduction, hold up when the analysis is done at the SKU-daily level, using real product-store-day sales records rather than aggregated data,

and using inventory outcomes that are directly measured through simulation rather than estimated analytically. This is a test of transferability, not a repetition of prior work, and Section 6 reports what that test found.

The four research gaps this thesis addresses can be summarised as follows.

Gap	The Problem	What This Thesis Does About It
1. Forecast-to-decision disconnect	Most studies report forecast accuracy and stop there, without measuring the resulting inventory cost or service outcome.	Builds the full pipeline from forecast to inventory policy, and measures the resulting cost and service level directly by simulation.
2. Limited real-world validation	Most studies validate against historical data only, without testing how the resulting policy behaves when actually run forward in time.	Runs the inventory policy forward, day by day, over real test-period demand, and includes a dedicated engine-validation exercise (Section 5.4).
3. SME underrepresentation	Most research and commercial tooling assumes enterprise-scale infrastructure and data science staffing that SMEs do not have.	Uses only free, open-source tools; runs on a standard laptop; requires only data an SME point-of-sale system already produces.
4. Uncertainty ignored in inventory calibration	Standard safety-stock formulas assume a single point forecast and a symmetric error distribution that intermittent retail demand does not follow.	Builds a probabilistic (quantile) forecasting model and tests an inventory formula based on the real shape of the demand distribution, not an assumed one.

5. Research Methodology and Approach

This section summarises how the research was conducted, written for a reader who wants to understand the credibility and rigour of the approach without needing the full technical specification found in the thesis itself.

5.1 The Data

The research uses the M5-Forecasting (Walmart) dataset, a large, publicly available benchmark comprising daily unit sales for 3,049 products across ten Walmart stores in three US states, spanning more than five years (2011-2016), together with calendar information (day of week, holidays, promotional events, SNAP benefit-eligibility days) and weekly pricing data. This dataset was chosen because it is large-scale, hierarchically structured in a way that mirrors real SME retail operations, and because it is a recognised academic benchmark, enabling direct comparison with the published studies discussed in Section 4. For computational practicality, a representative sample of 502 product-store combinations was selected, proportionally balanced across product categories and states, so that the results generalise across the different types of product and geography represented in the full dataset.

5.2 Building the Forecasting Models

Rather than building and testing a single forecasting model, the research constructed and compared five tiers of increasingly sophisticated approaches, so that any claimed benefit from machine learning could be measured against a full spectrum of alternatives, not just a single, potentially weak, baseline.

Classical baselines. The six classical statistical methods most commonly used in retail practice today, including moving averages and a specialised method (Croston's method) designed specifically for the intermittent, gap-filled demand patterns described in Section 3.1.

Standard machine learning. Three widely used machine learning algorithms (Random Forest, XGBoost, and LightGBM), each trained on a rich set of engineered features capturing recent sales trends, calendar effects, promotional signals, and pricing information.

Deep learning. A deep learning model (a Long Short-Term Memory neural network) capable of learning complex temporal patterns directly from sequences of past sales.

Ensemble model. A combination model that blends the predictions of the three machine learning models above through a statistical weighting layer, representing the state-of-the-art approach reported in the academic literature.

Probabilistic model. A specialised model that forecasts not just a single expected demand figure but a full range of plausible outcomes (from the 10th to the 99th percentile), directly addressing Gap 4 from Section 4.

All models were trained strictly on historical data and evaluated exclusively on a later, entirely unseen eight-week test period, so that the reported results reflect genuine out-of-sample forecasting performance, the same standard a business would need to trust before deploying a model live.

5.3 Translating Forecasts into Inventory Decisions

The forecasting outputs were integrated with a standard, industry-recognised inventory control policy (the Order-Up-To-Level, or periodic review, policy), which specifies how much buffer stock to hold, at what inventory level to trigger a reorder, and how much to order, based on a target customer service level (the percentage of demand the business commits to fulfilling from stock on hand). Three variants of this policy were compared, differing only in how the buffer stock (safety stock) is calculated: one using the classical forecast with the textbook safety-stock formula, one using the machine learning forecast with the same textbook formula, and one using the machine

learning forecast with the alternative, distribution-aware formula described in Section 3.3. This design isolates two separate questions that are usually conflated in commercial claims: does a better forecast help, and does a better safety-stock formula help, independently of one another?

5.4 The Critical Methodological Decision: Measuring, Not Assuming, Inventory Performance

This is the single most important methodological choice in the research, and the one that ultimately determined the honesty and credibility of its conclusions. An earlier draft of this research, like most of the published literature it reviews, calculated the expected inventory cost and service level of each policy analytically, using a mathematical formula to project what should happen, based on assumptions about how much of the time the policy would run short of stock. On the advice of the academic supervisor overseeing this work, that approach was replaced with something more rigorous: each inventory policy was run forward, day by day, against the real 56-day test-period demand for every one of the 502 products studied, genuinely tracking stock levels, genuinely recording every day on which a product ran out, and genuinely calculating the resulting costs from what actually happened rather than from what a formula predicted should happen.

This distinction, between assuming performance and measuring it, is exactly the kind of rigour a business should demand before trusting any forecasting or inventory vendor's claims, and it is precisely this rigour that produced the thesis's most important and most commercially relevant finding, reported in Section 6 below. To confirm that this simulation was itself trustworthy, that it would correctly detect a stockout if one occurred, rather than silently failing to register it, the research includes a dedicated validation exercise in which the simulation was deliberately stress-tested under conditions engineered to force stockouts (removing all safety stock, tripling demand, and cutting off resupply entirely). In every stress test, the simulation correctly detected the resulting service failures and calculated the associated cost, confirming that the simulation engine is measuring real outcomes and can be trusted.

5.5 Building the Decision-Support Dashboard

Finally, the entire analysis was translated into an interactive dashboard, built using Microsoft Power BI's free tier, organised into four pages: a forecast overview showing predicted versus actual demand for any selected product; an inventory status page showing safety stock, reorder point, and order-up-to level recommendations; a cost analysis page breaking down the cost of holding, ordering, and stocking-out for each policy; and a sensitivity page allowing a manager to explore how the results change under different assumptions about lead time and the cost of a

stockout. This dashboard is the practical deliverable of the research: a tool a non-technical manager can open and use, without needing to understand the underlying statistics.

6. What We Found

The findings fall into two parts, corresponding to the two separate questions identified in Section 5.3: what happened to forecast accuracy, and what happened to inventory cost. The two answers turned out to be quite different from each other, and that difference is itself the thesis's central finding.

6.1 Forecast Accuracy: A Real but Modest Improvement

The ensemble model, the combination of the three machine learning algorithms, was the most accurate forecaster among all twelve models tested. It improved on the best classical method by 4.5 percent on average forecasting error and by 9.5 percent on a metric that penalises large errors more heavily (relevant for promotional spikes and other high-impact events). It did not, however, improve meaningfully on the specific percentage-error metric (MAPE) most commonly quoted in the marketing literature, and it therefore fell short of the 10-25 percent improvement range reported in the published studies discussed in Section 4.

This shortfall is not a failure of the modelling work; it is itself an important and honest finding. As anticipated in Section 4, the published 10-25 percent figures were obtained on data aggregated to the weekly or category level. At the daily, individual-product level, which is the granularity at which real inventory decisions actually get made, demand is far more erratic, and even the best available machine learning model cannot close that gap to the same degree. A retail decision-maker evaluating any forecasting technology, whether from this thesis or from a commercial vendor, should specifically ask at what level of data aggregation quoted accuracy figures were measured, because the answer materially changes what improvement should realistically be expected in day-to-day operation.

6.2 Inventory Cost: The More Important, and More Surprising, Finding

This is where the rigorous, simulation-based measurement described in Section 5.4 produced a result that a purely analytical, assumption-based study would never have revealed. When the three inventory policies were run forward through real demand, all three, the classical policy, the machine-learning-with-standard-formula policy, and the machine-learning-with-distribution-aware-formula policy, achieved a 100 percent realised service level across the test period. In plain terms: none of the three policies ever ran out of stock, for any of the 502 products, at any point during the eight-week test window.

The practical consequence of this is significant and somewhat counter-intuitive: because no policy ever experienced a stockout, none of the three policies had any stockout cost to save, and the more accurate forecast and the more sophisticated safety-stock formula therefore had nothing to improve upon in terms of avoided lost sales. The machine-learning policy using the standard safety-stock formula performed essentially identically to the classical policy. The two differed in total annual cost by less than 0.05 percent, an amount too small to be commercially meaningful. The machine-learning policy using the more sophisticated, distribution-aware safety-stock formula actually cost approximately 1.5 percent more than the classical policy, because it allocated a larger safety-stock buffer, appropriately reflecting the true, skewed shape of the demand distribution, but that extra buffer was protecting against stockouts that, in this specific demand regime, simply did not occur.

This is the central, headline finding of the thesis, and it directly overturns an earlier, less rigorous version of this same analysis, which, using the conventional analytical approach most of the published literature relies on, had estimated a 21 percent inventory cost saving from adopting the machine-learning and distribution-aware approach. That 21 percent figure was an artefact of an assumption embedded in the analytical formula, not a measured outcome; once the same policies were actually run against real demand, the saving disappeared. This is precisely the kind of gap between claimed and actual performance that a decision-maker should be alert to when evaluating any forecasting or inventory technology that has not been validated through genuine simulation or a live pilot.

6.3 Why This Happened, in Plain Terms

The explanation is straightforward once stated plainly. The classical inventory policy, using the simplest available forecasting method, was already holding enough buffer stock to meet demand in full over the test period. When a policy is already meeting 100 percent of demand, there is, by definition, no stockout cost left to eliminate, and therefore no room for a more sophisticated forecast or a more sophisticated safety-stock formula to demonstrate a cost advantage through avoided stockouts. The additional safety stock held by the more sophisticated policy becomes pure additional carrying cost, with no offsetting benefit.

This does not mean machine learning forecasting or distribution-aware safety stock are without value in general. It means their value depends on the specific demand and inventory situation an organisation is in. Where a business's current, simpler forecasting and inventory approach is already adequate, where stockouts are rare under the existing policy, a more sophisticated approach is unlikely to save meaningful money, and may in fact add small unnecessary holding

costs. Where a business's current approach is genuinely exposed to stockout risk, running short of stock with meaningful frequency under existing rules, the same sophisticated approach that showed no benefit in this test window would be expected to show a real one, because there would be actual stockout cost available to reduce. The research includes supporting evidence for this: when the test-period demand was deliberately increased threefold in the validation exercise described in Section 5.4, all three policies began experiencing real stockouts, and while the distribution-aware policy then performed on par with the classical one rather than worse, even a threefold demand shock was not sufficient, in this specific dataset, to make the more sophisticated approach clearly superior. This indicates that the conditions under which the more complex approach pays off are more specific and more stringent than the general marketing narrative around machine learning forecasting suggests.

6.4 Which Products and Which Model Features Mattered Most

Two further findings are of direct practical relevance to any retailer considering this kind of investment. First, an interpretability analysis (using a technique called SHAP, which quantifies how much each input variable contributed to each individual forecast) showed that the best-performing model relied overwhelmingly on straightforward, intuitive signals. Recent sales trends and each product's typical seasonal pattern accounted for nearly 80 percent of the model's predictive power. This is a reassuring finding for a non-technical stakeholder: the model is not relying on opaque or hard-to-explain logic, but on the same kind of signals an experienced store manager already intuitively tracks, simply applied more systematically and consistently across every product and every day.

Second, a fairness check confirmed that forecasting performance did not systematically favour or disadvantage any particular product category or store location beyond what would be expected from underlying differences in how predictable different types of products naturally are. This matters for any retailer operating a diverse product range or multiple locations: the model is not silently underperforming for a subset of the business.

7. What This Means for Decision-Makers: Central Recommendations

The findings in Section 6 translate into a specific, actionable set of recommendations, organised by the type of decision a reader of this summary is likely to be facing.

7.1 Do Not Invest Based on Published Percentage Improvements Alone

The single most important recommendation from this research is a note of caution about how to read the forecasting and inventory-optimisation literature, and by extension, vendor marketing

material that draws on it. Reported accuracy improvements of 10-25 percent and cost reductions of 5-15 percent are real, published findings, but they were typically measured under conditions, aggregated data, and in many cases analytically assumed rather than measured inventory outcomes, that do not automatically transfer to the day-to-day, product-level reality of most retail operations. Any organisation evaluating a forecasting or inventory technology investment should ask two specific questions of any vendor or study: at what level of data aggregation was the accuracy figure measured, and was the claimed cost saving measured by actually running the resulting policy against real demand, or calculated from a formula?

7.2 Test Your Own Situation Before Investing — Use the Diagnostic Delivered by This Research

The most practically useful output of this thesis, for an SME decision-maker specifically, is not a single number but a reusable diagnostic tool: the forward-simulation engine built and validated in this research. Before committing budget to a more sophisticated forecasting and inventory system, an organisation can run this same simulation approach against its own historical sales data to answer the one question that actually determines whether the investment will pay off: under our current, simpler approach, how often do we actually run out of stock, and how much does that cost us? If the answer is "rarely, and not much," as it was in this specific test dataset, the additional investment in machine learning forecasting and sophisticated safety-stock formulas is unlikely to be cost-justified, and resources are better spent elsewhere. If the answer is "frequently, and at real cost," which is a common situation for fast-moving, high-margin, or highly seasonal product categories, the same evidence base built in this thesis indicates the investment is likely to be worthwhile.

This reframes the investment decision from a generic "should we adopt AI" question into a specific, testable, low-cost diagnostic exercise that any organisation with basic sales records can run before spending anything further.

7.3 Where the Investment Is Most Likely to Pay Off

Based on the mechanism identified in Section 6.3, the organisations and product categories most likely to see a genuine return from machine learning forecasting and distribution-aware safety stock are those where: existing stockouts are frequent or costly (high-margin items, fast fashion, or products where a stockout risks losing the customer to a competitor); demand is highly seasonal, promotional, or volatile, making the current, simpler forecasting approach more likely to be inadequate; and service-level targets are deliberately set high (99 percent or above), a regime in which the cost of holding enough buffer stock to meet demand in full is greatest, so the

case for distribution-aware safety stock is correspondingly stronger where genuine stockout risk is present (this study did not itself sweep service level, so this is a reasoned extension of the boundary condition in Section 6.3, not a measured result). Conversely, stable, slow-moving, low-margin categories where the current approach is already adequate are less promising candidates for this specific investment.

7.4 Deploy the Dashboard as a Low-Cost, Low-Risk First Step

Regardless of the outcome of the diagnostic in Section 7.2, the Power BI dashboard delivered alongside this research is a low-cost, low-risk way to begin building forecasting and inventory-management capability. It requires no paid software licence, runs on standard office computing hardware, and requires only the sales, calendar, and pricing data that a typical point-of-sale system already produces. Adopting it does not commit an organisation to the more expensive machine-learning-plus-quantile-forecasting approach; it provides visibility into current inventory parameters and cost drivers regardless of which forecasting method ultimately feeds it, and it can be upgraded to more sophisticated forecasting later if the diagnostic in Section 7.2 indicates that upgrade would pay off.

7.5 Treat This as an Ongoing Measurement Discipline, Not a One-Time Decision

Demand regimes change: a product category that shows no stockout risk today may become promotional, seasonal, or supply-constrained in the future, changing the answer to the diagnostic question in Section 7.2. The recommendation is not a single, one-time investment decision but an ongoing practice: periodically re-running the simulation-based diagnostic against current data, particularly ahead of known high-stakes periods (major promotional events, seasonal peaks, new product launches, or supply chain disruptions), so that the decision to invest further in forecasting sophistication is always grounded in current, measured evidence rather than a decision made once and never revisited.

8. A Practical Path Forward

For an organisation persuaded by the case above, the following is a realistic, low-risk sequence for putting this research into practice, requiring no specialist data science hires and no capital software investment.

Step	Action	Resource Required
1	Export 12-24 months of daily sales, calendar/promotional, and price data from the existing point-of-sale system.	One analyst; standard POS export functionality

2	Run the open-source diagnostic pipeline delivered with this research against that data to identify which product categories currently experience meaningful stockout risk.	One laptop (16 GB RAM, no specialist hardware); free, open-source software
3	For categories showing low stockout risk, retain the existing, simpler forecasting and reorder approach — no further investment is indicated.	No additional resource
4	For categories showing meaningful stockout risk, pilot the machine-learning forecasting and distribution-aware safety-stock approach on that subset only, before considering a wider rollout.	Continued use of the same open-source pipeline; no new licences
5	Deploy the Power BI dashboard for ongoing visibility into forecast accuracy, inventory parameters, and cost drivers across the business.	Power BI Desktop or Power BI Service free tier
6	Re-run the diagnostic periodically (for example, quarterly, or ahead of major seasonal or promotional periods) to keep the investment decision grounded in current data.	Recurring, low time cost using the same free pipeline

This sequence is deliberately designed so that no organisation needs to commit significant budget before it has direct, measured evidence, from its own data, not from a published benchmark, that the investment is likely to be worthwhile.

9. Conclusion

This research set out to answer a question every retail decision-maker eventually faces: does investing in machine learning-based demand forecasting and inventory optimisation actually pay off, and is it accessible to an organisation without enterprise-scale resources? The answer this thesis provides is more useful, and more honest, than a simple yes: the value of the investment is real but conditional, and this research delivers both the evidence for that conclusion and a practical, low-cost tool that lets any organisation test which side of that condition it falls on, using its own data, before spending anything further.

The technology built to demonstrate this, a full forecasting and inventory pipeline, a validated simulation engine, and an accessible dashboard, is delivered entirely in free, open-source form, runs on standard office hardware, and requires only the data a typical point-of-sale system already produces. The barrier to entry this research removes is not financial; it is informational. What most retail decision-makers have lacked is not access to machine learning tools, but a trustworthy, rigorously validated way of knowing whether those tools are worth adopting for their specific business. This thesis provides that.

CHAPTER 1

Introduction

Retail organisations operating in contemporary markets face a persistent and costly operational dilemma: how to align stock levels with customer demand when that demand is inherently volatile, shaped by promotions, seasonal patterns, pricing decisions, public holidays, and broader macroeconomic conditions. Underestimating demand results in stockouts, immediate revenue loss, and erosion of customer trust. Overestimating demand ties up working capital in excess inventory, increases holding costs, and in the case of perishable goods generates significant waste with both financial and environmental consequences. This tension between under-supply and over-supply places demand forecasting and inventory optimisation at the centre of modern retail operations management.

Traditional approaches to this problem, including Moving Average (MA) methods, Exponential Smoothing, Autoregressive Integrated Moving Average (ARIMA), and its seasonal extension SARIMA have long formed the backbone of retail forecasting practice. However, these methods carry fundamental limitations in environments characterised by non-linear demand dynamics, multiple interacting variables, and abrupt structural shifts driven by promotions or external shocks. A growing body of academic research demonstrates that machine learning approaches, and ensemble and hybrid architectures in particular, substantially outperform classical statistical methods across diverse retail contexts (Barghi, 2025; Nasser et al., 2023; Mitra et al., 2022; Razzak et al., 2025). These models can simultaneously process historical sales data, promotional calendars, temporal indicators, and external factors to learn complex demand patterns that classical models cannot represent.

This thesis presents the design, implementation, and evaluation of a machine learning-based framework for retail demand forecasting and inventory optimisation. Using the M5-Forecasting (Walmart) dataset; a publicly available, large-scale hierarchical dataset comprising daily sales records for 3,049 products across 10 Walmart stores in the United States over five years, enriched with calendar and price information. This study develops and compares multiple forecasting approaches and translates their outputs into operational inventory recommendations through the Order-Up-To-Level (OUTL) policy framework. Although the M5 dataset originates from a large retail enterprise, the methodology and decision framework developed in this thesis are explicitly

designed and evaluated for transferability to small and medium-sized retailers (SMEs) who lack enterprise-scale analytics infrastructure but face identical demand planning challenges. The full pipeline is presented through an interactive Power BI dashboard, offering a practical and accessible tool for inventory decision-makers.

The remainder of this chapter is structured as follows. Section 1.1 states the general objectives of the research. Section 1.2 provides the justification and relevance of the study, grounding it in both practical industry challenges and identified gaps in the academic literature. Section 1.3 presents an overview of the research methodology. Section 1.4 maps the structure of the thesis. Section 1.5 outlines the theoretical and conceptual foundations. Section 1.6 articulates the contributions of the study.

1.1 General Objectives of the Study

The primary objective of this research is to develop and evaluate a data-driven framework that integrates machine learning-based demand forecasting with the Order-Up-To-Level (OUTL) inventory policy, producing measurable improvements in operational efficiency for retail businesses. The study will use the full M5-Forecasting (Walmart) dataset as an empirical benchmark and design the resulting framework to be applicable in SME retail environments. Specifically, the research pursues five objectives:

1. Implement and comparatively evaluate a range of forecasting models including classical statistical baselines (ARIMA, Exponential Smoothing, Moving Average), standard machine learning algorithms (Random Forest, XGBoost, LightGBM), time-series deep learning models (LSTM), and hybrid or ensemble architectures to identify the most accurate approach for multi-product, multi-store retail demand forecasting on the M5 dataset.
2. Develop a systematic data preparation and feature engineering pipeline incorporating lag variables, rolling averages, promotional indicators, calendar features, and price signals from the M5 dataset that improves model performance and can be replicated in SME retail contexts with comparable data structures.
3. Integrate the best-performing forecasting model with the Order-Up-To-Level (OUTL) inventory policy to compute safety stock levels, reorder points, and order-up-to levels at a target

95% service level, and evaluate the resulting inventory parameters across varying lead time scenarios.

4. Test whether the operational gains reported in closely related studies transfer to this setting. Prior work on aggregated retail data (Barghi, 2025; Seyedan et al., 2023) reports improvements of 10–25% in forecast accuracy (MAPE, MAE, RMSE) and inventory cost reductions of 5–15% relative to classical baselines. Because those benchmarks were obtained at coarser levels of aggregation, this study treats them as hypotheses to be tested rather than targets to be met, and evaluates whether they hold at the SKU-daily, intermittent-demand granularity of the M5 series. Where results fall outside these ranges, the gap is itself treated as a finding and examined directly in the Discussion and Conclusion chapters.

5. Develop a Power BI dashboard that translates forecasting outputs and inventory recommendations into an accessible, interactive decision-support tool suitable for deployment in resource-constrained retail environments, including SMEs.

1.2 Justification and Relevance of the Research Topic

1.2.1 Practical Industry Context

The retail sector is one of the most data-intensive and operationally demanding environments for inventory management. Retailers must place replenishment orders days or weeks in advance of actual customer demand, under conditions where that demand is influenced by a complex web of factors: weekly shopping cycles, promotional campaigns, holiday calendars, competitor pricing, weather conditions, and longer-term economic trends. When forecasts are inaccurate, the consequences are immediate and financially significant. Industry estimates suggest that out-of-stock events alone cost global retailers hundreds of billions of dollars annually in lost sales, while excess inventory ties up capital and inflates holding costs across the supply chain (industry sources: Tredence; Omniful).

Small and medium-sized retailers are disproportionately exposed to these risks. Unlike large enterprises such as Walmart or Amazon, SMEs rarely have dedicated data science teams, advanced enterprise resource planning systems, or the computational infrastructure needed to implement sophisticated forecasting tools. They typically rely on manual inventory decisions, spreadsheet-

based planning, or simple moving average heuristics methods that consistently underperform in the face of promotional volatility and seasonal complexity (Wahedi et al., 2023; Keith, 2023). The result is a structural disadvantage: SMEs bear inventory risks comparable to those of large retailers but lack the analytical tools to manage them effectively.

This study addresses that gap directly. By developing and validating a machine learning forecasting framework using the M5-Forecasting (Walmart) dataset as a large-scale empirical proxy, and by designing the methodology and outputs including the OUTL inventory policy integration and Power BI dashboard to be replicable with the kinds of data SMEs already collect (daily sales records, promotional flags, and calendar features), this research provides a practical and transferable solution to a real and quantifiable industry problem.

1.2.2 Academic Gaps and Research Motivation

The academic literature on demand forecasting and inventory optimisation has grown substantially over the past decade, with a clear consensus emerging around the superiority of machine learning and ensemble approaches over classical statistical methods in retail contexts (Barghi, 2025; Nasser et al., 2023; Punia et al., 2020; Kilimci et al., 2019). However, a review of the existing literature reveals four critical gaps that this study is specifically designed to address:

- Gap 1: Disconnect between forecast accuracy and inventory decision-making. Most of the published studies report forecasting accuracy metrics MAPE, MAE, RMSE without translating improved forecasts into operational inventory decisions. Seyedan, Mafakheri, and Wang (2023) note explicitly that most comparative forecasting studies focus on accuracy metrics rather than their implications for inventory control. This study addresses that gap by integrating forecasting outputs with the OUTL policy and measuring, through forward simulation, whether improved forecasts and a distributional safety-stock formulation actually translate into lower inventory cost.
- Gap 2: Limited focus on real-world deployment and operational validation. Most proposed ML forecasting frameworks are validated through historical simulation without addressing live deployment, model maintenance, or the practical constraints of real operations. This study incorporates simulation-based inventory evaluation across multiple lead time scenarios and delivers a Power BI dashboard as a step toward operational deployment.

- Gap 3: Underrepresentation of SME retailers and practical constraints. Existing literature is disproportionately focused on large enterprises with extensive data infrastructure. Wahedi et al. (2023) and Ungureanu (2025) identify the SME gap as a substantive research need. This study frames the Walmart M5 dataset as a methodological benchmark and explicitly designs the resulting framework for transferability to resource-constrained retail environments.
- Gap 4: Insufficient integration of uncertainty in inventory calibration. Classical safety stock formulas assume point-forecast inputs, which systematically underrepresent demand variability. This study incorporates demand standard deviation derived from forecast residuals into the OUTL safety stock calculation, producing inventory parameters that account for forecast uncertainty rather than treating the mean forecast as certain.

1.2.3 SME Transferability: Data, Tools, and Technical Requirements

A recurring claim in this thesis is that the analytical framework developed here is transferable to small and medium-sized retail environments. This section specifies concretely what that claim requires in practice, so that the transferability argument functions as a verifiable design commitment rather than an aspirational statement. A typical SME retailer seeking to replicate this pipeline would need the following.

Data requirements. The pipeline uses three data sources that are standard outputs of any commercial point-of-sale or ERP system: (1) daily unit sales records at the product-store level (equivalent to M5's sales file); (2) a calendar or events file encoding holidays, promotions, and special trading days (equivalent to M5's calendar file); and (3) a weekly price file recording the current selling price per product-store combination (equivalent to M5's sell_prices file). No proprietary data feeds, real-time API integrations, or customer-level transaction data are required. A retailer using any mid-market POS system; such as Lightspeed, Square for Retail, or Shopify POS, can export these three file types as CSV without custom development.

Tools and software. The modelling pipeline is implemented entirely in Python using open-source libraries (pandas, scikit-learn, LightGBM, XGBoost, TensorFlow/Keras) that are freely available and run on standard laptop or desktop hardware. No GPU infrastructure, cloud computing subscription, or enterprise data warehouse is required for the gradient boosting models that constitute the core of the recommended framework. The Power BI dashboard used to present

outputs is available under Microsoft’s free tier (Power BI Desktop), which is sufficient for the visualisations produced in this study.

Technical capacity. Deployment would require one person with intermediate Python competency (equivalent to a data analyst with 12–18 months of scripting experience) to run the preprocessing and modelling pipeline and refresh the dashboard on a weekly or monthly cycle. No dedicated data science team or machine learning engineering capability is assumed. The feature engineering and model training code is designed to be reusable across different product assortments and store configurations with minimal modification. This baseline capacity is consistent with the technical profile of the SME retail operators studied by Wahedi et al. (2023) and Ungureanu (2025), both of whom explicitly designed their frameworks around the same resource constraint.

1.3 Methodology Overview

This research adopts a data science methodology structured around the CRISP-DM (Cross-Industry Standard Process for Data Mining) framework, which provides a systematic, iterative progression from business understanding through data preparation, modelling, evaluation, and deployment. The methodology is organised into five phases:

Phase 1: Data Understanding and Preparation

The primary data source is the M5-Forecasting (Walmart) dataset, sourced from Kaggle. The dataset contains daily unit sales for 3,049 products organised across three product categories (Hobbies, Foods, and Household) and three US states (California, Texas, and Wisconsin) spanning 1,913 consecutive days from 29 January 2011 to 24 April 2016. The dataset is accompanied by a calendar file (encoding day-of-week, SNAP eligibility flags, event labels, and other special days) and a sell-prices file (encoding weekly product prices by store). A stratified random subsample of 502 product-store series, drawn proportionally across category $\times$ state strata, is used to enable rigorous walk-forward cross-validation within the project’s compute budget. Data preparation involves merging and reshaping the three source files, forward-fill and backward-fill imputation of missing prices, removal of closure-day zero-sale records, winsorisation of unexplained outliers, and label encoding of categorical variables. Target-encoded mean features (item_mean, dept_mean, store_mean, state_mean, item_dow_mean, store_dow_mean, cat_month_mean, item_month_mean) are computed exclusively on the training window to prevent data leakage.

Feature engineering produces six categories of derived variables: lag features (7-, 14-, and 28-day lags capturing weekly, bi-weekly, and monthly autocorrelation); rolling statistics (7- and 28-day rolling means plus 7-day rolling standard deviation for volatility); calendar and temporal features (day-of-week, week-of-year, month, year, is_weekend); event proximity features (days_since_last_event, days_to_next_event); promotional and event indicators (is_snap_day, is_sporting_event, is_cultural_event, is_national_event, is_religious_event); and price features (sell_price, price_change).

Phase 2: Model Development and Comparison

- **Classical Statistical Baselines:** Moving Average (MA), Simple Exponential Smoothing (SES), AutoARIMA (with weekly seasonality), Croston’s Classic, the Syntetos–Boylan Approximation (Croston-SBA), and Seasonal Naïve, six methods providing a performance floor. Croston-class methods are included specifically because SKU-daily retail demand is intermittent and zero-inflated, conditions under which ARIMA and SES are known to underperform. All Tier 1 models are fitted per series via the statsforecast library.
- **Standard Machine Learning Models:** Random Forest (RF), XGBoost, and LightGBM gradient boosting and bagging methods trained on the full engineered feature set. XGBoost and LightGBM are configured with the Tweedie loss function (variance power 1.1), suited to non-negative count data with many zero observations, consistent with M5 retrospective best practice.
- **Deep Learning Time-Series Model:** A two-layer stacked LSTM (128 and 64 units, dropout 0.2, 28-day lookback) with log1p target transformation and early stopping capturing long-range temporal dependencies and multi-day promotional dynamics.
- **Ensemble and Hybrid Architecture:** A stacking ensemble in which LightGBM, XGBoost, and Random Forest serve as base learners and a non-negative Ridge regression meta-learner learns optimal combination weights from validation-set predictions. This architecture is preferred over the sequential residual-correction design used in some prior work (Barghi, 2025) because stacking with non-negative Ridge weights is the more conventional heterogeneous ensemble approach and avoids overfitting successive layers to earlier-layer noise (Mitra et al., 2022; Seyedan et al., 2023).
- **LightGBM Quantile Regression (Tier 5):** Five independent LightGBM models trained at quantile levels P10, P50, P90, P95, and P99 providing probabilistic demand forecasts that feed the

empirical safety-stock calculation. This tier does not compete with Tiers 1–4 on point-forecast accuracy; it addresses the Gaussian formula’s assumption of normally distributed residuals, which intermittent retail demand systematically violates.

Phase 3: Model Evaluation

Forecasting accuracy is assessed using five complementary metrics: Mean Absolute Error (MAE), measuring average error in demand units directly interpretable by store managers; Root Mean Squared Error (RMSE), penalising large errors more heavily and sensitive to promotional spikes; Mean Absolute Percentage Error (MAPE), normalising errors as a percentage of actual demand to enable cross-product comparisons; Weighted Mean Absolute Percentage Error (WMAPE), which divides total absolute error by total actual demand and is robust to the zero-actual rows that distort per-row MAPE on intermittent SKU-daily data; and Pred(10%), the proportion of forecasts falling within 10% of actual demand. Improvement figures of 10–25% in MAPE (and MAE/RMSE) reported in prior work on aggregated data are used as a benchmark against which the measured results are tested, rather than as targets the study commits to meeting, given the granularity constraints of SKU-daily intermittent demand.

Phase 4: Inventory Policy Integration

Forecasts are integrated with a periodic (R, s, S) Order-Up-To-Level (OUTL) policy (Silver, Pyke & Thomas, 2017) with review period $R = 7$ days. Three forecasting policies are compared on a like-for-like basis: (1) Classical: Croston-SBA forecasts feeding a Gaussian safety-stock formula $SS_2 = z \cdot \sigma \cdot \sqrt{(R+L)}$, where σ is the standard deviation of forecast residuals; (2) ML-Gaussian: stacking-ensemble forecasts feeding the same Gaussian formula; and (3) ML-Empirical-Quantile: LightGBM quantile forecasts feeding an empirically calibrated safety stock $SS_2 = (q_{0.95} - q_{0.50}) \cdot \sqrt{(R+L)}$, which captures the actual right-skewed demand distribution rather than assuming normality. Reorder point $r = \bar{d} \cdot (R+L) + SS$ and order-up-to level $S = r + Q$ are computed analytically, with Q from the EOQ formula to decouple ordering frequency from forecast level. Total annual cost decomposes into holding cost $HC = (SS + Q/2) \cdot h$, ordering cost $OC = (D/Q) \cdot K$, and stockout cost $SC = (\text{target_SL} - \text{achieved_SL}) \cdot D \cdot \text{unit_cost} \cdot m$ evaluated across three stockout cost multipliers ($m = 0.4, 1.0, 2.0$). Parameters are evaluated across four lead-time scenarios ($L = 7, 10, 14, 21$ days) at a 95% service-level target (the 90% and 99% sweeps of the earlier draft are not reported, as the service-level sweep was an artefact of the assumed-service-level formulation

and is superseded by the simulation of Chapter 5). The achieved service level of each policy is obtained from the forward simulation of Chapter 5 rather than assigned analytically; the 5–15% inventory cost reduction reported in prior work serves as a benchmark against which the simulated outcome is assessed, not as a target.

Phase 5: Dashboard Development and Validation

A Power BI dashboard (free Desktop tier) is developed to translate modelling outputs into an interactive decision-support tool for non-technical retail managers. The dashboard is structured across four pages driven by CSV data feeds exported from the Python pipeline: Forecast Overview (actual sales, ML ensemble forecast, classical baseline forecast, and LightGBM quantile fan P50/P90/P95/P99 with slicers for category, state, and store); Inventory Status (safety stock, reorder point, and order-up-to-level per product-store at the user-selected lead time and service level, with conditional formatting for items approaching reorder); Cost Analysis (total annual cost decomposed into holding, ordering, and stockout components across the three forecasting policies and four lead-time scenarios); and Sensitivity Analysis (cost curves across lead time and the stockout-cost multiplier). The sensitivity analysis in the revised Chapter 5 focuses on lead time (7–21 days) and the stockout cost multiplier (0.4–2.0×); the MAPE and service-level sweeps of the earlier draft are retired, as both were artefacts of the assumed service-level gap rather than measured effects.

1.4 Structure of the Thesis

- Chapter 2 Literature Review: Provides a comprehensive synthesis of existing research on demand forecasting and inventory optimisation in retail, covering classical statistical methods, individual machine learning models, ensemble and hybrid architectures, deep learning approaches, and the integration of forecasting with inventory control policies. Key gaps in the literature are identified and mapped to the research objectives of this study.
- Chapter 3 Methodology: Details the research design, data sources, preprocessing pipeline, feature engineering strategy, modelling approaches, evaluation framework, and inventory policy integration procedure. Justifies the selection of the CRISP-DM framework, the M5-Forecasting dataset, the OUTL inventory policy, and the specific models chosen for comparison.

- Chapter 4 Model Development and Results: Presents the implementation and comparative evaluation of all forecasting models, reports accuracy metrics (MAE, RMSE, MAPE) across models and product categories, and identifies the best-performing approach for integration with the inventory framework.
- Chapter 5 Inventory Optimisation and Dashboard: Demonstrates how forecasting outputs are integrated with the OUTL policy to produce safety stock levels, reorder points, and total cost estimates. Presents the Power BI dashboard and evaluates operational impact through simulation and sensitivity analysis.
- Chapter 6 Discussion: Interprets results in the context of existing literature, evaluates the extent to which research objectives have been achieved, discusses the practical implications for SME retailers, and acknowledges study limitations.
- Chapter 7 Conclusion and Future Work: Summarises key findings and contributions, and outlines directions for future research including real-time streaming data integration, probabilistic demand forecasting, multi-echelon supply chain extension, and live deployment validation with retail partners.

1.5 Theoretical and Conceptual Foundations

1.5.1 Classical Inventory Theory

Classical inventory theory, developed from the foundational Economic Order Quantity (EOQ) model of Harris (1913) and extended through periodic review systems (Silver, Pyke & Thomas, 2017), provides the operational backbone for inventory decision-making in this study. The core concepts of safety stock, reorder point, service level, and order-up-to level formalised in the OUTL (R, s, S) policy translate demand forecasts into actionable replenishment decisions. The Gaussian safety stock formula $SS = z \cdot \sigma \cdot \sqrt{R+L}$ bridges the statistical properties of the demand forecast (σ , the standard deviation of forecast errors) with the operational constraint (L, the supply lead time) and the business requirement (z, the service-level z-score). However, this formula assumes forecast residuals are normally distributed an assumption that intermittent retail demand, which is right-skewed and zero-inflated, systematically violates. A central theoretical motivation of this study is therefore the extension of classical inventory theory through an empirical quantile safety stock

formulation, $SS_{2h5} = (q_{0.95} - q_{0.50}) \cdot \sqrt{(R+L)}$, derived from LightGBM quantile-regression forecasts, which captures the true shape of the demand distribution rather than imposing parametric normality. The comparison of Gaussian and empirical-quantile formulations across realistic operational scenarios is the primary analytical contribution of Chapter 5.

1.5.2 Machine Learning and Ensemble Learning Theory

Ensemble learning theory; the principle that combining predictions from multiple models reduces bias and variance relative to any single constituent model underpins the choice of gradient boosting ensembles as the primary modelling architecture. The theoretical result by Zhou et al. (2002) that combining diverse partial predictors yields superior generalisation motivates the stacking ensemble design of Chapter 4, in which LightGBM, XGBoost, and Random Forest base learners are combined through a non-negative Ridge regression meta-learner that learns optimal combination weights from validation-set predictions. This stacking architecture is preferred over the sequential residual-correction design used in some prior work (Barghi, 2025) because it avoids overfitting successive layers to the noise of earlier layers and is the standard heterogeneous ensemble approach in the modern forecasting literature (Mitra et al., 2022; Seyedan et al., 2023). The specific models selected LightGBM for computational efficiency and histogram-based binning on large datasets, XGBoost for L1/L2 regularisation and generalisation, Random Forest for stable variance reduction, and LSTM for temporal dependency modelling are each grounded in established theoretical properties validated in the retail forecasting literature (Barghi, 2025; Nasser et al., 2023; Punia et al., 2020). Quantile regression theory provides the probabilistic foundation for Tier 5: by minimising the asymmetric pinball loss at each target quantile α , LightGBM quantile models produce calibrated distributional forecasts that, unlike point forecasts, directly encode the upper tail of demand required for safety-stock calculations. This probabilistic framework is the theoretical bridge between the forecasting and inventory layers of the study.

1.5.3 Supply Chain and Demand Uncertainty Theory

Demand uncertainty theory, which recognises that real-world retail demand cannot be forecast with certainty and that inventory policies must be designed to perform robustly across the distribution of possible demand outcomes, motivates several key design choices in this study. The finding that the Gaussian safety-stock formula systematically under-allocates buffer inventory for

intermittent demand covering approximately the 83rd rather than the targeted 95th percentile of empirical demand is grounded in this theoretical framework: the formula's parametric assumption of normality fails precisely because the true demand distribution is right-skewed and heavy-tailed relative to a normal distribution with the same mean and variance. Whether this distributional under-allocation translates into realised stockouts and higher cost, or remains a property of the distribution with no operational effect is measured directly in Chapter 5 rather than assumed. The sensitivity analysis in Chapter 5, which examines how inventory cost outcomes vary across lead-time scenarios (7, 10, 14, 21 days) and stockout cost assumptions (multiplier 0.4, 1.0, 2.0×), is grounded in the recognition that supply chain resilience requires robustness to parameter uncertainty rather than optimisation for a single point estimate (Seyedan et al., 2023; Silver, Pyke & Thomas, 2017).

1.5.4 CRISP-DM as a Research Methodology Framework

The Cross-Industry Standard Process for Data Mining (CRISP-DM) provides the overarching methodological structure that connects the three theoretical domains above into a coherent research workflow. By progressing iteratively from business understanding through data preparation, modelling, evaluation, and toward deployment, CRISP-DM ensures that each technical decision in this study from feature engineering choices to model selection to dashboard design is anchored in the operational business problem rather than pursued as a purely academic exercise. This framework has been adopted in several closely related retail forecasting studies (Ungureanu, 2025; Bastos, 2023) and its iterative structure is particularly well-suited to the cycles of model refinement and inventory policy calibration that characterise this research.

1.6 Contributions of the Study

1.6.1 Methodological Contribution

The study makes two methodological contributions. First, it provides the first quantitative measurement on the M5 benchmark of the Gaussian safety-stock under-allocation gap for intermittent retail demand: the empirical 95th-percentile demand spread, derived from LightGBM quantile regression, is 1.36 times the Gaussian-implied buffer at the same nominal service level, corresponding to an effective (analytical) coverage of about 83% of the daily demand distribution under a 95% nominal target, a property of the demand skew, not the service level realised in

operation, which Chapter 5 measures directly by simulation and finds to be 100% for all policies in this regime. This measurement, and the empirical quantile safety-stock formulation built upon it, represents a principled extension of the textbook (R, s, S) policy to non-Gaussian demand distributions; whether that extension yields an operational cost benefit is shown in Chapter 5 to depend on the demand regime. Second, the study develops and validates a five-tier comparative forecasting pipeline spanning classical statistical baselines, gradient boosting models with Tweedie loss, a stacked LSTM, a Ridge-regression stacking ensemble, and a five-quantile LightGBM model integrated end-to-end with the OUTL inventory policy and evaluated on a stratified subsample of 502 product-store series from the M5-Forecasting (Walmart) dataset. The pipeline is implemented exclusively in free and open-source tooling, runs on a standard 16 GB-RAM laptop in under three hours, and is documented to the level required for independent replication providing a reproducible methodological template for future retail forecasting and inventory research.

1.6.2 Practical Contribution

The study delivers a complete, end-to-end analytical workflow from raw M5 data ingestion through preprocessing, feature engineering, model training, OUTL inventory calculation, and Power BI dashboard export that is explicitly designed for deployment in SME retail environments. The Power BI dashboard is structured across four operational pages: Forecast Overview (actual sales, ML ensemble and classical forecasts, and LightGBM quantile fan); Inventory Status (safety stock, reorder point, and order-up-to-level per product-store with conditional formatting for items approaching reorder); Cost Analysis (total annual cost decomposed by policy and lead-time scenario); and Sensitivity Analysis (cost curves across lead time and the stockout-cost multiplier). The SME accessibility commitment is implemented through three concrete constraints: all tooling is free and open-source (Python, scikit-learn, LightGBM, XGBoost, TensorFlow, Power BI Desktop free tier); the full pipeline runs in under three hours on a standard 16 GB-RAM laptop with no GPU; and the only data inputs required daily SKU-store sales, weekly prices, and a promotional/event calendar are standard outputs of off-the-shelf point-of-sale systems. These specifications collectively ensure the framework is accessible to SMEs without enterprise-scale infrastructure, licensing, or data science expertise.

1.6.3 Research Contribution

The study provides three empirical findings that advance the research literature. First, it shows that the operational value of ML forecasting in retail inventory management is conditional on the demand regime rather than automatic. When inventory outcomes are measured by forward simulation rather than assumed, the ML-Gaussian policy is indistinguishable from the classical baseline and the empirical-quantile policy does not lower cost on the M5 test window, because the standard safety-stock buffers already meet demand in full, a 100% realised service level, with no stockouts to avert. Second, and relatedly, the study quantifies a granularity mismatch in the published benchmarks: the 10–25% accuracy and 5–15% cost-reduction figures reported by Barghi (2025) and Seyedan et al. (2023) were obtained on aggregated retail and supply-chain data, where demand is comparatively smooth; at SKU-store-day granularity the attainable accuracy gain is far smaller (approximately 4.5% in MAE for the best ensemble over the strongest classical baseline) and the inventory benefit is regime-dependent. The shortfall relative to those benchmarks is therefore not a deficiency of the pipeline but evidence that the benchmarks are conditional on aggregation level. Third, the study delimits the boundary condition under which the quantile-based safety-stock reformulation pays off: it reduces cost only when the baseline policy is exposed to genuine stockout risk, and adds idle holding cost when service is already saturated, a result that gives practitioners an explicit test for whether the more complex policy is worth adopting in their own setting. Together, these findings reframe the contribution from a claimed cost saving to a rigorous, simulation-based characterisation of when ML forecasting and distributional safety stock affect inventory outcomes, and when they do not.

In summary, this thesis is positioned at the intersection of data science, operations management, and retail analytics. It moves deliberately beyond the production of accurate forecasts as an end in themselves, toward a fully integrated decision-driven framework that connects predictive modelling to operational inventory policy, delivers outputs through an accessible dashboard, and is explicitly designed for applicability in the SME retail environments where the need for improved demand forecasting is greatest and the current gap between need and capability is widest.

CHAPTER 2

LITERATURE REVIEW

2.1 Introduction

Retail inventory management sits at the intersection of two persistent operational problems: the difficulty of predicting what customers will buy and when, and the cost of holding the wrong quantity of stock in anticipation of that demand. When forecasts are inaccurate, retailers face a binary penalty stockouts that erode revenue and customer loyalty, or overstock that ties up capital, increases holding costs, and, in the case of perishables, generates waste with financial and environmental consequences. The scale of these costs across the global retail sector has motivated decades of academic research into demand forecasting methodologies.

The emergence of machine learning as a mainstream analytical tool has fundamentally changed the terms of this research programme. Where earlier studies debated the relative merits of ARIMA variants and exponential smoothing families, contemporary research compares gradient boosting ensembles, long short-term memory (LSTM) networks, transformer-based models, and hybrid architectures that combine statistical and machine learning components. The volume of published work has accelerated sharply: of the 52 sources reviewed here, the majority were published after 2020, reflecting both the maturation of machine learning tools and the growing recognition that retail data; rich in promotional dynamics, seasonal patterns, and external shocks provides an ideal testing ground for advanced predictive models.

This review synthesises these 52 sources to provide a structured, critical account of the state of knowledge. Sources include two master's theses (Barghi, 2025; Ungureanu, 2025; Stanelyte, 2021; Bastos, 2023), one doctoral dissertation (Seyedan, 2023), peer-reviewed journal articles published in outlets including Applied Sciences, International Journal of Production Research, Operations Research Forum, Logistics, and various engineering and computer science journals, and conference papers from IEEE and other professional bodies. The synthesis is thematic rather than chronological, organised around the questions the literature collectively addresses.

2.2 Problem Framing: What the Literature Collectively Addresses

Across 52 sources, four interconnected research problems recur with sufficient regularity to constitute the defining concerns of this literature.

2.2.1 The Inadequacy of Traditional Forecasting Methods

The most consistent framing across the literature is the contrast between the limitations of classical statistical methods and the capabilities of machine learning. Traditional methods; Moving Averages, Exponential Smoothing, ARIMA, SARIMA, and their variants are characterised as adequate for stable, univariate, linear demand environments but fundamentally ill-equipped for the multivariate, non-linear, promotion-driven, and seasonally complex demand patterns typical of modern retail.

Barghi (2025) documents this precisely on the Rossmann dataset: ARIMA produces a Mean Absolute Error (MAE) of 3,167.62 and Root Mean Squared Error (RMSE) of 4,210.67, compared to 416.77 and 608.66 for the proposed hybrid gradient boosting ensemble, a reduction of 86.8% in MAE. Nasser et al. (2023, applsci1311112) similarly find that tree-based ensemble methods substantially outperform classical baselines across three perishable product categories using over six years of daily demand data covering more than 330 products and 5.2 million records. Razzak et al. (2025) arrive at the same conclusion through a systematic review of 72 peer-reviewed articles: traditional statistical models demonstrate continued relevance only in scenarios marked by stable demand patterns, short forecasting horizons, and limited data availability.

Kilimci et al. (2019) were among the earlier studies to formally establish this comparison, developing an intelligent demand forecasting system applied to SOK Market in Turkey (6,700 stores, 1,500 products, 23 distribution centres) that blends time series analysis, support vector regression (SVR), and deep learning through a novel decision integration strategy. Their extensive experiments demonstrate that the proposed system delivers noteworthy results compared to state-of-the-art statistical methods alone. Haque, Amin, and Miah (2023, 2308_11939v1) extend this critique by identifying a specific gap: most existing studies focus on historical sales data while ignoring the influence of macroeconomic conditions on consumer behaviour, incorporating CPI, Index of Consumer Sentiment, and unemployment rates into multivariate retail demand models.

2.2.2 The Challenge of Demand Volatility and Uncertainty

A second problem concerns not just average forecast accuracy but the capacity to model and respond to demand volatility, the unpredictable spikes and troughs driven by promotions, holidays, weather, macroeconomic shocks, and competitive actions. The literature frames this as distinct from the baseline accuracy problem: even a model with low average error may fail catastrophically during high-stakes periods when demand deviates sharply from its trend.

Nasseri et al. (2023) explicitly design their study around three perishable product categories, fresh meat, fruit, and another category, finding that the performance disparity between extra tree regressors and LSTM is particularly pronounced for fresh meat (the highest volatility category) and marginal for fruit (more stable). Kalisetty (2023) proposes a scalable AI-driven approach using decomposable local-global network design and an ensemble of traditional time series and CNN-based predictors specifically to handle complex non-linear temporal patterns with multi-scale characteristics in large retail datasets. Barghi (2025) addresses uncertainty at the data level through fuzzy generalisation, a preprocessing technique that discretises features into 12 representative midpoint categories to smooth the input space before training the boosting ensemble, reducing MAPE from 6.48% in Case Study 1 to 3.53% in Case Study 2, a 45.5% improvement attributable entirely to this preprocessing step.

Kalusivalingam et al. (2020) propose a different approach to volatility, integrating deep reinforcement learning (DRL) with neural networks to create systems that learn optimal inventory policies through trial and error and adjust dynamically to fluctuating market conditions. Mittal (2025) introduces the most statistically sophisticated treatment of uncertainty, presenting an end-to-end framework that uses Temporal Fusion Transformers (TFT), LSTM networks, and spatiotemporal Graph Neural Networks (ST-GNN), combined with a quantile-based uncertainty-to-inventory translation mechanism that converts predictive distributions into safety stock and reorder parameters, achieving 22–35% improvements in forecast accuracy and 14–21% reductions in total inventory cost versus statistical baselines.

2.2.3 Integration of Forecasting with Inventory Policy

A third problem concerns the disconnect between forecasting research and operational inventory decision-making. Much of the earlier literature, as Seyedan et al. (2023) observe, compared forecasting accuracy metrics without translating improved forecasts into inventory decisions. The more operationally grounded strand of the literature explicitly bridges this gap by coupling forecasting models with inventory policies, most commonly the Order-Up-To-Level (OUTL) framework and measuring impact in terms of total cost, safety stock levels, reorder points, and service levels.

Barghi (2025) and Seyedan, Mafakheri, and Wang (2023) both implement the (R, s, S) OUTL policy, using ensemble demand predictions to compute safety stock $SS = z \cdot \sigma \cdot \sqrt{R+L}$, reorder point $r = d(R+L) + z \cdot \sigma \cdot \sqrt{R+L}$, and order-up-to level $S = r + Q$, targeting a 95% service level. Wahedi et al. (2023, applsci1308581) investigate this integration in a small and medium-sized enterprise (SME) context, partnering with Svanehøj Danmark A/S to evaluate whether ML forecasts translate into improvements in inventory Key Performance Indicators (KPIs), finding that predictive ML models perform superior to statistical forecasting approaches but the advantage is not always consistent when industrial KPIs rather than pure accuracy metrics are the criterion. Sihotang (2023) takes the simpler approach of exponential smoothing combined with a Decision Support System (DSS) for inventory ordering decisions in retail, demonstrating that even classical methods deliver tangible operational improvements when properly integrated with replenishment logic.

2.2.4 Scalability, Accessibility, and Context-Specificity

A fourth research problem concerns who can benefit from advanced forecasting. Much of the literature focuses on large enterprises; Walmart, Amazon, Rossmann, large grocery chains while a smaller but growing strand addresses smaller operators. Ungureanu (2025) explicitly targets mid-sized grocery retailers, observing that existing literature largely focuses on either large grocery chains deploying advanced systems or small boutique stores, leaving mid-sized operators underserved, and designing the entire framework around the constraint that the solution must be deployable without enterprise-scale budgets or IT infrastructure. Wahedi et al. (2023) similarly focus on SMEs, citing limited resource capabilities and financial constraints as amplifying the need for accessible data-driven decision support systems. Obi (2024, 11010) applies a stacking

ensemble framework specifically to retail demand forecasting using Walmart outlet data, demonstrating that ensemble approaches can deliver value in operationally realistic settings without requiring proprietary infrastructure.

2.3 Theoretical Frameworks and Conceptual Foundations

2.3.1 Classical Inventory Theory

The inventory optimisation strand of the literature is grounded in classical operations research. The periodic versus continuous review distinction, the Economic Order Quantity (EOQ) model, the Newsvendor framework, and the Order-Up-To-Level policy all appear across multiple sources as foundational building blocks. Barghi (2025), Seyedan et al. (2023), and Ungureanu (2025) all adopt periodic review systems, the (R, s, S) system in the former two and a daily P-system in the latter, on the grounds that periodic review is compatible with the batch data updates characterising retail sales systems.

The safety stock formula $SS = z \cdot \sigma \cdot \sqrt{R+L}$ and its variants appear across Barghi (2025), Seyedan et al. (2023), and Seyedan (2023, Chapter 5), where z is the z-score for the desired service level and σ is the standard deviation of demand. The total cost model $TC = PC + HC + OC$, decomposing into purchase, holding, and ordering costs, is similarly shared. Namwad et al. (2024, Optimizing_Inventory_Management_with_Seasonal_Dema) extend the classical inventory model into a fuzzy environment, treating degradation rates and faulty percentages as fuzzy variables within a finite planning horizon that also incorporates carbon pollution costs, producing an inventory model for deteriorating products that accounts for imperfect quality and partial backlogging.

2.3.2 Ensemble Learning Theory

Ensemble learning theory; the principle that combining predictions from multiple models reduces variance and corrects individual model biases provides the second major theoretical pillar. Seyedan, Mafakheri, and Wang (2022) ground this in the result by Zhou et al. (2002) that combining partial basic predictors yields comparable or superior generalisation to combining all predictors simultaneously. Mitra et al. (2022, 43069_2022_Article_166) motivated their hybrid

RF-XGBoost-LR model by observing that no single model achieves complete forecasting accuracy and that hybrid frameworks strategically combining the merits of individual models reduce individual weaknesses. Islam et al. (2024, Paper42024) similarly find that a hybrid RF-XGBoost-LR model outperforms all individual component models in terms of MAE, MSE, and R^2 on retail sales data, confirming the ensemble advantage across geographically diverse datasets.

The literature distinguishes between homogeneous ensembles (combining models of the same type, as in Random Forest) and heterogeneous ensembles (combining architecturally distinct models). Seyedan et al. (2023) explicitly critique the dominance of homogeneous ensemble approaches, arguing that heterogeneous DNNs; MLP for global patterns, LSTM for temporal dependencies, 1D-CNN for local patterns capture a broader feature space and arguing this is why their heterogeneous ensemble outperforms the best individual model by 22.3% in MAPE for sports products.

2.3.3 Machine Learning Methodology Frameworks

Several sources adopt formal research methodology frameworks. Barghi (2025) and Ungureanu (2025) both use Design Science Research (DSR) or CRISP-DM (Cross Industry Standard Process for Data Mining) to structure the iterative development, evaluation, and refinement of the proposed artefact. Stanelyte (2021, 80509157) adopts a business management perspective in a Lithuanian retail chain context, framing the research as strategic investment analysis. Bastos (2023, content) applies CRISP-DM in a sportswear company context at NOVA Information Management School, using Random Forest, XGBoost, and LightGBM for demand forecasting across multiple SKUs.

2.3.4 Supply Chain Theory

Supply chain theory; the study of information flows, material flows, and coordination mechanisms across multi-stage networks informs the framing of several papers. The bullwhip effect, where demand variability amplifies upstream in a supply chain, motivates the accuracy improvement agenda across Nasution, Matondang, and Ishak (2022, jsti8637) and Mohammed and Mandal (2022, MohammedIJSCM). Pasupuleti et al. (2024, logistics0800073v2) situate machine learning within a broader supply chain agility and sustainability framework, arguing that ML-driven logistics and inventory management delivers not only efficiency gains but environmental benefits through reduced emergency restocks and more efficient transportation. Farooq, Abbey, and Onukwulu (2024, 20241227183207) develop a conceptual framework that explicitly links

inventory optimization with sustainability goals, circular inventory systems, and resource-efficient procurement practices.

2.4 Data Sources, Characteristics, and Preprocessing

2.4.1 Benchmark Retail Datasets

The Rossmann Store Sales dataset, sourced from Kaggle, is the most widely referenced benchmark across the reviewed literature. Barghi (2025) uses all 1,115 stores over 942 days (January 2013 to July 2015), producing 1,017,209 records with 19 variables including store attributes, daily sales, customer counts, temporal indicators, and promotional flags. Seyedan, Mafakheri, and Wang (2022, 2023) reference the Rossmann dataset as a key benchmark in their literature review and motivation, while Bastos (2023) works with a proprietary sportswear company dataset structured similarly to Rossmann in terms of its multivariate SKU-level character.

Walmart sales data constitutes the second major benchmark cluster. Obi (2024) uses historical sales from ten Walmart outlets across Texas, California, and Wisconsin. Haque et al. (2023) use publicly available Walmart demand data from Kaggle, enriched with macroeconomic variables (CPI, ICS, unemployment rates). The Favorita Grocery Sales dataset from Kaggle, used by Ungureanu (2025), provides store-level daily sales for a large Ecuadorian grocery chain, with promotional flags, oil price indicators, and holiday calendars, approximately 5 years of data across multiple stores and product families.

The UCI Online Retail Dataset (~541,000 transaction records from a UK e-commerce retailer) is used by Ungureanu (2025) alongside the Supermarket Sales Data (daily sales 2020–2023, focused on perishables). Sivilai (2025, 5Application) and Bendhi (2023, IJAIDSMLV4I2P108) use retail time-series datasets structured around SKU-level historical sales. Loa, Wiratama, and Halim (2024, 1571158092stampede) use a real-world dataset of 78,000 transaction records from January 2024 to January 2025 covering 10 attributes including product ID, quantity, price, and store ID. The 'Store Item Demand Forecasting Challenge' Kaggle dataset (50 items, 10 stores, January 2013 to December 2017) is used in Seyedan (2023, Chapter 5) for RMDN evaluation.

2.4.2 Industry and Proprietary Datasets

Several studies use proprietary or industry-specific data unavailable for public download, providing ecological validity at the cost of reproducibility. Nasser et al. (2023) use over six years of historical demand data from a prominent Austrian retail entity, covering more than 330 products and 5.2 million records, enriched with meteorological and COVID-19 data; one of the largest and most richly featured datasets in the review. Wahedi et al. (2023) partner with Svanehøj Danmark A/S, a Danish SME manufacturer, using internal ERP data. Khan et al. (2020, Effective_Demand_Forecasting) apply their business intelligence empowered ML system to real-time organisational data, achieving up to 92.38% accuracy in intelligent demand forecasting for a specific store.

Kilimci et al. (2019) use real-world data from SOK Market in Turkey, one of the largest Turkish grocery chains. Zohdi et al. (2022) use customer information data from an unspecified company, focusing on the role of customer behavioural data (demographics, purchase history, preferences) in improving demand forecast accuracy. Stanelyte (2021) works within a Lithuanian retail network with proprietary sales data, building a demand prediction model for a specific retail chain's inventory optimisation programme.

2.4.3 Data Characteristics and Preprocessing

Across the literature, common data quality challenges include missing values, negative entries (returns), extreme outliers from promotional spikes, zero-sale records from store closures, and categorical variables requiring encoding. Preprocessing approaches are largely consistent: mean or forward-fill imputation for missing values; outlier detection using percentile-based methods or unsupervised SVM followed by capping or interpolation; removal of zero-sales records; and encoding of categorical variables using label encoding or one-hot encoding. Feature engineering; the construction of lag features, rolling averages, holiday proximity indicators, and promotional interaction terms, is identified consistently as a major contributor to model performance.

Nasser et al. (2023) explicitly demonstrate the value of external feature integration, incorporating meteorological and COVID-19 data and showing improved accuracy compared to models using only historical sales. Haque et al. (2023) similarly demonstrate that enriching time series with CPI, consumer sentiment indices, and unemployment rates improves forecast accuracy for retail demand, bridging the gap between micro-level sales data and macro-level economic conditions.

Taparia et al. (2023, 02_Kumar) use exponential smoothing predictions as one input feature to machine learning regression models alongside trend and seasonality components; an approach that leverages the best of both classical and ML paradigms.

Table 1: Summary of Key Datasets Used Across the Literature

Dataset / Source	Papers Using It	Size / Scope	Key Features
Rossmann Store Sales (Kaggle)	Barghi (2025); Seyedan et al. (2022, 2023)	1,115 stores; 942 days; ~1M records	Sales, promotions, store type, holidays, temporal
Walmart Sales (Kaggle/Public)	Obi (2024); Haque et al. (2023)	10 stores; multiple years	Sales, macroeconomic indicators, CPI, ICS
Favorita Grocery (Kaggle)	Ungureanu (2025)	Multi-store; 2013–2017; ~5 years	Sales, promotions, oil prices, holiday calendars
SOK Market (Turkey)	Kilimci et al. (2019)	6,700 stores; 1,500 products	Sales, demand history, store features
Svanehøj ERP Data (Proprietary)	Wahedi et al. (2023)	SME manufacturer; internal ERP	Spare parts; sales; inventory levels
UCI Online Retail	Ungureanu (2025)	~541K transactions; UK e-commerce	Transaction-level sales, product, customer
Constante Supply Chain Data	Seyedan et al. (2022, 2023)	Sports & electronics; Jan 2015–Sep 2017	52 features; sales, price, holidays, demand
Store Item Demand (Kaggle)	Seyedan (2023, Ch. 5)	50 items; 10 stores; 2013–2017	Date, store ID, item ID, daily sales
Sportswear Company (Proprietary)	Bastos (2023)	Multiple SKUs; internal company data	Sales, SKU attributes, temporal features
Retail Transaction Data	Loa et al. (2024)	78,000 records; Jan 2024–Jan 2025	Product ID, quantity, price, store ID

2.5 Forecasting Methodologies: A Layered Overview

2.5.1 Classical Statistical Methods (Baseline Tier)

Classical methods appear primarily as benchmarks against which machine learning approaches are evaluated, not as methodological contributions in themselves. Moving Average (MA), Exponential Smoothing (ES and its Holt-Winters variant), ARIMA, SARIMA, and SARIMAX are the most commonly used baselines. Barghi (2025) compares against ARIMA (MAPE: 9.79%), MA (MAPE: 9.72%), and SES (MAPE: 13.79%). Ungureanu (2025) uses SARIMAX as the primary baseline (MAPE: approximately 13.5% at the aggregate level for the comparative model table), noting it incorporates exogenous variables (promotions, holidays, day-of-week) and provides interpretable parameter estimates through ACF/PACF-guided configuration.

Sihotang (2023) stands out as one of the few studies that treats classical methods not as a failed baseline but as a practical, deployable solution for resource-constrained retailers, arguing that exponential smoothing combined with a DSS provides actionable inventory ordering improvements even without machine learning. Keith (2023, *OptimizingInventoryManagement*) similarly argues that advanced forecasting techniques including both statistical and ML approaches improve inventory management performance across different global contexts (US, Canada, Europe, Africa), but notes that implementation quality matters as much as algorithm choice.

2.5.2 Standard Machine Learning Algorithms

2.5.2.1 *Tree-Based Methods*

Random Forest (RF), XGBoost, LightGBM, Gradient Boosting, and AdaBoost constitute the most widely deployed machine learning tier across the reviewed literature. Their consistent advantages in handling of non-linear relationships, robustness to outliers and missing data, interpretability through feature importance, and compatibility with mixed feature types makes them the default first choice when moving beyond statistical methods.

Nasseri et al. (2023) conduct the most rigorous comparative study of tree-based methods, evaluating Extra Tree Regressors (ETR), XGBoost, Random Forest Regression (RFR), and Gradient Boosting Regression (GBR) alongside LSTM across 5.2 million records, finding that ETR outperforms LSTM in MAPE, MAE, RMSE, and R^2 particularly for fresh meat, with comparable performance across tree-based methods for less volatile categories. Mitra et al. (2022)

evaluate RF, XGBoost, gradient boosting, AdaBoost, and ANN on US retail weekly sales data including temperature and store size features, finding XGBoost and RF to be the strongest individual models before introducing their hybrid RF-XGBoost-LR framework. Ismail et al. (2025, 973) propose a hybrid KNN-SVM model for retail market inventory forecasting, achieving accuracy of 0.9999 and F1-score of 0.9808 on real-world sales data, demonstrating that even computationally lighter hybrid models can deliver near-perfect classification accuracy on well-structured retail data.

Loa, Wiratama, and Halim (2024) apply XGBoost directly to retail inventory management, achieving a 15% reduction in overstock and stockouts compared to traditional methods on a dataset of 78,000 records. Taparia et al. (2023, 02_Kumar) demonstrate that the choice of the best ML regressor varies by SKU and that a hybrid model trained on top of initial exponential smoothing predictions improves accuracy further, introducing the concept of SKU-level model selection based on a feasibility index. Taparia, Mishra, Gupta, and Chandiramani (2023, 03_Hitesh) extend this to price optimisation, using linear regression for demand incorporating price as an input and then optimising prices considering inventory and perishability costs.

2.5.2.2 Support Vector Methods

Support Vector Regression (SVR) and Support Vector Machine (SVM) classifiers appear as intermediate-tier methods more powerful than moving averages but generally outperformed by tree-based ensembles and deep learning on large datasets. Kilimci et al. (2019) include SVR as one component of their decision integration strategy alongside time series analysis and deep learning, finding that the combination outperforms any single method. Wahedi et al. (2023) evaluate SVR among other methods for the Svanebjerg SME context, finding it performs reasonably but ranks below ANN and LSTM on their specific KPI measures.

2.5.2.3 Neural Networks and Basic Deep Learning

Multi-layer Perceptron (MLP), basic Artificial Neural Networks (ANN), and K-Nearest Neighbours (KNN) regression constitute an intermediate tier positioned above classical methods and below the more sophisticated architectures. Zohdi et al. (2022) find that ANN-based methods outperform KNN, decision trees, gradient boosting, and extreme learning machines on customer information-based demand forecasting, with the extreme learning machine performing as well as or better than other predictors while requiring less computational time. Obi (2024) compares a

Stacking Ensemble against an MLP, finding the stacking approach outperforms the deep learning baseline on Walmart retail data.

2.5.3 Deep Learning Architectures

2.5.3.1 LSTM and Recurrent Neural Networks

Long Short-Term Memory (LSTM) networks are the most widely used deep learning architecture across the review, valued for their capacity to capture temporal dependencies and long-range patterns that simpler models miss. Punia, Nikolopoulos, Singh, Madaan, and Litsiou (2020, 2020_Deep_learning) provide an early and influential demonstration of LSTM applied to multi-channel retail demand forecasting, published in the International Journal of Production Research, finding LSTM and Random Forest complementary. LSTM excels at pattern learning while RF handles structured features and proposing their combination. Abbasimehr et al. (referenced across multiple sources including Seyedan, 2023, and Barghi, 2025) use LSTM with grid search hyperparameter optimization for furniture company monthly sales, achieving superior performance over ARIMA, SES, ANN, and SVM.

Kalisetty (2023) proposes three scalable LSTM-based deep learning architectures specifically designed for SKU-level retail demand at scale: SoxNeT (integrating mobility data into extreme learning machines with LSTM cells), DemandNet (applying CNN and RNN to model demand with urban land-use context effects), and EKO (a deep learning ensemble combining boosting models with categorical embedding layers). Marripudugala (2024, IRJMETs60700053312) combines CNN and LSTM specifically for inventory demand prediction, integrating real-time sales data, seasonal trends, and external market and promotional factors, validated across multiple retail environments.

2.5.3.2 Advanced Deep Learning: Transformers, GRUs, and Probabilistic Networks

The most recent strand of the literature introduces architecturally sophisticated deep learning models. Mittal (2025, ssrn5386001) evaluates Temporal Fusion Transformers (TFT), LSTM, and spatiotemporal Graph Neural Networks (ST-GNN) as probabilistic forecasting models for multi-channel retail, using a quantile-based mechanism to translate predictive distributions into safety stock and reorder parameters. Chetlapalli and Vinayagam (2019, Paper10459466) propose a BERT-based demand forecasting model integrated with the Squirrel Search Algorithm (SSA) for

hyperparameter optimisation in e-commerce, one of the earliest attempts to apply transformer architectures to retail demand. Maharana, Choudhury, and Vimaladevi (2025, ssrn5832182) integrate machine learning, deep learning, and computer vision within a cloud-native infrastructure for fashion retail inventory optimisation, where rapid trend shifts driven by social media require near-real-time predictive capabilities.

Seyedan (2023, Chapter 5) introduces Recurrent Mixture Density Networks (RMDNs) which combine LSTM's sequential memory with Mixture Density Networks' probabilistic output for distribution-level demand forecasting. Rather than predicting mean demand, RMDNs parameterise a Gaussian mixture distribution (weights π_i , means μ_i , standard deviations σ_i for K components), enabling calculation of weighted mean μ_{LTD} and weighted variance σ^2_{LTD} for safety stock and reorder point computation. The RMDN approach yields safety stock values of 20 units (90% service level, 7-day lead time) compared to 85 units under LSTM, reflecting a more confident and precise characterisation of demand that avoids the conservative over-buffering inherent in point-forecast-based inventory models.

2.5.4 Ensemble and Hybrid Architectures

2.5.4.1 Gradient Boosting Sequential Ensembles

Barghi (2025) develops the most architecturally distinctive gradient boosting ensemble: a sequential pipeline where LightGBM generates initial predictions (exploiting computational efficiency through GOSS and leaf-wise tree growth), CatBoost corrects residuals (natively handling categorical features through ordered boosting and target statistics), and XGBoost performs final regularisation-based refinement (L1 and L2 regularisation preventing overfitting). The final prediction $\hat{y} = f_{LGBM}(x) + f_{CatBoost}(x) + f_{XGB}(x)$ accumulates sequential corrections. The learning rate $\eta = 0.05$ is consistent across all three models to ensure balanced contribution. The Generic Hybrid Ensemble achieves MAE of 416.77 and MAPE of 6.48%; with fuzzy generalisation preprocessing, the Fuzzy Hybrid Ensemble achieves MAE of 209.21 and MAPE of 3.53%, with 97.11% of predictions within 10% of actual sales.

Mitra et al. (2022) and Islam et al. (2024) independently develop a structurally similar RF-XGBoost-LR hybrid, combining random forest's bagging with XGBoost's boosting and using linear regression as a meta-learner to combine their outputs, finding this hybrid consistently

outperforms all individual components across MAE, MSE, and R^2 metrics on US retail weekly sales data. Loa et al. (2024) apply XGBoost as a single model but demonstrate that gradient boosting alone, without ensembling, reduces overstock and stockouts by 15% providing an important lower bound on the standalone contribution of gradient boosting methods.

2.5.4.2 Heterogeneous Deep Learning Ensembles with Meta-Learning

Seyedan, Mafakheri, and Wang (2023) build a stacking ensemble of three architecturally distinct deep learning models; MLP (global patterns), LSTM (temporal dependencies), 1D-CNN (local patterns) combined through a meta-learner MLP trained on predictions from the base models. Data is split into grid-search training set (60%), meta-model training set (30%), and holdout test set (10%). The ensemble achieves MAPE of 5.22% for sports products and 9.58% for electronics, outperforming the best individual base model (LSTM) by 22.3% and 24.4% respectively. Seyedan, Mafakheri, and Wang (2022) extend this with customer segmentation using RFM analysis and five clustering algorithms (K-means, single linkage, complete linkage, centroid linkage, Ward's linkage), combined through majority voting before applying LSTM and Prophet forecasters per segment, with Bayesian Model Averaging (BMA) as the ensemble combination method achieving MAPE of 8.12% with the full cluster-ensemble approach versus 8.17% without clustering.

2.5.4.3 Classical-ML Hybrid Frameworks

A growing strand hybridises statistical models with machine learning components to capture the interpretable trend or seasonality modelling of classical methods alongside the non-linear pattern recognition of ML. Punia et al. (2020) combine LSTM with Random Forest, finding their outputs complementary. Ungureanu (2025) uses SARIMAX as a baseline and XGBoost as an advanced model, comparing them directly and finding XGBoost outperforms SARIMAX in RMSE (38.7 vs. 45.2), MAE (27.1 vs. 31.7), and MAPE (11.2% vs. 13.5%). Taparia et al. (2023) layer ML regression on top of exponential smoothing predictions, creating a hybrid where the statistical model's output becomes one input feature to the ML model capturing both the interpretability of exponential smoothing and the non-linear correction capacity of the ML layer.

2.5.4.4 Reinforcement Learning for Inventory Policy

A smaller but methodologically distinct strand applies deep reinforcement learning (DRL) directly to inventory policy optimization, rather than using ML only for forecasting. Kalusivalingam et al.

propose a hybrid DRL and neural network approach where the system learns optimal ordering policies through simulated inventory scenarios, dynamically adjusting to fluctuating market conditions without requiring explicit demand forecasts. Wahedi et al. (2023) evaluate Q-learning and Deep Q Network (DQN) reinforcement learning approaches for inventory management alongside predictive ML models, finding that predictive models are superior on accuracy metrics but noting that RL approaches offer the advantage of directly optimising inventory policy objectives rather than minimising forecast error.

2.5.5 Specialised Applications and Extensions

Several papers apply ML forecasting to specific product categories or business contexts. Khayyat and Gupta (2025, Smartlogistic_paper) develop an advanced predictive model for inventory management and demand forecasting specifically for smart logistics environments, integrating IoT-generated data streams with ML models. Bendhi (2023, IJAIDSMLV4I2P108) demonstrates ML-driven demand forecasting reducing excess stock by up to 20% in a retail chain implementation, using a comprehensive methodology incorporating time-series sales data, promotional calendars, and economic indicators. Farooq et al. (2024) develop a conceptual framework linking inventory optimization with sustainability goals, arguing that data-driven resource management reduces food waste and carbon emissions as well as costs. Mweshi (2025, OptimizingSmartCommerce) applies ML to dynamic demand forecasting for both private retail/e-commerce and public sector logistics (healthcare, defence), one of the few studies to extend beyond commercial retail contexts. Yadav (2025, IJARCCE) proposes incorporating social media trends and local event data into ML-based demand forecasting, extending the feature set beyond traditional sales and promotional variables to capture real-world demand influencers.

Aryyama (2021, optimizationofecommercesupplychain) addresses the specific challenge of demand prediction for new products in e-commerce supply chains, where the absence of historical sales data requires different approaches to feature construction and model initialisation. Singhal, Singh, and Kaul (2024, Smart_Retail) demonstrate a Random Forest Regressor achieving an MSE of 8.15% for demand forecasting and 1.11% for price prediction in a smart retail context, highlighting the dual use of ML for both demand and pricing optimization.

2.6 Inventory Optimisation Frameworks and Operational Integration

2.6.1 The Order-Up-To-Level Policy

The OUTL (R, s, S) policy is the most widely used inventory framework across the operationally integrated studies. Under this policy, inventory is reviewed at fixed intervals R ; if the inventory position falls at or below the reorder point s , an order is placed to raise inventory to the order-up-to level S . The key formulas $SS = z \cdot \sigma \cdot \sqrt{(R+L)}$; $r = d(R+L) + z \cdot \sigma \cdot \sqrt{(R+L)}$; $S = r + Q$ appear consistently across Barghi (2025) and Seyedan et al. (2023), with L the lead time, d the average forecasted demand, z the service-level z -score, and σ the standard deviation of demand. Total cost $TC = PC + HC + OC$ is the optimisation target.

Barghi (2025) compares the Generic and Fuzzy Hybrid Boosting Ensembles within this framework at a 7-day lead time: the fuzzy model reduces total inventory cost from \$25.60 million to \$22.40 million (12.5%), lowers the reorder point from 116,062 to 104,963 (9.6%), and reduces the order quantity from 97,416 to 85,220 (12.5%), while slightly increasing safety stock from 18,646 to 19,742 (5.9%) as a more conservative buffer against residual uncertainty. These cost savings persist consistently across lead times of 5, 10, 15, 20, and 25 days. Seyedan et al. (2023) demonstrate that the heterogeneous ensemble consistently achieves lower total cost than any individual base model across 7, 10, and 14-day lead time scenarios for both sports (cost savings of 97–115 units over MLP) and electronics (savings of 90–175 units).

2.6.2 Safety Stock and Service Level Calibration

Seyedan (2023, Chapter 5) provides the most detailed analysis of safety stock calibration across different service levels and lead times, comparing RMDN and LSTM approaches. At a 90% service level with a 7-day lead time, RMDN requires safety stock of 20 units and an order quantity of 234 units, compared to 85 units of safety stock and 425 units order quantity for LSTM, a difference attributable to RMDN's distributional forecasting providing more confident demand characterisation than LSTM's point estimates. At a 99% service level with a 14-day lead time, RMDN safety stock rises to 52 units versus 219 for LSTM, demonstrating that the relative advantage of distributional forecasting is largest at high service levels and long lead times, precisely where safety stock decisions are most consequential.

Mittal (2025, [ssrn5386001](#)) operationalises probabilistic forecasting-to-inventory translation through a quantile-based mechanism, converting predictive distribution quantiles into safety stock and reorder parameters that explicitly target service levels while minimising total inventory cost. This approach achieves 14–21% reductions in total inventory cost versus statistical baselines while maintaining target service levels, demonstrating that uncertainty modelling and inventory policy can be integrated more tightly than the safety stock formula approach allows.

2.6.3 Simulation and Operational Validation

Ungureanu (2025) uses inventory simulation as the bridge between forecasting accuracy and operational impact, implementing a daily ordering rule ($\text{order_qty}(t) = \text{forecast}(t) + \text{safety_buffer}$) and tracking stock-outs and waste over a November 2016–January 2017 test period including holiday peaks and promotional events. The XGBoost-based reorder policy reduces stock-out rate from 8.5% (naive 7-day moving average baseline) to 2.7%, a 68% relative reduction while simultaneously decreasing perishable waste from 1,200 to 750 units (37.5%), reducing estimated disposal costs by approximately €900. These results exceed the study's 15% stock-out reduction target. Pasupuleti et al. (2024) report even larger operational improvements from a ML-driven logistics and inventory system applied to a multinational retail corporation: 15% improvement in demand forecasting accuracy, 10% reduction in overstock and stockouts, and 95% accuracy in predicting order fulfilment timelines.

Nasution, Matondang, and Ishak (2022) design an inventory optimisation model for a feed mill company using ML-based demand forecasts, demonstrating that ML-driven demand forecasting improves inventory performance metrics and reduces the costs associated with overstock and understocking in an industrial rather than retail context. Wahedi et al. (2023) evaluate inventory KPIs directly for the Svanehøj SME, finding that ML forecasts do not consistently outperform statistical methods when measured against industrial KPIs rather than pure accuracy metrics, an important qualification that highlights the gap between forecasting accuracy and operational value in constrained industrial settings.

2.7 Evaluation Metrics and Comparative Results

2.7.1 Standard Accuracy Metrics

Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and Mean Absolute Percentage Error (MAPE) are the universal accuracy metrics across the review. MAE measures average absolute error in demand units, directly interpretable by store managers. RMSE weights large errors more heavily through squaring, making it sensitive to promotional spikes and holiday surges where under-forecasting is most damaging. MAPE normalises errors as percentages of actual demand, enabling cross-product comparisons. Several studies add supplementary metrics: Pred(x=10%); the percentage of predictions within 10% of actuals used by Barghi (2025); R^2 (coefficient of determination) used by Mitra et al. (2022) and Islam et al. (2024); and SMAPE (Symmetric MAPE) used by Abbasimehr et al. (referenced in Seyedan, 2023, and Barghi, 2025).

2.7.2 Comparative Performance Across Studies

Table 2: Key Forecasting Performance Results Across Selected Studies

Study / Method	Dataset	Best Model MAPE (%)	Best Model MAE	Best Model RMSE
Barghi (2025) – Fuzzy Hybrid Ensemble	Rossmann (1,115 stores)	3.53	209.21	272.99
Barghi (2025) – Generic Hybrid Ensemble	Rossmann	6.48	416.77	608.66
Seyedan et al. (2023) – Ensemble DL (Sports)	Constante Supply Chain	5.22	19.46	24.02
Seyedan et al. (2023) – Ensemble DL (Electronics)	Constante Supply Chain	9.58	3.56	4.99
Seyedan et al. (2022) – Cluster-BMA Ensemble	Constante Sports Data	8.12	27.10	32.80
Nasseri et al. (2023) – ETR	Austrian Retail (5.2M records)	~8–12 (by category)	—	—
Ungureanu (2025) – XGBoost	Favorita/UCI Grocery	11.2	27.1	38.7
Ungureanu (2025) – SARIMAX Baseline	Favorita/UCI Grocery	13.5	31.7	45.2
Mitra et al. (2022) – RF-XGBoost-LR Hybrid	US Retail (Walmart-style)	Best across MAE/MSE/ R^2	—	—

Study / Method	Dataset	Best Model MAPE (%)	Best Model MAE	Best Model RMSE
Khan et al. (2020) – BI + ML System	Real-time org data	~7.62% error	—	—
Mittal (2025) – TFT/LSTM/ST-GNN	Synthetic multi-channel retail	22–35% gain vs baseline	—	—

2.7.3 Feature Importance and Key Demand Drivers

Feature importance analysis across studies converges on a consistent hierarchy of demand drivers. Promotional flags are the dominant predictor across tree-based models: Ungureanu (2025) reports that promotional days boost daily sales by 25–40% for perishable items and 10–20% for non-perishables, while Seyedan et al. (2022) find that holiday periods increase daily demand by 40–60% above typical weekday levels. Lagged sales (particularly 7-day lags reflecting weekly seasonality) rank highly across all ML approaches. Day-of-week indicators are consistently significant: Saturday sales exceed midweek averages by 25–30% in Ungureanu (2025). Store type and cluster attributes matter for models trained across heterogeneous store networks.

Haque et al. (2023) add macroeconomic variables to this picture, finding that CPI, ICS, and unemployment rates improve retail demand predictions by capturing macro-level shifts in consumer purchasing power that internal sales data alone cannot anticipate. Nasseri et al. (2023) integrate meteorological and COVID-19 data, demonstrating that external shocks significantly improve model accuracy during periods of unusual demand. Kalisetty (2023) incorporates urban land-use context effects through DemandNet, capturing the spatial dimensions of retail demand that pure temporal models miss.

2.8 Research Gaps and Limitations

2.8.1 Single-Echelon Focus and Multi-Echelon Neglect

The dominant research design across the review restricts analysis to a single-echelon retail system, the retailer's interaction with customers excluding upstream supplier and distributor dynamics. Barghi (2025) explicitly acknowledges this restriction, noting the study simplifies the supply chain

by excluding upstream echelons. Seyedan et al. (2023) conduct their analysis within an online retail context without modelling the full supply chain network. While this simplification enables methodological focus, it means the proposed frameworks cannot capture the bullwhip effect demand variability amplifying upstream which is identified by Mohammed and Mandal (2022), Nasution et al. (2022), and Razzak et al. (2025) as a major motivation for improved demand forecasting. Future research integrating demand forecasting models across multi-echelon networks, with explicit coordination mechanisms, represents an important unsatisfied need.

2.8.2 Point Forecasting Dominance and Probabilistic Neglect

The vast majority of studies reviewed generate point forecasts single-value predictions of expected demand without quantifying the uncertainty around those predictions. As Seyedan (2023) argues, this is a fundamental limitation because safety stock calculations require knowledge of demand variability, and using point forecast error statistics as proxies for this variability introduces systematic inaccuracies. Mittal (2025) and Seyedan (2023, Chapter 5) are notable exceptions, both generating distributional forecasts. The literature has not yet fully explored probabilistic forecasting architectures including Bayesian neural networks, conformal prediction, and natural gradient boosting that could provide calibrated uncertainty estimates across the full demand distribution.

2.8.3 Absence of Real-Time Deployment Validation

A cross-cutting limitation is that virtually all studies validated through historical data simulation rather than live operational deployment. Barghi (2025) acknowledges that empirical validation was confined to simulations due to time constraints and that real-world implementation was not feasible. Ungureanu (2025) explicitly recommends a controlled pilot deployment in 3–5 stores as the next step. Wahedi et al. (2023) are somewhat closer to operational validation through their SME partnership, but even they do not implement the ML-based recommendations in live inventory systems. The gap between in-sample simulation performance and out-of-sample operational performance under live data streams, model drift, and organisational constraints remains unaddressed.

2.8.4 Exclusion of Shortage and Perishability Costs

Most inventory-integrated studies exclude shortage costs and, for non-perishable-focused studies, obsolescence and waste costs. Barghi (2025) notes that stockout costs are excluded and that the dataset lacks stock-level data to confirm whether shortages occurred on zero-sales days. Namwad et al. (2024) are among the few who incorporate carbon pollution costs and product deterioration explicitly, but their fuzzy model is a theoretical construction rather than an empirically validated deployment. The financial and environmental costs of food waste, central to Ungureanu (2025) and Farooq et al. (2024), are still not consistently quantified in the broader literature, and the trade-off between stockout and waste costs for perishable categories remains underexplored.

2.8.5 SME and Mid-Market Accessibility

The literature remains disproportionately focused on large enterprise contexts. Wahedi et al. (2023) and Ungureanu (2025) explicitly identify the SME gap and design for resource-constrained operators, but they represent a minority of the reviewed studies. Several papers present technically sophisticated frameworks requiring substantial computational infrastructure (GPU training, distributed Spark environments) that would be inaccessible to typical mid-market retailers. The challenge of deploying advanced ML in operationally realistic SME contexts, with limited data engineering capacity, absence of real-time data pipelines, and dependence on legacy ERP systems, is acknowledged but not systematically addressed.

2.8.6 Ethical, Governance, and Sustainability Dimensions

The vast majority of studies treat demand forecasting as a pure technical optimisation problem, ignoring questions of algorithmic fairness, data governance, and sustainability impact. Ungureanu (2025) is the exception, devoting substantial attention to GDPR compliance, data minimisation, bias auditing (checking whether models systematically underforecast in lower-income or rural store locations), and the use of SHAP analysis for transparency. Farooq et al. (2024) develop a sustainability-oriented conceptual framework linking inventory optimization to circular economy principles and environmental stewardship. Mohammed and Mandal (2022) touch on the ethical dimensions of dynamic pricing optimization through reinforcement learning. But none of the studies systematically investigate whether ML-based inventory systems produce equitable outcomes across different customer demographics, geographic markets, or store formats; a gap

that will become more pressing as regulatory scrutiny of algorithmic decision-making in supply chains increases.

2.8.7 Concept Drift and Model Maintenance

The challenge of model degradation over time, as demand patterns evolve, new products are introduced, and external conditions shift, receives limited attention. Ungureanu (2025) recommends periodic retraining and automated error monitoring, noting the dataset's 2015–2017 vintage means consumer behaviour, product assortments, and promotional strategies may have changed substantially. Barghi (2025) mentions regular model retraining with fresh data and hyperparameter tuning as mitigation strategies. But none of the studies empirically examine how quickly model performance degrades after training, what drift detection mechanisms are appropriate for retail demand, or how to balance the cost of retraining against the cost of performance degradation, questions that are central to any real-world deployment.

2.9 Directions for Future Research

The gaps identified in Section 8 generate a clear research agenda. Across the 52 sources, the most consistently recommended directions are:

Multi-echelon supply chain integration. Barghi (2025), Seyedan et al. (2023), and Mohammed and Mandal (2022) all identify the extension of retail-level demand forecasting to multi-echelon contexts, incorporating supplier lead times, distributor dynamics, and the bullwhip effect, as a natural and important next step. Joint optimisation of demand forecasting and replenishment decisions across networked inventory systems represents a methodologically challenging but high-impact research direction.

Probabilistic and distributional demand forecasting. The RMDN framework introduced by Seyedan (2023) and the TFT/ST-GNN probabilistic forecasting framework of Mittal (2025) represent early steps toward a more principled treatment of demand uncertainty. Future work extending these approaches to gradient boosting frameworks, incorporating conformal prediction, or applying Bayesian neural networks to retail demand would fill a significant methodological gap and directly improve safety stock calibration.

Real-time and streaming data integration. Ungureanu (2025), Kalisetty (2023), and Khayyat and Gupta (2025) all point to real-time data streams, POS feeds, IoT sensors, social media signals, local event calendars, as the next frontier for demand forecasting. Ungureanu (2025) recommends converting batch ingestion routines to Spark Structured Streaming flows. The development of forecasting models that update continuously as new data arrives, rather than retraining periodically on fixed historical windows, would represent a significant practical advance.

Multi-SKU and joint inventory optimization. Barghi (2025) and Seyedan (2023) both note that their frameworks address single products, while real-world retail involves thousands of interdependent SKUs. Coordinated replenishment across related items, joint ordering to reduce logistics costs, and ABC-analysis-informed prioritisation of forecasting effort across SKU portfolios are all identified as important extensions.

Ethical AI, fairness, and sustainability. Ungureanu (2025) and Farooq et al. (2024) point to the need for systematic investigation of equitable outcomes across store formats and geographic markets, transparent and auditable algorithmic decision-making, and explicit sustainability accounting that treats food waste and carbon emissions as objectives rather than afterthoughts. Future research should incorporate fairness constraints into model training, test for systematic bias in inventory allocation across demographic groups, and develop environmental impact metrics alongside economic efficiency measures.

Transfer learning and cross-retailer generalisation. Several studies acknowledge dataset specificity as a limitation. Models trained on Rossmann data may not generalise to grocery or fashion retail. Transfer learning approaches that pre-train on large retail datasets and fine-tune on specific retailer data would reduce data requirements and improve accessibility for smaller operators. This is identified by Razzak et al. (2025) as a frontier area where deep learning's capacity for knowledge transfer could be particularly valuable.

2.10 Synthesis: Convergent Findings and Key Contributions

2.10.1 What the Literature Has Established

The 52 sources reviewed provide robust convergent evidence for several conclusions. First, ensemble machine learning methods; whether gradient boosting sequential ensembles,

heterogeneous deep learning meta-ensembles, or stacking frameworks; consistently and substantially outperform both classical statistical methods and individual ML models across diverse retail datasets, sectors, and geographic contexts. Second, data preprocessing and feature engineering contribute as much to predictive performance as model architecture: the fuzzy generalisation step in Barghi (2025) reduces MAPE by 45.5% without changing the ensemble structure; lag features and rolling averages consistently improve accuracy across tree-based models. Third, the integration of accurate forecasts with inventory replenishment policy is essential to translating predictive performance into operational value, improved MAPE without OUTL integration produces no direct cost savings. Fourth, key demand drivers are consistent across studies: promotional flags, holiday proximity indicators, day-of-week features, and lagged sales constitute the core feature set that explains most of the forecasting advantage of ML over statistical baselines.

2.10.2 Where the Literature Converges and Diverges

The literature converges clearly on the superiority of ensemble approaches, the importance of multivariate feature sets, and the need for forecast-inventory integration. It diverges on the optimal ensemble architecture: gradient boosting sequential ensembles (Barghi, 2025; Mitra et al., 2022) versus heterogeneous deep learning stacks (Seyedan et al., 2023; Punia et al., 2020) versus classical-ML hybrids (Nasseri et al., 2023; Ungureanu, 2025), reflecting dataset-specific advantages that do not generalise universally. The literature also diverges on the value of LSTM relative to tree-based methods: Nasseri et al. (2023) find Extra Tree Regressors outperform LSTM on their large dataset, while Seyedan et al. (2023) find LSTM provides the best individual base model and contributes distinctively to ensemble performance. This divergence suggests that the optimal architecture is context-dependent, and that systematic meta-analysis across datasets and methods is needed to establish reliable guidance.

2.10.3 Contribution of This Review

This literature review contributes to the field by providing the first synthesis that draws on 52 sources spanning the full methodological range from classical exponential smoothing to probabilistic deep learning, across retail contexts from grocery to fashion to industrial supply chains, and from single-paper algorithmic contributions to book-length thesis treatments. By synthesising across rather than within these sources, the review identifies cross-cutting patterns,

the consistent value of ensemble approaches, the underexplored terrain of distributional forecasting, the persistent SME accessibility gap, the absence of real-time deployment validation that are not visible from any single study. The review also provides a structured reference for practitioners seeking to select and implement ML-based forecasting and inventory systems, offering a framework for understanding which approaches are appropriate for which contexts.

2.11 Conclusion

The machine learning revolution in retail demand forecasting is well underway. The 52 sources reviewed here document a consistent, robust, and growing body of evidence that ensemble machine learning methods, applied to multivariate retail datasets with careful feature engineering and integrated with operational inventory replenishment policies, deliver substantial improvements in forecast accuracy, cost efficiency, service levels, and waste reduction relative to classical statistical approaches. The Rossmann benchmark demonstrates reductions in MAE of up to 86% from ARIMA to hybrid gradient boosting ensemble. Applied to mid-sized grocery retail, XGBoost-based reorder policies cut stock-out rates from 8.5% to 2.7% and reduce perishable waste by 37.5%. Probabilistic forecasting through RMDNs reduces required safety stock by 75% at identical service levels compared to LSTM point forecasting.

Yet the field retains important blind spots. The single-echelon focus, the dominance of historical simulation validation over live deployment, the neglect of distributional uncertainty in most forecasting frameworks, the inaccessibility of many proposed systems to SME operators, and the almost complete absence of fairness and sustainability accounting are not peripheral concerns, they are central to whether the forecasting improvements documented in the laboratory translate into equitable, sustainable, and economically beneficial supply chain operations in the real world.

The research agenda is rich. Multi-echelon integration, real-time streaming architecture, distributional forecasting, multi-SKU coordination, cross-retailer transfer learning, and ethically governed algorithmic decision-making all represent high-impact directions where the current literature has opened doors without walking through them. The present study contributes to this agenda by providing a comprehensive, critical, and structured synthesis of the existing literature, establishing both the strong empirical foundation on which future work can build and the precise gaps that most urgently need to be filled.

CHAPTER 3

Methodology

3.1 Introduction

This chapter presents the research design, data sources, preprocessing pipeline, feature engineering strategy, modelling framework, evaluation methodology, and inventory policy integration procedure that constitute the empirical backbone of this study. Every decision documented here is anchored in the four research gaps identified in Chapter 2 and guided by the five objectives stated in Chapter 1: comparative evaluation of forecasting models, systematic feature engineering, integration of the best-performing model with the Order-Up-To-Level (OUTL) inventory policy, quantification of operational impact relative to classical baselines, and delivery of results through an accessible Power BI dashboard.

The overarching methodological framework is the Cross-Industry Standard Process for Data Mining (CRISP-DM), which structures the research as an iterative, six-phase process; Business Understanding, Data Understanding, Data Preparation, Modelling, Evaluation, and Deployment, progressing from problem definition through to operational output (Barghi, 2025; Ungureanu, 2025). CRISP-DM was selected over alternative frameworks because its iterative structure accommodates the cycles of feature refinement and model calibration that characterise retail demand forecasting research, because it has been adopted in closely related studies (Bastos, 2023; Ungureanu, 2025), and because its explicit deployment phase anchors every technical decision in the operational business problem.

The remainder of this chapter is structured as follows. Section 3.2 describes the research design philosophy and the role of CRISP-DM. Section 3.3 provides a detailed description of the M5-Forecasting (Walmart) dataset and the stratified subsample used for modelling. Section 3.4 documents the data preprocessing pipeline. Section 3.5 presents the feature engineering strategy. Section 3.6 describes model development and the comparative evaluation framework across five model tiers. Section 3.7 details the OUTL inventory policy integration procedure, including a quantile-based robustness check and stockout cost modelling. Section 3.8 outlines the Power BI dashboard development approach. Section 3.9 explains the sensitivity analysis design. Section 3.10

addresses ethical considerations governing the study. Section 3.11 specifies the SME transferability requirements. Section 3.12 summarises the chapter.

3.2 Research Design and Methodological Framework

3.2.1 Research Philosophy

This study is positioned within a post-positivist research philosophy. It seeks to identify causal relationships between modelling architecture choices and operational inventory outcomes, using quantitative measurement of forecast accuracy and inventory cost as primary evidence. The study does not claim to generate universal laws but rather to produce replicable, empirically grounded findings that can inform decision-making in comparable retail environments (Seyedan, 2023). A deductive approach is adopted: hypotheses about model superiority and inventory cost reduction are formulated from the literature in Chapter 2 and tested against the M5 dataset.

The research strategy is a quantitative computational experiment, a form of empirical investigation in which multiple modelling approaches are systematically compared on a shared dataset under controlled conditions (Nasseri et al., 2023; Mitra et al., 2022). This strategy was chosen over survey or case study designs because it permits direct, metric-grounded comparison of forecasting and inventory outcomes, which is precisely what the literature gaps identified in Chapter 2 require.

3.2.2 CRISP-DM as the Structuring Framework

The CRISP-DM framework organises the research into six iterative phases. Phase 1 (Business Understanding) is documented in Chapter 1. Phase 2 (Data Understanding) is addressed in Sections 3.3 and 3.4. Phase 3 (Data Preparation) covers Sections 3.4 and 3.5. Phase 4 (Modelling) is addressed in Section 3.6. Phase 5 (Evaluation) spans Sections 3.6 and 3.7 and Chapter 4. Phase 6 (Deployment) is addressed in Section 3.8 through the Power BI dashboard. The iterative character of CRISP-DM is reflected in feedback loops between feature engineering and model evaluation: features were refined based on pilot model performance, and the inventory policy formulation was extended with a quantile-based robustness check after preliminary results revealed the limitations of the Gaussian assumption on intermittent demand series.

3.3 Data Source: The M5-Forecasting (Walmart) Dataset

3.3.1 Dataset Description and Provenance

The primary data source for this study is the M5-Forecasting (Walmart) dataset, sourced from the Kaggle platform (Obi, 2024; Haque et al., 2023). The dataset was originally prepared for the M5 Accuracy and Uncertainty competitions hosted by the Makridakis Open Forecasting Center and provides one of the most comprehensive publicly available benchmarks for hierarchical retail demand forecasting.

The dataset contains daily unit sales for 3,049 individual product-store series, organised across three product categories, Hobbies, Foods, and Household, and seven product departments. Products are sold across ten Walmart stores in three US states: California (CA1–CA4), Texas (TX1–TX3), and Wisconsin (WI1–WI3). The publicly distributed `sales_train_validation.csv` file spans 1,913 consecutive days from 29 January 2011 to 24 April 2016. The dataset is distributed across three raw files: `sales_train_validation.csv` (daily unit sales), `calendar.csv` (date, day-of-week, SNAP eligibility flags, event labels), and `sell_prices.csv` (weekly selling prices per product-store combination).

3.3.2 Justification of Dataset Choice

The M5 dataset was preferred over alternatives such as the Rossmann Store Sales dataset (Barghi, 2025) or the Favorita Grocery Sales dataset (Ungureanu, 2025) for four specific reasons. First, its scale, 3,049 product-store series over more than five years, provides a richer basis for model generalisation than smaller benchmarks. Second, its hierarchical structure mirrors the organisational structure of real SME retail operations, making the methodology transferable. Third, the inclusion of enriched contextual variables, SNAP benefit eligibility flags, granular event categorisation, and weekly price data, reflects the types of operational data that medium-sized retailers already collect, fulfilling the SME accessibility principle stated in Chapter 1. Fourth, it constitutes an established academic benchmark, permitting direct comparison with published results from Obi (2024), Haque et al. (2023), and the broader M5 retrospective literature.

3.3.3 Stratified Subsample for Computational Tractability

The full M5 dataset expands to approximately 56.6 million product-store-day records once reshaped from wide to long format. While computationally tractable for tree-based gradient

boosting, this scale prohibits rigorous walk-forward cross-validation and recurrent neural network training within the project's compute budget. A precedent for working with subsamples exists in the literature: Nasser et al. (2023) report results on a subset of approximately 5.2 million records, and Seyedan et al. (2023) restrict their evaluation to selected product categories.

This study uses a stratified random subsample of 502 product-store series, drawn proportionally across the nine strata defined by the three product categories crossed with the three states. The sampling is performed with a fixed random seed (42) to ensure reproducibility. The subsample preserves the categorical structure of the parent dataset while reducing total observations to approximately 933,000, a 60-fold reduction that enables identical model architectures to be trained across all five tiers within consistent computational constraints, and permits the LSTM and Random Forest models to be evaluated rigorously rather than approximately. Table 3.1 reports the subsample composition.

Table 3.1: Subsample composition by category and state

Category	California	Texas	Wisconsin	Row total
Foods	94	71	71	236
Hobbies	37	28	28	93
Household	69	52	52	173
Column total	200	151	151	502

3.4 Data Preprocessing Pipeline

3.4.1 Data Integration

The three source files are merged into a single unified modelling dataset prior to any preprocessing. The sales file is first reshaped from wide format (one column per day) to long format (one row per product-store-day combination). The calendar file is joined on the date field, adding day-of-week, month, year, event type, event name, and SNAP flag variables. The sell_prices file is joined on the item_id, store_id, and wm_yr_wk composite key, adding the sell_price variable. The integrated long-format dataset is then filtered to the stratified subsample of 502 series.

3.4.2 Handling Missing Values

Missing values in the `sell_price` variable arise because prices are only recorded for weeks in which a product was actively stocked. Forward-fill imputation is applied first, using the most recent observed price for that product-store combination. Where no prior price observation exists, backward-fill imputation is applied from the first available price. This approach preserves the temporal structure of price variation and is consistent with retail practice, where products carry a stable price between changes (Nasseri et al., 2023; Barghi, 2025). Missing values in calendar event fields are treated as absence of an event and imputed with the label "No Event". No missing values are present in the core sales columns, as zero-sales records are valid observations.

3.4.3 Handling Zero-Sale Records

Zero-sale records may represent genuine zero demand, temporary stockouts, or store closures. Following Barghi (2025), zero-sales records corresponding to confirmed store closure days, identifiable from the `calendar_event_type` field, are removed from the training set, as they do not represent purchasable demand. Zero-sale records on non-closure days are retained, as they represent genuine demand signals that are particularly important for safety stock calibration on intermittent series.

3.4.4 Outlier Detection and Treatment

Extreme sales outliers, defined as daily unit sales exceeding three standard deviations above the rolling 28-day mean for a given product-store combination, are identified and flagged. Outliers corresponding to confirmed promotional events or special calendar events are retained, as they represent genuine demand spikes that the model should learn to predict. Outliers that cannot be explained by any calendar or promotional event are winsorised to the 99th percentile for that product-store series, preventing extreme values from distorting gradient-based model training (Nasseri et al., 2023).

3.4.5 Categorical Variable Encoding

Identifier and categorical variables, `item_id`, `dept_id`, `cat_id`, `store_id`, `state_id`, `event_type`, and `event_name`, are encoded using label encoding. Label encoding is preferred over one-hot encoding for tree-based models because it preserves ordinality within compact integer representations without generating high-dimensional sparse matrices, which would substantially increase memory

consumption. LightGBM natively handles label-encoded categoricals through its histogram-based splitting mechanism (Barghi, 2025; Obi, 2024).

3.4.6 Target-Encoded Mean Features

To complement the label-encoded identifiers, eight target-encoded mean features are constructed: `item_mean`, `dept_mean`, `store_mean`, `state_mean`, `item_dow_mean`, `store_dow_mean`, `cat_month_mean`, and `item_month_mean`. Each is the mean of the target variable computed exclusively over the training data, grouped by the relevant identifier(s). Target encoding is applied because it provides a strong prior for the expected demand level of each product-store combination, separating that signal cleanly from the temporal lag and rolling features. To prevent data leakage, all target encodings are fitted on the training window only and propagated unchanged to the validation and test windows. Missing combinations (categories not seen in training) are imputed with the global training mean. This technique is reported by M5 competition retrospectives as a significant contributor to top-performing solutions (Obi, 2024).

3.5 Feature Engineering

Feature engineering is the process of constructing derived variables from raw data that capture the demand patterns the models need to learn. The literature review identified feature engineering as contributing as much to predictive performance as model architecture choice: Barghi (2025) demonstrated that preprocessing alone reduced MAPE by 45.5% in a comparable retail forecasting study. The feature engineering strategy in this study is organised across six categories.

3.5.1 Lag Features

Lag features capture temporal autocorrelation by including historical sales values as predictors. Three lag windows are constructed for each product-store series: `lag_7` (sales 7 days prior, capturing weekly cyclical patterns), `lag_14` (sales 14 days prior, capturing bi-weekly patterns), and `lag_28` (sales 28 days prior, capturing monthly demand cycles and promotional recurrence). These windows are consistent with those used by Nasser et al. (2023) and Obi (2024) on the M5 dataset.

3.5.2 Rolling Average Features

Rolling averages smooth demand signals across multiple time windows, providing estimates of recent trend and mean demand levels. Three rolling statistics are constructed: `rolling_mean_7` (7-

day rolling mean, capturing short-term trend), `rolling_mean_28` (28-day rolling mean, capturing medium-term trend and reducing noise from single-day spikes), and `rolling_std_7` (7-day rolling standard deviation, providing an estimate of recent demand volatility for safety stock calibration). All rolling statistics are computed over strictly historical windows to prevent data leakage.

3.5.3 Calendar and Temporal Features

Calendar features decompose the date variable into components that capture known demand seasonality patterns: `day_of_week` (0–6), `week_of_year` (1–52), `month` (1–12), `year`, `is_weekend` (binary indicator), `days_since_last_event` (count of days since the most recent calendar event, capturing pre-event demand run-up), and `days_to_next_event` (count of days until the next calendar event, capturing anticipation effects). Day-of-week is consistently the strongest temporal predictor across the literature; Seyedan et al. (2022) found Saturday sales exceeded midweek averages by 25–30% in comparable retail data.

3.5.4 Promotional and Event Features

Promotional and event indicators are among the strongest predictors of demand spikes in retail contexts (Ungureanu, 2025; Nasser et al., 2023). The following binary indicators are constructed from the calendar file: `is_snap_day` (SNAP benefit eligibility in the relevant state, which drives demand for food products), `is_sporting_event`, `is_cultural_event`, `is_national_event`, and `is_religious_event` (derived from the `event_type` field).

3.5.5 Price Features

Two price-related features are constructed from the `sell_prices` file: `sell_price` (the weekly selling price of each product-store combination, a direct demand driver through price elasticity effects) and `price_change` (the week-on-week price difference, capturing the direction and magnitude of price movements and their potential demand impact).

3.5.6 Target-Encoded Mean Features

As specified in §3.4.6, eight target-encoded mean features are added to the feature set: `item_mean`, `dept_mean`, `store_mean`, `state_mean`, `item_dow_mean`, `store_dow_mean`, `cat_month_mean`, and `item_month_mean`. Table 3.2 summarises the complete engineered feature set.

Table 3.2: Complete feature engineering summary

Category	Features	Demand signal captured
Lag	lag_7, lag_14, lag_28	Historical sales at 7-, 14-, 28-day offsets
Rolling statistics	rolling_mean_7, rolling_mean_28, rolling_std_7	Short- and medium-term trend; volatility
Temporal	day_of_week, week_of_year, month, year, is_weekend	Weekly, seasonal and annual cycles
Event proximity	days_since_last_event, days_to_next_event	Pre- and post-event demand run-up and decay
Promotional / event	is_snap_day, is_sporting_event, is_cultural_event, is_national_event, is_religious_event	SNAP eligibility; event-type spikes
Price	sell_price, price_change	Price level and week-on-week movement
Identifier (label-encoded)	item_id, dept_id, cat_id, store_id, state_id	Cross-series generalisation
Target-encoded mean	item_mean, dept_mean, store_mean, state_mean, item_dow_mean, store_dow_mean, cat_month_mean, item_month_mean	Demand-level prior per item, store, day-of-week and month

3.6 Model Development and Evaluation Framework

3.6.1 Model Selection Rationale

Five model tiers are evaluated in this study, structured to provide a systematic performance progression from classical statistical baselines through to a quantile-regression model that supports the inventory policy integration. This tiered design reflects the comparative methodology adopted in the most rigorous studies reviewed in Chapter 2, notably Barghi (2025), Nasseri et al. (2023), Mitra et al. (2022), and Seyedan et al. (2023), and ensures that the performance improvement claimed for the proposed ensemble is substantiated against the full range of alternatives a practitioner might deploy.

Tier 1. Classical statistical baselines. Six classical models are implemented as performance baselines. The Moving Average (MA, 28-day window) provides a naïve demand estimate with no external features. Simple Exponential Smoothing (SES) applies exponentially decreasing weights to historical observations. AutoARIMA performs automatic parameter selection by AIC minimisation and is configured with a seasonal component of length seven to capture weekly cycles, encompassing both ARIMA and SARIMA as required by the data. Croston's Classic

method and the Syntetos-Boylan Approximation (Croston-SBA) are added in recognition that retail SKU-daily demand is intermittent and contains many zero-sale days, conditions under which conventional methods like ARIMA and SES are known to underperform (Croston, 1972; Syntetos & Boylan, 2005). Seasonal Naive is included as a transparent reference baseline that simply forecasts the value observed seven days prior. These baselines represent methods that many SME retailers currently deploy in some form (Wahedi et al., 2023; Keith, 2023). All Tier 1 models are implemented via the statsforecast Python library, which provides parallel per-series fitting.

Tier 2. Standard machine learning models. Random Forest (RF) is a bagging ensemble of decision trees providing variance reduction and tolerance for non-linear interactions (Nasseri et al., 2023), implemented using scikit-learn. XGBoost is a gradient boosting algorithm with L1/L2 regularisation; LightGBM is a gradient boosting algorithm with leaf-wise tree growth and histogram-based binning that is well-suited to large tabular datasets (Obi, 2024; Barghi, 2025). For this study, both XGBoost and LightGBM are configured with the Tweedie loss function (variance power = 1.1), which is designed for non-negative count data with many zero observations and is consistently reported by M5 retrospective analyses as a significant contributor to top-performing solutions on intermittent demand.

Tier 3. Deep learning time-series model. A Long Short-Term Memory (LSTM) network captures long-range temporal dependencies through gated memory cells, learning multi-day promotional dynamics and seasonal rhythms that tree-based models treat as independent feature interactions (Punia et al., 2020; Nasseri et al., 2023). The LSTM architecture comprises two stacked recurrent layers (128 and 64 units respectively), each followed by a dropout layer with rate 0.2 for regularisation, and a single dense output unit. Input sequences use a 28-day lookback window. The target variable is transformed using $\log(1+p)$ prior to training to handle the zero-inflated demand distribution and stabilise gradient updates; predictions are inverse-transformed using $\exp(m)$ and clipped at zero. Training uses the Adam optimiser with mean-squared-error loss in log space, with early stopping on validation loss (patience = 5 epochs). Implementation uses TensorFlow/Keras.

Tier 4. Stacking ensemble (primary methodological contribution). The primary modelling contribution of this study is a stacking ensemble in which the three Tier 2 base learners, LightGBM, XGBoost, and Random Forest, produce predictions over the validation window, and

these predictions serve as features for a Ridge regression meta-learner that learns the optimal combination weights:

$$\hat{y}_{ensemble} = w_{LGBM} \cdot \hat{y}_{LGBM} + w_{XGB} \cdot \hat{y}_{XGB} + w_{RF} \cdot \hat{y}_{RF} + b$$

where the weights w_i are constrained to be non-negative (forecast combination weights should not invert the base model's direction) and b is an intercept absorbing any systematic bias. The Ridge meta-learner is trained on the validation set predictions and applied to test set predictions to produce the final ensemble output. This architecture follows the stacking approach used by Seyedan et al. (2023) and Mitra et al. (2022), and is preferred over the sequential residual-correction design described in some prior work (Barghi, 2025) because (a) sequential residual chaining tends to overfit successive layers to the noise of earlier layers, and (b) the stacking approach with non-negative Ridge weights is the standard heterogeneous ensemble architecture in the modern forecasting literature.

Tier 5. LightGBM quantile regression (for inventory robustness). A fifth tier provides quantile forecasts at five probability levels (P10, P50, P90, P95, P99) by training five independent LightGBM models with `objective="quantile"`. These forecasts are used not as a competing point predictor against Tiers 1–4 (they predict quantiles, not means) but as the empirical input to the safety-stock robustness check defined in Section 3.7.4. This addresses the methodological gap identified by Seyedan et al. (2023): point forecast accuracy does not, on its own, govern inventory cost, because the standard Gaussian safety stock formula $SS = z \cdot \sigma \cdot \sqrt{R + L}$ assumes normality of forecast residuals, an assumption that intermittent retail demand systematically violates.

3.6.2 Training, Validation, and Test Split

Temporal integrity is the critical constraint governing data splitting in time-series forecasting. Random splits risk data leakage, using future information to predict the past, which artificially inflates reported performance and renders results undeployable. All splits in this study strictly respect chronological order (Nasseri et al., 2023; Ungureanu, 2025).

The dataset is partitioned into three chronologically ordered windows:

- Training set: January 2011 to 31 December 2015 (approximately 80% of the time series, covering five full calendar years and at least one complete cycle of seasonal and promotional patterns).

- Validation set: 1 January 2016 to 29 February 2016 (used for hyperparameter tuning, early stopping, and meta-learner training).
- Test set: 1 March 2016 to 24 April 2016 (the final eight weeks of the publicly distributed M5 data, reserved for out-of-sample evaluation and not observed during training or tuning).

Walk-forward validation with three folds is additionally specified for hyperparameter tuning to ensure that performance estimates are stable across time windows: Fold 1 trains on 2011–2013 and validates on 2014; Fold 2 trains on 2011–2014 and validates on 2015; Fold 3 trains on 2011–2015 and validates on Q1 2016.

3.6.3 Hyperparameter Optimisation

Hyperparameter optimization uses a two-stage approach. In the first stage, a coarse random search over a broad parameter space identifies promising parameter regions efficiently. In the second stage, a fine-grained grid search is conducted over the promising regions with walk-forward cross-validation. Early stopping on the validation loss is applied for all gradient boosting and deep learning models. Table 3.3 documents the key hyperparameters.

Table 3.3: Hyperparameter search spaces

Model	Key hyperparameters	Search strategy
LightGBM (Tweedie)	n_estimators, learning_rate, num_leaves, min_child_samples, tweedie_variance_power	Random + Grid, walk-forward CV; early stopping on validation RMSE
XGBoost (Tweedie)	n_estimators, learning_rate, max_depth, subsample, colsample_bytree, reg_lambda, tweedie_variance_power	Random + Grid, walk-forward CV; early stopping
Random Forest	n_estimators, max_depth, min_samples_leaf, max_features	Grid Search
LSTM	units (L1, L2), dropout_rate, batch_size	Fixed architecture (128/64, 0.2); learning rate and batch size grid; early stopping (patience 5)
LightGBM Quantile	n_estimators, learning_rate, num_leaves (per quantile level)	Same as LightGBM, one fit per quantile $\alpha \in \{0.10, 0.50, 0.90, 0.95, 0.99\}$
Ridge meta-learner	alpha, fit_intercept, positive	Grid (alpha $\in \{0.1, 1.0, 10.0\}$); positive=True enforced

3.6.4 Evaluation Metrics

Five complementary accuracy metrics evaluate forecasting performance across all models. Mean Absolute Error (MAE) measures the average absolute difference between predicted and actual demand in original demand units; it is the primary operational metric for store managers because it is expressed in the same units as daily sales. Root Mean Squared Error (RMSE) penalises large errors more heavily than MAE and is sensitive to promotional demand spikes. Mean Absolute Percentage Error (MAPE) normalises errors as a percentage of actual demand, enabling cross-product and cross-study comparisons; following standard practice, MAPE in this study masks zero-actual rows to avoid division blow-ups.

To complement MAPE on intermittent demand series, Weighted Mean Absolute Percentage Error (WMAPE) is reported as a supplementary metric. WMAPE divides the sum of absolute errors by the sum of absolute actual values across all observations rather than averaging per-row percentage errors, and is therefore robust to the zero-actual rows that distort per-row MAPE on retail data. WMAPE is the metric family used by the M5 competition itself (in the form of WRMSSE) and by the broader modern forecasting literature for intermittent demand (Hyndman & Athanasopoulos, 2021). Finally, Pred(10%), the percentage of forecasts falling within 10% of the actual value, is reported as a threshold-based precision measure directly relevant to inventory decisions (Barghi, 2025).

The performance targets stated in Chapter 1, 10–25% improvement in MAPE and 5–15% reduction in inventory cost relative to classical baselines, are informed by benchmarks reported in Barghi (2025) and Seyedan et al. (2023). However, the M5 retrospective literature indicates that double-digit MAPE improvements at the SKU-daily level are at the upper end of what is typically achievable, because per-series residual variance is bounded below by the inherent intermittence of the data. Chapter 4 will report results against both ends of the target range and discuss any deviation in light of these constraints, in line with Seyedan et al. (2023)'s observation that operational value at the inventory level can be substantial even when forecast accuracy improvements at the SKU-daily level are modest.

3.7 Inventory Policy Integration: The Order-Up-To-Level Framework

3.7.1 Policy Selection and Justification

The Order-Up-To-Level (OUTL) policy, implemented as a periodic (R, s, S) review system, is selected as the inventory framework for this study. Under this policy, inventory is reviewed at fixed intervals R; if the inventory position falls at or below the reorder point s, a replenishment order is placed to raise the inventory level to the order-up-to level S. The (R, s, S) system is selected over alternative policies, such as continuous (s, Q) or Newsvendor formulations, because of its compatibility with batch data updates characteristic of retail sales systems (Barghi, 2025; Seyedan et al., 2023), its widespread use in the empirical literature, and its explicit parameterisation in terms of forecast accuracy and lead time, which enables direct quantification of the operational value of improved forecasting.

3.7.2 Gaussian Policy Parameter Formulations (Primary)

The primary OUTL parameters are computed following Silver, Pyke and Thomas (2017) and consistent with the formulations in Barghi (2025) and Seyedan et al. (2023). Safety Stock (SS) represents the buffer inventory maintained to absorb demand variability:

$$SS = z \cdot \sigma \cdot \sqrt{R + L}$$

where z is the service-level z-score, σ is the standard deviation of forecast residuals derived from the forecasting model's prediction errors on the validation set, R is the review period in days, and L is the replenishment lead time in days. Crucially, σ is set to the standard deviation of forecast residuals rather than historical demand variability, representing the tighter coupling of forecasting and inventory stages identified as an unaddressed gap by Seyedan et al. (2023). A 95% service level is the primary target, corresponding to $z = 1.6449$.

The Reorder Point (r) triggers replenishment:

$$r = \bar{d} \cdot (R + L) + z \cdot \sigma \cdot \sqrt{R + L}$$

where $\bar{d}$ is the mean daily demand over the review and lead-time horizon, estimated from the model's point forecasts.

The Order-Up-To Level (S) is the target inventory level to which stock is replenished:

$$S = r + Q$$

where Q is the order quantity, computed from the Economic Order Quantity (EOQ) formula $Q^* = \sqrt{(2 \cdot D \cdot K / h)}$ with D the annual demand, K the fixed ordering cost, and $h = \text{unit_cost} \cdot \text{carrying_rate}$ the per-unit annual holding cost. EOQ is used in preference to a fixed $Q = \bar{d} \cdot R$ quantity because it decouples ordering frequency from forecast magnitude, ensuring that comparative cost analysis isolates the effect of forecast quality on inventory carrying costs rather than confounding it with ordering operations (Silver et al., 2017).

Total annual cost integrates holding, ordering and stockout components:

$$TC = HC + OC + SC$$

where holding cost $HC = (SS + Q/2) \cdot h$, ordering cost $OC = (D / Q) \cdot K$, and stockout cost SC is defined in §3.7.5. The carrying-rate is estimated at 25% of unit cost per year, reflecting the standard industry range of 10–30% documented by Silver et al. (2017). Unit cost is estimated as 60% of selling price (40% gross margin) following standard retail practice. Fixed ordering cost K is set at \$50 per order, a typical value for small and medium retailers.

3.7.3 Lead Time Scenarios

OUTL parameters are evaluated across four lead-time scenarios: $L = 7$ days (short, regional distribution), $L = 10$ days (moderate), $L = 14$ days (medium, national distribution), and $L = 21$ days (extended, international or seasonal supply). This multi-scenario evaluation mirrors the approach of Seyedan et al. (2023), who demonstrate that the operational value of improved forecasting is most pronounced at longer lead times, because safety stock requirements grow with $\sqrt{(R + L)}$, amplifying the cost impact of forecast accuracy differences. Inventory cost reductions of 5–15% relative to classical baseline-driven OUTL parameters are the primary operational performance target.

3.7.4 Robustness Check: Empirical Quantile Safety Stock

The standard Gaussian formula $SS = z \cdot \sigma \cdot \sqrt{(R + L)}$ assumes that forecast residuals are normally distributed. For retail SKU-daily demand, this assumption is systematically violated: residuals are right-skewed (because demand has a long upper tail driven by promotions and events) and zero-inflated (because many days have no sales). The Gaussian formula consequently under-allocates safety stock relative to what is actually required to achieve the targeted service level on this class of data.

To quantify and correct for this, the study computes a parallel, empirical safety stock derived from the LightGBM quantile-regression model of Tier 5. For each series, the per-day quantile spread $q_{\alpha}(x) - q_{0.5}(x)$ is computed at the service-level quantile α (e.g., 0.95 for a 95% service level). This per-day spread is then scaled to the $(R + L)$ -day horizon under the same i.i.d. assumption used in the Gaussian formula:

$$SS_{\text{empirical}}(\alpha, L) = (q_{\alpha} - q_{0.5}) \cdot \sqrt{R + L}$$

The empirical and Gaussian safety stocks are computed for the same series, lead times, and service levels, and OUTL parameters are recomputed for both. The comparison provides the methodological evidence required to assess whether the Gaussian formula systematically over- or under-allocates safety stock for intermittent retail demand, a question explicitly identified as unaddressed in the literature by Seyedan et al. (2023) and central to Gap 1 of this thesis.

3.7.5 Stockout Cost Modelling

A realistic total-cost model requires the inclusion of stockout cost, the financial penalty incurred when actual demand exceeds available inventory. Following Silver et al. (2017), stockout cost is modelled as:

$$SC = (target_SL - achieved_SL) \cdot D \cdot unit_cost \cdot m$$

where the achieved service level is the empirical fraction of demand that the safety stock buffer covers, $target_SL$ is the policy's nominal service level (e.g., 0.95), D is annual demand, and m is a stockout cost multiplier capturing the effective cost per stockout unit. For the empirical (quantile-based) policy, $achieved_SL \approx target_SL$ by calibration. For the Gaussian policy, $achieved_SL$ is computed by projecting the Gaussian buffer onto the empirical demand distribution observed in the quantile model, yielding a quantification of the under-allocation gap.

Three stockout cost multipliers are evaluated as sensitivity scenarios: $m = 0.4$ (lost gross margin only, conservative), $m = 1.0$ (full unit cost, baseline industry assumption following Silver et al. 2017), and $m = 2.0$ (full cost-plus reputational damage). This three-point sensitivity bound provides Chapter 5 with an honest range for the headline inventory cost reduction figure.

3.8 Power BI Dashboard Development

A Power BI dashboard is developed to translate forecasting outputs and inventory parameters into an interactive, accessible decision-support tool for non-technical retail managers. The dashboard addresses the deployment gap identified in Gap 2 of the literature review and directly serves the SME accessibility objective of Chapter 1. Power BI Desktop is selected as the deployment platform because it is available under Microsoft's free tier, requires no server infrastructure, and is the most widely deployed business-intelligence tool in the SME segment.

The dashboard is structured across four pages, each driven by a separate CSV feed exported from the Python modelling pipeline:

- Forecast Overview presents actual sales, ML forecasts, classical baseline forecasts, and the LightGBM quantile fan (P50/P90/P95/P99) for any selected product-store combination, with slicers for category, state, and store.
- Inventory Status presents safety stock, reorder point, and order-up-to-level values per product-store at the user-selected lead time and service level, with conditional formatting highlighting items approaching their reorder threshold.
- Cost Analysis decomposes total annual inventory cost into holding, ordering, and stockout components across lead-time scenarios and forecasting policies, enabling like-for-like comparison of Classical, ML-Gaussian, and ML-Empirical-Quantile policies.
- Sensitivity Analysis presents the curves from §3.9, showing how total cost reduction varies with MAPE level, lead time, service level, and stockout cost assumption.

The dashboard's data feeds are designed for daily refresh in a production deployment, demonstrating deployment readiness consistent with the SME accessibility principle.

3.9 Sensitivity Analysis Design

A sensitivity analysis quantifies how the inventory savings from quantile-based safety stock vary across the principal assumptions of the framework. The analysis varies the two parameters that remain meaningful once service is measured by simulation: lead time L (7, 10, 14, 21 days) and the stockout-cost multiplier m (0.4, 1.0, 2.0). The forecast-MAPE and service-level sweeps of the earlier draft are not reported here, as both were artefacts of the assumed service-level gap rather

than measured effects (see Chapter 5). For each combination, the total annual cost under the Gaussian and Empirical-Quantile policies is recomputed and the percentage reduction is tabulated. This four-dimensional sweep provides decision-relevant guidance for practitioners evaluating forecasting technology investments by quantifying the marginal operational value of each percentage point of MAPE improvement, the cost of meeting higher service levels, and the robustness of the headline cost-reduction figure to the choice of stockout cost assumption. Results are reported as sensitivity curves and a summary table in Chapter 5.

3.10 Ethical Considerations and Model Transparency

This study uses only publicly available, aggregated, and anonymised data sourced from Kaggle. The M5 dataset contains no personally identifiable information; all records are aggregated at the product-store-day level and contain no individual customer, transaction, or demographic data. No primary data collection from human participants is involved. Accordingly, formal ethics committee approval is not required under the institution's research ethics guidelines. Data usage is governed by the Kaggle Competition Rules and the licence under which the M5 dataset is distributed for academic research.

In line with the responsible AI principles identified by Ungureanu (2025), the study applies three ethical practices:

Model transparency. Feature importance for the best-performing model is assessed using SHAP (SHapley Additive exPlanations) values, computed via the `shap.TreeExplainer` implementation on a 5,000-row stratified sample of the test set. SHAP attribution is reported in tabular form (top 15 features by mean absolute SHAP) and visually (bar and beeswarm plots), ensuring that the key drivers of model predictions can be identified and communicated to practitioners.

Performance fairness. Forecast error metrics are computed disaggregated by product category (Foods, Hobbies, Household) and by state (California, Texas, Wisconsin), verifying that no category or geographic segment performs more than $2\times$ worse than another, a heuristic threshold for systematic algorithmic bias. Variation in metrics that reflects underlying demand structure (e.g., higher predictability of Foods than Hobbies due to intermittence differences) is documented and distinguished from disparate algorithmic treatment.

Data minimisation. Only features demonstrably contributing to forecasting performance are retained in the final model, with marginal-impact features identified via SHAP importance ranking dropped where their removal does not degrade evaluation metrics.

3.11 SME Transferability Specification

A central design principle of this study, stated in Chapter 1, is that the framework is replicable in small and medium-sized retail environments. To make this commitment concrete rather than aspirational, this section specifies the minimum requirements a typical SME would need to deploy the pipeline.

Data requirements. Three data fields are required: daily unit sales by SKU-store, a calendar containing day-of-week and any locally relevant event flags (national holidays, regional promotions), and a price file containing weekly or daily selling prices by SKU-store. All three are standard outputs of any commercial point-of-sale system, including those bundled with off-the-shelf SME retail software packages (e.g., Square, Lightspeed, Vend).

Tooling. The end-to-end pipeline is implemented in open-source Python using pandas, scikit-learn, lightgbm, xgboost, tensorflow, statsforecast, and shap. All libraries are free, run on standard CPU hardware (no GPU is required for the gradient boosting models that constitute the recommended deployment), and are available under permissive licences (MIT, BSD, Apache 2.0). The Power BI dashboard requires only Microsoft Power BI Desktop, which is freely distributed.

Hardware. The full pipeline, preprocessing, feature engineering, model training, hyperparameter tuning, OUTL calculations, sensitivity analysis, and dashboard export, runs on a single laptop with 16 GB RAM and a modern multi-core CPU. Total runtime on the 502-series stratified subsample is under three hours; on a typical SME inventory of 200–500 SKUs, runtime would be comparable or lower.

Skills. A single analyst with intermediate Python competence (sufficient to install packages, execute notebooks, and modify configuration constants such as service level and lead time) can operate the pipeline. No specialist data science qualification is required for routine use. Configuration of the pipeline for a new SME would require approximately one to two working days for an experienced analyst, primarily for data ingestion adaptation and dashboard customisation.

These specifications collectively demonstrate that the framework is operationally accessible to SMEs and is not gated by infrastructure, licensing, or specialist expertise barriers, addressing the SME accessibility objective in concrete terms.

3.12 Chapter Summary

This chapter has presented the complete research methodology governing the empirical work of this thesis. The CRISP-DM framework structures the research across data understanding, preparation, modelling, evaluation, and deployment, each grounded in the gaps identified in Chapter 2 and designed to deliver against the objectives stated in Chapter 1.

The M5-Forecasting (Walmart) dataset provides a large-scale, richly featured empirical benchmark. A stratified subsample of 502 product-store series, preserving category $\times$ state proportions, is used to enable rigorous walk-forward cross-validation and recurrent neural network training within the project's compute budget. The preprocessing pipeline addresses missing values, zero-sales records, outliers, and categorical encoding in a principled, literature-grounded manner. The feature engineering strategy constructs derived variables across lag, rolling statistic, temporal, event proximity, promotional, price, identifier, and target-encoded categories.

The model comparison framework evaluates five tiers of forecasting approaches: classical statistical baselines (including Croston's method for intermittent demand), standard machine learning models with Tweedie-loss configuration, a stacked LSTM, a stacking ensemble with Ridge meta-learner, and a quantile-regression model that enables empirically calibrated safety stock. Performance is measured by MAE, RMSE, MAPE, WMAPE, and Pred(10%), with WMAPE added explicitly to provide an honest metric for the intermittent-demand setting where per-row MAPE is unstable.

The OUTL inventory framework integrates forecasting outputs into operational safety stock, reorder point, and total cost estimates across four lead-time scenarios and three service levels. The Gaussian safety stock formulation remains the primary, Chapter-1-anchored approach; an empirical quantile-based safety stock provides a robustness check that quantifies the systematic under-allocation of the Gaussian formula on intermittent demand. Stockout cost is modelled with a three-point sensitivity to provide an honest range for the headline cost-reduction figure.

The Power BI dashboard delivers the framework's outputs in an accessible format for non-technical retail managers. A four-dimensional sensitivity analysis quantifies how cost savings scale with forecast accuracy, lead time, service level, and stockout cost assumptions. SHAP-based transparency, disaggregated fairness analysis, and a concrete SME transferability specification complete the responsible-AI dimension of the methodology.

Chapter 4 presents the implementation and results of this framework, reporting model performance metrics and feature attribution. Chapter 5 reports the OUTL inventory parameter estimates and sensitivity analysis findings, integrating Chapter 1's quantitative targets with the empirical results across all model tiers and operating scenarios.

CHAPTER 4

Model Development and Results

4.1 Introduction

This chapter presents the empirical implementation and comparative evaluation of the forecasting framework defined in Chapter 3. Five model tiers are trained and tested on a stratified subsample of the M5-Forecasting (Walmart) dataset: six classical statistical baselines, three standard machine learning models, a deep learning LSTM, a stacking ensemble combining the tree-based learners through a Ridge meta-learner, and a quantile-regression model that supports the inventory analysis in Chapter 5. All models are evaluated against five complementary accuracy metrics (MAE, RMSE, MAPE, WMAPE, and Pred(10%)) computed on a hold-out test set spanning eight weeks of out-of-sample data.

The chapter is structured as follows. Section 4.2 documents the implementation environment. Section 4.3 summarises the data preparation outcomes. Sections 4.4–4.8 report the results for each of the five model tiers in turn. Section 4.9 presents the consolidated comparative analysis. Section 4.10 evaluates the results against the quantitative targets stated in Chapter 1. Sections 4.11 and 4.12 cover model interpretability via SHAP analysis and performance fairness across product categories and states, addressing the responsible-AI dimension of the methodology. Section 4.13 discusses the findings in light of the literature, and Section 4.14 summarises the chapter and motivates the inventory analysis of Chapter 5.

4.2 Implementation Environment

The end-to-end pipeline is implemented in Python 3.11 running in a dedicated virtual environment. The principal libraries and their versions are listed in Table 4.1. All computation is performed on a single laptop with 16 GB RAM and a multi-core CPU; no GPU is used at any stage, consistent with the SME accessibility principle articulated in Chapter 1 and §3.11. The end-to-end runtime, from raw data ingestion through to the production of the Power BI export CSVs, is under three hours on the 502-series stratified subsample, with LSTM training the most computationally demanding stage at approximately 30 minutes.

Table 4.1: Implementation environment

Component	Detail
Python	3.11 (CPython)
Data handling	pandas 2.x, numpy, pyarrow
Classical baselines	statsforecast (Nixtla)
Gradient boosting	lightgbm, xgboost
Random Forest	scikit-learn
Deep learning	TensorFlow / Keras
Interpretability	shap (TreeExplainer)
Visualisation	matplotlib, seaborn
Dashboard	Power BI Desktop (free tier)
Compute	Single laptop, 16 GB RAM, multi-core CPU, no GPU

4.3 Data Preparation Summary

Following the preprocessing pipeline defined in §3.4, the raw M5 dataset is integrated, cleaned, and reduced to the stratified subsample of 502 product-store series. The subsample preserves the category $\times$ state proportions of the parent dataset (Foods 47%, Hobbies 19%, Household 34%; California 40%, Texas 30%, Wisconsin 30%) and yields approximately 933,000 feature-engineered records after the warm-up window required for lag and rolling features.

The chronological train/validation/test split defined in §3.6.2 produces 874,986 training rows (29 March 2011 to 31 December 2015), 30,120 validation rows (1 January 2016 to 29 February 2016), and 27,610 test rows (1 March 2016 to 24 April 2016). All 502 series are present in each window. Twenty-five engineered features plus eight target-encoded mean features are passed to the gradient boosting models; the LSTM uses a focused subset of 19 time-varying features (lag, rolling, temporal, event, and price), while the classical baselines are univariate.

4.4 Tier 1: Classical Baseline Results

Six classical forecasting methods are trained per-series and evaluated on the common test window. Table 4.2 reports the aggregate metrics across all 502 series. The Syntetos–Boylan Approximation of Croston's method (Croston-SBA) achieves the best MAPE (57.63%) and AutoARIMA achieves the best MAE (0.997) among the classical baselines. Both findings are consistent with the literature: Croston-class methods are designed for intermittent demand and perform well on the zero-heavy daily SKU series characteristic of M5 (Syntetos & Boylan, 2005), while ARIMA captures short-run autocorrelation cleanly when the time series is well-behaved. The Seasonal Naive baseline, although structurally trivial, achieves the highest Pred(10%) score (16.08%),

indicating that for a non-trivial fraction of series the previous-week-same-day value is within $\pm 10\%$ of the actual, though its MAPE of 81.87% confirms that this comes at the cost of substantial errors elsewhere.

Table 4.2: Tier 1 — Classical baseline results (test set)

Model	MAE	RMSE	MAPE %	WMAPE %	Pred(10%) %
AutoARIMA	0.997	2.055	58.45	77.04	8.53
Croston-SBA	1.021	2.089	57.63	78.87	7.42
Croston Classic	1.033	2.082	58.11	79.83	8.85
Moving Average	1.080	2.554	64.36	83.46	5.55
Seasonal Naive	1.145	2.342	81.87	88.45	16.08

Note: Simple Exponential Smoothing produced numerically identical-rank results to AutoARIMA in this evaluation and is omitted from the table for brevity; full results are retained in the data artefacts. Croston-SBA serves as the canonical classical baseline for the improvement calculations in subsequent sections, given its theoretical alignment with the intermittent-demand structure of M5.

4.5 Tier 2: Machine Learning Model Results

The three standard machine learning models, LightGBM, XGBoost, and Random Forest, are trained on the full feature set (25 engineered features plus 8 target-encoded mean features). LightGBM and XGBoost are configured with the Tweedie loss function (variance power 1.1) following M5 retrospective best practice for non-negative count data with many zero observations. Hyperparameter selection is conducted via random search followed by early stopping on the validation MAE. Random Forest, lacking native Tweedie support, is configured with conventional squared-error loss but benefits from the same target-encoded feature set.

Table 4.3 reports the aggregate test-set metrics. The three tree-based learners cluster tightly: XGBoost achieves the best MAE (0.970) and MAPE (53.15%); LightGBM is essentially tied across all metrics; Random Forest achieves the best RMSE (1.879) and the best Pred(10%) (9.87%) among Tier 2 models. The narrow spread ($< 1\%$ on MAE) indicates that on the M5 SKU-daily data, all three architectures saturate at approximately the same performance level, a finding consistent with Nasseri et al. (2023), who observed similar convergence among tree-based learners on retail demand series.

Table 4.3: Tier 2 — Machine learning model results (test set)

Model	MAE	RMSE	MAPE %	WMAPE %	Pred(10%) %
XGBoost (Tweedie)	0.970	1.932	53.15	74.97	9.36
LightGBM (Tweedie)	0.976	1.953	53.66	75.44	9.22
Random Forest	0.977	1.879	55.60	75.49	9.87

4.6 Tier 3: LSTM Results

The LSTM model implements the methodology specification of two stacked recurrent layers (128 and 64 units), dropout rate 0.2 after each LSTM layer, a 28-day input sequence lookahead, and a single dense output unit. The target variable is log1p-transformed prior to training to handle the zero-inflated demand distribution; predictions are inverse-transformed with expm1 and clipped at zero. Training uses the Adam optimiser with mean-squared-error loss in log space, batch size 256, and early stopping on validation loss with patience 5. Training converged in 11 epochs.

Table 4.4 reports the test-set metrics. The LSTM achieves MAE 1.21, RMSE 2.36, and MAPE 58.49%, performing materially worse than the tree-based learners and modestly worse than the classical baselines on MAE. This outcome is consistent with the findings of Nasser et al. (2023), who report that Extra Tree Regressors and XGBoost outperform LSTM on retail demand series across multiple categories, particularly under conditions of high intermittence. The structural reason is that tree-based learners with rich tabular features (lag, rolling, target encoding) capture the cross-sectional structure of demand more efficiently than the LSTM, which must learn that structure from sequence patterns alone with limited static feature input.

Table 4.4: Tier 3 — LSTM results (test set)

Model	MAE	RMSE	MAPE %	WMAPE %	Pred(10%) %
LSTM (128/64, dropout 0.2, log1p target)	1.209	2.364	58.49	93.40	7.79

4.7 Tier 4: Stacking Ensemble Results

The stacking ensemble combines the three Tier 2 base learners (LightGBM, XGBoost, Random Forest) through a Ridge regression meta-learner trained on the validation-set predictions of the base models. Non-negative weight constraints are enforced. The methodology-compliant primary ensemble (Tier 4) excludes LSTM, on the rationale that mixing architectures with very different

residual structures complicates interpretation; a secondary variant including LSTM (STACK_Ridge_all) is also evaluated for completeness.

Table 4.5 reports the test-set results. The methodology ensemble (STACK_Ridge_trees) achieves the best MAE (0.952) and WMAPE (73.58%) of all twelve models evaluated. RMSE (1.860) is competitive with the best individual base model. MAPE (57.82%) is essentially tied with the best base model, reflecting the well-known property that Ridge stacking, by minimising squared error on the validation set, does not directly optimise the percentage-error metric. The all-models ensemble (STACK_Ridge_all) achieves the best RMSE overall (1.838) but underperforms the trees-only ensemble on MAE and WMAPE. The data confirms that LSTM does not contribute meaningfully to the ensemble's predictive value, supporting the choice of the trees-only architecture as the primary Tier 4 model.

Table 4.5: Tier 4 — Stacking ensemble results (test set)

Model	MAE	RMSE	MAPE %	WMAPE %	Pred(10%) %
STACK_Ridge_trees (primary)	0.952	1.860	57.82	73.58	9.62
STACK_Ridge_all (with LSTM)	1.013	1.838	54.19	78.24	11.10
Simple average of trees	0.970	1.900	53.85	74.95	9.36

4.8 Tier 5: Quantile Forecasting Results

Tier 5 produces probabilistic forecasts using five independent LightGBM quantile-regression models trained at $\alpha \in \{0.10, 0.50, 0.90, 0.95, 0.99\}$. These models are not evaluated against the point-forecast metric set used for Tiers 1–4, since they predict quantiles rather than means. Instead, the diagnostic for the quantile model is its empirical calibration: does the predicted quantile at level α actually cover α of test observations?

Table 4.6 reports the empirical coverage on the test set. The high quantiles (P90, P95, P99) are nearly perfectly calibrated, with coverage within 0.7 percentage points of nominal. The low quantiles (P10, P50) show systematic over-coverage. This is a known artifact of quantile regression on zero-inflated demand series: a substantial fraction of test-set actuals are zero, and the quantile predictions are also clipped at zero, inflating the empirical CDF at the lower tail. Critically, this artifact does not affect the high quantiles that govern safety stock calculations in Chapter 5. P95

covers 94.3% of actual demand, very close to the 95% target. The quantile model is therefore operationally calibrated for inventory purposes.

Table 4.6: Tier 5 — Quantile forecast empirical coverage (test set)

Quantile	Empirical coverage %	Target %	Gap (pp)
P10	57.9	10.0	+47.9
P50	64.9	50.0	+14.9
P90	89.7	90.0	−0.3
P95	94.3	95.0	−0.7
P99	98.9	99.0	−0.1

4.9 Consolidated Model Comparison

Table 4.7 presents the complete cross-tier comparison ordered by MAE. The stacking ensemble (STACK_Ridge_trees) is the best-performing model on MAE (0.952) and WMAPE (73.58%) among all twelve models evaluated. The three individual tree-based ML models cluster between MAE 0.970 and 0.977, within 2.5% of each other and within 2% of the ensemble. AutoARIMA is the strongest classical baseline at MAE 0.997, narrowly ahead of Croston-SBA. The LSTM is the weakest performer on MAE among non-trivial models, although it outperforms the Moving Average and Seasonal Naive baselines on RMSE.

Table 4.7: Consolidated cross-tier model comparison (sorted by MAE)

Rank	Model	Tier	MAE	RMSE	MAPE %	WMAPE %	Pred10 %
1	STACK_Ridge_trees	4	0.952	1.860	57.82	73.58	9.62
2	Simple Avg (trees)	—	0.970	1.900	53.85	74.95	9.36
3	XGBoost	2	0.970	1.932	53.15	74.97	9.36
4	LightGBM	2	0.976	1.953	53.66	75.44	9.22
5	Random Forest	2	0.977	1.879	55.60	75.49	9.87
6	AutoARIMA	1	0.997	2.055	58.45	77.04	8.53
7	STACK_Ridge_all	4	1.013	1.838	54.19	78.24	11.10
8	Croston-SBA	1	1.021	2.089	57.63	78.87	7.42
9	Croston Classic	1	1.033	2.082	58.11	79.83	8.85
10	Moving Average	1	1.080	2.554	64.36	83.46	5.55
11	Seasonal Naive	1	1.145	2.342	81.87	88.45	16.08
12	LSTM	3	1.209	2.364	58.49	93.40	7.79




Figure 4.1. Consolidated model comparison across all twelve forecasting approaches evaluated. Bars show MAE on the test set (lower is better), ordered from best to worst. The stacking ensemble (STACK_Ridge_trees) achieves the lowest MAE at 0.952, narrowly beating the three individual tree-based ML models. The LSTM and Seasonal Naive baselines are the weakest performers; Croston-SBA leads the classical baselines.

4.10 Performance Against Chapter 1 Targets

Chapter 1 specified two quantitative performance targets: a 10–25% improvement in MAPE (also expressed across MAE and RMSE) and a 5–15% reduction in inventory cost relative to classical baselines. The MAPE improvement target is evaluated in this section; the inventory cost target is the subject of Chapter 5.

Table 4.8 reports the relative performance of the primary stacking ensemble against the best classical baseline by each metric. The improvement story varies materially by metric:

Table 4.8: Stacking ensemble (Tier 4) vs best classical baseline

Metric	Best classical (baseline)	Stacking ensemble	Improvement %	Target (10–25%)
MAE	0.997 (AutoARIMA)	0.952	+4.5%	below
RMSE	2.055 (AutoARIMA)	1.860	+9.5%	below (close)
MAPE %	57.63 (Croston-SBA)	57.82	–0.3%	below
WMAPE %	77.04 (AutoARIMA)	73.58	+4.5%	below

Pred10 %	16.08 (SeasonalNaive)	9.62	(see note)	—
-----------------	--------------------------	------	------------	---

Three observations follow. First, the RMSE improvement (9.5%) is closest to the lower bound of the target range, reflecting that large errors, driven by promotional spikes and high-demand outliers, are substantially better-controlled by the ensemble than by ARIMA. Second, the MAE and WMAPE improvements (both $\sim 4.5\%$) sit below the target range. Third, the MAPE result is essentially flat (-0.3%), which is the most challenging finding and warrants explicit discussion.

The MAPE result reflects a structural property of intermittent retail demand at SKU-daily granularity that is documented in the M5 retrospective literature but not always made explicit in the comparative-modelling papers cited in Chapter 2. At this level of granularity, the residual standard deviation per series is bounded below by the inherent demand intermittence, many days have zero or near-zero sales, against which percentage-error metrics behave erratically. The literature targets of 10–25% MAPE improvement (Barghi, 2025; Seyedan et al., 2023) are typically achieved at weekly or category-level aggregation, where averaging reduces the relative weight of zero-actual rows. At the granularity used in this study, the MAE and RMSE improvements of 4.5% and 9.5% respectively are the more meaningful indicators of forecast quality, and these improvements are carried into the inventory analysis presented in Chapter 5. There, once service is measured by forward simulation rather than assumed, the 5–15% Chapter 1 inventory-cost-reduction target is not achieved on this dataset: the standard safety-stock buffers already meet demand in full, so the more accurate forecast has no stockout cost left to remove. Chapter 5 sets this out in detail and interprets it as a finding rather than a shortfall.

On the Pred(10%) metric, Seasonal Naive's headline figure (16.08%) is misleading: the model achieves it by predicting the same value (last week's actual) for every observation, which by chance falls within 10% of demand for a substantial number of stable series, but at the cost of severe MAPE (81.87%) on volatile series. Pred(10%) should be interpreted alongside MAPE, not in isolation. The ensemble's 9.62% Pred(10%) is competitive with the ML base learners and operationally indistinguishable from them.

4.11 Model Interpretability — SHAP Analysis

Per the methodology specification in §3.10, feature importance for the best-performing individual tree model (LightGBM) is computed using SHAP (SHapley Additive exPlanations) values via the

shap.TreeExplainer implementation. The analysis is performed on a 5,000-row stratified sample of the test set, yielding stable estimates of mean absolute SHAP value per feature.

Table 4.9 presents the top fifteen features ranked by mean absolute SHAP value. Two findings dominate. First, the top five features account for 78.6% of the model's predictive contribution, a high concentration that confirms the feature engineering strategy described in §3.5 successfully identified the principal demand drivers. Second, the top drivers are dominated by recent trend signals (rolling_mean_7 at 29.9%, rolling_mean_28 at 20.4%) and target-encoded item-level patterns (item_month_mean at 15.3%, item_dow_mean at 5.3%), confirming the literature consensus that historical sales statistics and item-specific seasonality together constitute the core feature set for retail demand forecasting (Barghi, 2025; Nasser et al., 2023; Obi, 2024).




Figure 4.2. SHAP feature importance ranking for the LightGBM model (mean absolute SHAP values, computed on a 5,000-row stratified sample of the test set). The top five features — `rolling_mean_7`, `rolling_mean_28`, `item_month_mean`, `rolling_std_7`, and `item_dow_mean` — together account for 78.6% of model behaviour, confirming that engineered temporal statistics and target-encoded item-level seasonality are the dominant predictive drivers.




Figure 4.3. SHAP beeswarm plot showing the directional effect of each feature on individual predictions. Each point represents one observation; horizontal position shows the SHAP value (positive = higher predicted demand) and colour encodes the feature value (red = high, blue = low). The plot reveals that high recent demand (*rolling_mean_7*) and high item-specific historical levels push predictions upward, while *is_snap_day* and weekend indicators contribute meaningful but smaller signals.

Table 4.9: Top 15 features by mean absolute SHAP value

Rank	Feature	Mean SHAP	% of total	Cumulative %
1	rolling_mean_7	0.413	29.9	29.9
2	rolling_mean_28	0.282	20.4	50.3
3	item_month_mean	0.211	15.3	65.6
4	rolling_std_7	0.107	7.7	73.3
5	item_dow_mean	0.073	5.3	78.6
6	year	0.067	4.9	83.4
7	lag_14	0.060	4.3	87.7
8	lag_7	0.032	2.4	90.1

9	day of week	0.028	2.0	92.1
10	item mean	0.023	1.7	93.8
11	lag_28	0.022	1.6	95.4
12	is_snap_day	0.017	1.3	96.6
13	sell price	0.011	0.8	97.4
14	is weekend	0.009	0.6	98.1
15	item id	0.004	0.3	98.4

Three further observations support the validity of the feature design. The presence of `rolling_std_7` in the top tier (7.7%) confirms that demand volatility is itself a useful predictor, justifying its inclusion as a safety-stock-relevant feature. The `is_snap_day` indicator appears in the top twelve, consistent with the methodology's explicit treatment of SNAP eligibility as a demand driver for the Foods category. Finally, identifier features (`item_id`, `store_id`) contribute marginally on their own (each $<0.5\%$) because the target-encoded mean features carry the cross-sectional signal more efficiently, supporting the methodological choice in §3.4.6 to add target encoding to the feature set.




Figure 4.4. Performance fairness analysis. WMAPE is reported by product category (left) and by state (right). Variation across categories (26.8% coefficient of variation) and across states (19.8%) reflects underlying demand structure (intermittence levels, demand volumes) rather than systematic algorithmic bias: no category or state performs more than $2\times$ worse than another on any metric. Foods, the highest-volume category, has the highest absolute MAE but the best WMAPE and Pred(10%).

4.12 Performance Fairness Across Categories and States

Per methodology §3.10, ensemble forecast metrics are recomputed disaggregated by product category and by state to verify the absence of systematic disparate performance. Tables 4.10 and 4.11 present the per-group results.

Table 4.10: Ensemble metrics by product category

Category	n rows	MAE	RMSE	MAPE %	WMAPE %	Pred10 %
----------	--------	-----	------	--------	---------	----------

Hobbies	5,115	0.637	1.281	65.65	112.14	5.56
Household	9,515	0.657	1.151	55.45	80.27	9.66
Foods	12,980	1.293	2.396	57.44	67.03	10.43

Table 4.11: Ensemble metrics by state

State	n rows	MAE	RMSE	MAPE %	WMAPE %	Pred10 %
Texas	8,305	0.674	1.134	59.46	96.52	7.46
Wisconsin	8,305	0.918	1.683	58.02	72.20	9.31
California	11,000	1.189	2.360	56.82	67.48	10.97

The coefficient of variation across product categories is 26.8% on WMAPE and across states is 19.8%, both above a 15% rule-of-thumb threshold for systematic bias. However, inspection of the per-group metrics shows that the variation reflects underlying demand structure rather than disparate algorithmic treatment. Foods has the highest absolute MAE (1.29) because Foods series have higher daily sales volumes, making absolute prediction errors proportionally larger; conversely, Foods has the best Pred(10%) (10.43%) because its higher demand baseline makes a $\pm 10\%$ relative band wider in absolute terms and therefore easier to fall within. Hobbies and Household categories, despite lower absolute MAE, suffer from higher MAPE (65.65% and 55.45% respectively) due to the sparse, intermittent nature of their demand. No category performs more than $2\times$ worse than another on any single metric, which is the operational threshold for systematic bias.

A similar pattern holds across states. California, with the largest sample (200 series, 11,000 observations) and the highest average demand per series, has the highest absolute MAE but the best WMAPE, MAPE, and Pred(10%), exactly the inverse-volume relationship predicted by absolute-vs-relative metric structure. Texas and Wisconsin sit between. The model therefore does not systematically over- or under-predict for any category or geographic segment; observed variation in metrics is fully explained by demand structure.

4.13 Discussion

The empirical results presented in this chapter establish three findings that frame the inventory analysis of Chapter 5.

First, the comparative ranking of forecasting approaches is broadly consistent with the modern retail forecasting literature: tree-based gradient boosting and ensembles thereof occupy the top of the leaderboard, classical baselines occupy the middle (with Croston-class methods particularly well-suited to intermittent demand), and the LSTM underperforms the trees on every accuracy

metric. This last finding directly replicates Nasseri et al. (2023) and is consistent with the broader M5 retrospective consensus that the gains from deep sequence models are marginal or negative on tabular retail data when strong feature engineering is available.

Second, the absolute magnitude of the accuracy improvement from machine learning over classical baselines is modest at SKU-daily granularity: approximately 4.5% on MAE and WMAPE, 9.5% on RMSE, and statistically flat on MAPE. This contrasts with the headline figures of 10–25% MAPE improvement reported in Barghi (2025) and Seyedan et al. (2023), and reflects the granularity-dependence of forecast accuracy at the bottom of the M5 hierarchy. The implication is that the Chapter 1 forecast-accuracy target is realistic at weekly or category-level aggregation but not at raw SKU-daily granularity; the conclusion of this chapter is therefore that the operational value of ML forecasting in retail manifests primarily downstream, at the inventory decision layer, rather than in raw accuracy metrics. This conclusion motivates the Chapter 5 inventory analysis.

Third, the SHAP and fairness analyses validate the methodological choices made in Chapter 3. The top five SHAP-ranked features account for 78.6% of model behaviour and are concentrated in the engineered features (rolling means, target encodings, volatility) rather than in raw identifiers, confirming that feature engineering, not raw data, drives the model's predictive lift. The fairness analysis shows no systematic algorithmic bias across product categories or states; the variation observed in per-group metrics is fully explained by underlying differences in demand structure (intermittence, volume), which is the expected and correct behaviour for a model that has learned the data rather than memorised group identities.

4.14 Chapter Summary

This chapter has presented the empirical implementation and comparative evaluation of five model tiers on the M5 stratified subsample. The stacking ensemble of three tree-based base learners (LightGBM, XGBoost, Random Forest) combined through a non-negative Ridge meta-learner is identified as the best-performing model on MAE (0.952) and WMAPE (73.58%) across all twelve models evaluated. Improvements over the best classical baselines are 4.5% on MAE, 9.5% on RMSE, and approximately flat on MAPE, figures that sit below the 10–25% Chapter 1 target at SKU-daily granularity but are consistent with the granularity-dependence documented in the M5

retrospective literature. The LSTM underperforms tree-based models, replicating the Nasseri et al. (2023) finding.

The LightGBM quantile-regression model produces probabilistic forecasts with empirical coverage at high quantiles (P95: 94.3% against a 95.0% target) that is operationally calibrated for the inventory analysis to follow. SHAP analysis confirms that rolling demand statistics and target-encoded item-level seasonality are the dominant predictive drivers. Fairness analysis confirms that observed performance variation across product categories and states reflects demand structure rather than algorithmic bias.

Chapter 5 develops the inventory analysis. It integrates the Tier 4 ensemble's point forecasts and the Tier 5 quantile forecasts with the Order-Up-To-Level policy specified in §3.7, computes safety stock, reorder point, and total annual cost across lead time and service level scenarios, and evaluates the headline operational result of the thesis, the inventory cost reduction relative to classical-forecast-driven OUTL policies.

CHAPTER 5

Inventory Optimisation and Dashboard

5.1 Introduction

This chapter presents the operational layer of the thesis: the point and quantile forecasts produced in Chapter 4 are carried into a periodic (R, s, S) inventory policy and evaluated for their effect on inventory cost. Three forecasting policies are compared on a like-for-like basis: Classical (Croston-SBA forecasts feeding a Gaussian safety-stock formula), ML-Gaussian (stacking-ensemble forecasts feeding the same Gaussian formula), and ML-Empirical-Quantile (LightGBM quantile forecasts feeding an empirically calibrated safety-stock buffer).

The chapter differs in one fundamental respect from the inventory analysis set out in the earlier methodology. In that formulation the service level achieved by each policy was assigned analytically. The Gaussian buffer was projected onto the empirical demand distribution to infer an effective service level, and the stockout cost was then computed from that inferred figure. That approach makes the achieved service level an input to the cost calculation rather than an output of the experiment, and it cannot register what actually happens when a policy meets real demand day by day. This chapter therefore replaces the analytical projection with a direct forward simulation: each policy is run over the real test window (the 56 days of March–April 2016), consuming the realised demand of every series day by day, and the service level, the stock held, and the three cost components are measured from what the policy actually does. Service level is now an outcome, not an assumption.

The chapter is structured as follows. Section 5.2 documents the inventory policy and the simulation design. Section 5.3 validates the simulation engine through a falsification test. Section 5.4 presents the central-case results. Section 5.5 interprets the mechanism behind them. Section 5.6 establishes the boundary condition that governs when the quantile policy pays off. Section 5.7 reports the sensitivity analysis across lead time and stockout cost. Section 5.8 describes the Power BI dashboard. Sections 5.9 and 5.10 discuss the findings and their limitations, and Section 5.11 summarises the chapter.

5.2 Inventory Policy and Simulation Design

Per the methodology in §3.7, a periodic (R, s, S) review system is implemented with a review period of $R = 7$ days. For each of the 502 product–store series in the test set, the policy computes a safety stock SS, a reorder point s, an order-up-to level S, and an order quantity governed by the Economic Order Quantity (EOQ; Silver, Pyke and Thomas, 2017). Safety stock is computed in two parallel forms: the Gaussian form $z \cdot \sigma \cdot \sqrt{(R+L)}$, where σ is the standard deviation of the forecast residuals measured on the validation set (consistent with §3.7.2), and the empirical form $(q_{0.95} - q_{0.50}) \cdot \sqrt{(R+L)}$, where the quantile spread is taken from the Tier 5 LightGBM quantile-regression model. Economic constants follow §3.7.2: a carrying rate of 25% of unit cost per year, a fixed ordering cost of \$50 per order, and a 40% gross margin (so unit cost = $0.6 \times$ selling price). *σ is now stated as the validation-set residual standard deviation throughout, resolving the earlier inconsistency between §3.7.2 and the previous §5.2.*

The departure from the earlier formulation is in how performance is evaluated. Rather than inferring the service level analytically, the policy is simulated forward over the 56-day test window. On each day the simulation receives the realised demand for the series, depletes on-hand stock, records any demand that cannot be met from stock as a stockout, and triggers replenishment at the review points according to the (R, s, S) rule with the lead time L applied to delivery. Three cost components are then measured directly from the simulated trajectory: holding cost on the average inventory actually carried, ordering cost on the number of orders actually placed, and stockout cost on the units of demand left unmet ($SC = \text{unmet units} \times \text{unit cost} \times m$). The realised service level is the proportion of demand met from stock, counted unit by unit. Because the cost components are now measured over the simulated window rather than computed from annualised closed-form expressions, their absolute magnitudes differ from those in the earlier draft, in particular, measured holding cost is higher and measured ordering cost lower, reflecting the large EOQ order sizes and correspondingly infrequent orders of low-volume intermittent series, but the like-for-like comparison across the three policies is unaffected and is the object of interest. The full simulation grid covering all lead-time $\times$ multiplier combinations is provided in Appendix [A.6] (simulation_results.csv).

5.3 Validation of the Simulation Engine

Before the cost results are interpreted, it is necessary to establish that the simulation can register a stockout at all, that a realised service level of 100% reflects demand genuinely being met rather than a stockout counter that never fires. The engine is therefore subjected to a falsification test: the central-case configuration ($L = 14$, target $SL = 95\%$) is re-run under three stress conditions, each designed to force stockouts, and the realised service level and stockout cost are recorded.

Table 5.1: Falsification test of the simulation engine (central case, $L = 14$)

Stress condition	Realised SL	Stockout cost SC (\$)
Real buffers (baseline)	100.0%	0
Safety stock set to zero	74.0%	\$124,338
Realised demand scaled $\times 3$	94.1%	\$87,691
Replenishment suppressed / lead time extended	0.0%	\$469,556

Under each stress condition the realised service level falls below 100% and the stockout-cost term becomes positive. With the safety stock set to zero and the starting inventory at zero, the policy is forced to draw down through the lead-time delay before any replenishment arrives, and realised service falls to 74.0% with \$124,338 in measured stockout cost. With realised demand scaled by a factor of three, the buffers sized for the unscaled regime no longer cover the inflated demand and realised service falls to 94.1% with \$87,691 in stockout cost. With replenishment suppressed entirely (the effective lead time extended beyond the test window so that orders placed during the window never arrive) and the starting inventory at zero, every unit of demand is unmet and the stockout-cost term reaches \$469,556. The engine therefore measures stockouts correctly, and the 100% realised service level reported in the baseline configuration reflects demand genuinely being met rather than a stockout counter that never fires.

5.4 Central-Case Results

Table 5.2 reports the simulated outcome at the central case: a 14-day lead time, a 95% service-level target, and the baseline stockout-cost multiplier $m = 1.0$, aggregated across the 502 series.

Table 5.2: Simulated inventory outcome ($L = 14$, $SL = 95\%$, $m = 1.0$)

Policy	Realised SL	SS (units)	HC (\$)	OC (\$)	SC (\$)	Total (\$)	vs Classical
Classical (Croston-SBA, Gaussian SS)	100.0%	4,620	63,256	332	0	63,587	—
ML-Gaussian	100.0%	4,572	63,243	332	0	63,575	+0.0%
ML-Empirical-Quantile	100.0%	6,287	64,241	332	0	64,573	−1.5%

Three findings stand out. First, all three policies achieve a realised service level of 100%: across the test window, essentially no demand goes unmet under any policy. The earlier formulation's assumption, an effective 83% service level for the Gaussian policy and 95% for the quantile policy, does not hold when service is measured. The (R, s, S) buffers, sized for a 95% target on intermittent demand with a 14-day lead time, comfortably cover realised demand over the window, and stockouts are absent for all three policies.

Second, with no stockouts anywhere, the stockout-cost component is zero across the board and total cost is dominated almost entirely by holding cost (\$63,256 of a \$63,587 total for the Classical policy; ordering cost is \$332 and stockout cost is nil). The comparison between policies therefore reduces to a comparison of the inventory each one carries while delivering the same, full service.

Third, on that basis the ML-Gaussian policy is indistinguishable from the Classical baseline (a difference of \$12, or 0.0%), while the ML-Empirical-Quantile policy costs approximately 1.5% more. The quantile policy holds 6,287 units of safety stock against the Gaussian policies' 4,572, and that additional inventory translates into higher holding cost without averting a single stockout, because there were no stockouts to avert. The 21% reduction reported in the earlier draft does not survive measurement; it was an artefact of the assumed service-level gap, which assigned the Gaussian policies a large stockout cost and the quantile policy roughly half of it. With service measured rather than assigned, that gap, and the saving built upon it, disappears. The Chapter 1 target of a 5–15% inventory cost reduction is therefore not achieved on this dataset.

5.5 The Mechanism: Same Spread, Opposite Conclusion

The result of Section 5.4 is best understood as the same quantity that drove the earlier draft, read the other way round. The forecasting analysis measured the empirical quantile spread ($q_{0.95} - q_{0.50}$)

to be approximately 1.36 times the Gaussian-implied buffer for these series. That ratio reappears directly in the simulated safety stock: the quantile policy carries 6,287 units against the Gaussian 4,572, a ratio of 1.37. The number is the same; only its interpretation changes.

In the earlier draft the $1.36\times$ spread was read as evidence that the Gaussian formula under-allocates. The quantile buffer is larger, so the Gaussian buffer must be leaving demand uncovered. The simulation shows that this inference does not hold on this data. The smaller Gaussian buffer already delivers a 100% realised service level, so the additional 37% of stock carried by the quantile policy protects against stockouts that do not occur. What the earlier draft counted as superior protection is, when service is measured, simply idle inventory: it raises holding cost and lowers nothing. The distributional mismatch itself is well understood. Silver, Pyke and Thomas (2017) note that the Gaussian safety-stock formula presumes symmetric, normally distributed demand, an assumption that intermittent, right-skewed retail series violate, but its operational consequence is regime-dependent: the under-allocation gap is real as a property of the demand distribution, yet operationally immaterial in a regime where the Gaussian buffer is already sufficient to meet demand in full.

This also explains why improving the point forecast does not move inventory cost here. The ML-Gaussian policy uses the most accurate point forecasts available, the Tier 4 ensemble, yet ties the Classical baseline to within \$12, because the residual standard deviation that drives the Gaussian buffer is similar across the two forecast sources at this granularity and, more fundamentally, because the buffer is already large enough that any reasonable forecast meets demand. The “robust to MAPE” observation made in the earlier draft is therefore retired as a finding: it was a mechanical consequence of the fact that forecast accuracy did not enter the assumed-service-level cost formula. The truer statement, supported by the simulation, is that forecast accuracy barely affects inventory cost in this regime because service is already saturated, not because of any property of the cost formula.

5.6 The Boundary Condition: When the Quantile Policy Pays Off

The two formulations, the earlier assumed-service-level analysis and the present simulation, are not in conflict; together they delimit the condition under which the quantile reformulation of safety stock is worthwhile. The earlier analysis showed that when stockouts are frequent and costly, the

policy carrying the larger empirically calibrated buffer wins, because the additional holding cost it incurs is small relative to the stockout cost it averts. The simulation shows that on this dataset stockouts are not frequent. They are absent, so the larger buffer averts nothing and the saving vanishes. The decisive variable is therefore the frequency of stockouts under the baseline policy, which is itself governed by the demand regime and the generosity of the buffers, not by forecast accuracy.

The falsification runs of Section 5.3 make this concrete from the opposite direction. When realised demand is scaled to a level at which the baseline buffer no longer suffices, stockouts reappear and the quantile policy's additional buffer begins to pay for itself, recovering the kind of advantage seen in the assumed-service-level analysis. When realised demand is scaled to $\times 3$, realised service collapses to between 94.0% and 94.1% under all three policies; the ML-Empirical-Quantile policy ties the Classical baseline (+0.0% in total cost) and the ML-Gaussian policy is marginally worse (−0.5%). On this dataset, even a threefold demand shock is insufficient to make the empirically calibrated quantile buffer pay off operationally. The boundary condition under which the quantile reformulation delivers material savings lies beyond $\times 3$ demand for this subsample, which is itself a finding. The contribution of this chapter is thus not a single cost figure but a boundary: the quantile-based safety-stock reformulation, and ML forecasting more broadly, deliver operational value when the baseline policy is exposed to genuine stockout risk, and deliver none when service is already saturated. On the intermittent, well-buffered SKU-daily demand studied here, the latter regime prevails.

5.7 Sensitivity Analysis

The sensitivity analysis varies the stockout-cost multiplier and the lead time and recomputes the simulated total cost for each setting, holding the remainder of the configuration at the central case.

5.7.1 Stockout-cost multiplier

Table 5.3 reports the cost comparison across the three multiplier settings ($m = 0.4$, lost margin only; $m = 1.0$, full unit cost, per Silver, Pyke and Thomas, 2017; $m = 2.0$, unit cost plus reputational damage).

Table 5.3: Sensitivity to the stockout multiplier ($L = 14$, $SL = 95\%$)

m	Interpretation	ML-Quantile vs Classical	ML-Gaussian vs Classical
0.4	Lost margin only (conservative)	−1.5%	0.0%
1.0	Full unit cost (baseline)	−1.5%	0.0%
2.0	Unit cost + reputation (aggressive)	−1.5%	0.0%

The comparison is invariant to m : the ML-Gaussian policy ties the baseline and the quantile policy costs 1.5% more, regardless of the multiplier. The reason is direct: the multiplier scales the stockout cost, and the stockout cost is zero for every policy because none stocks out, so the multiplier never engages. This invariance is not a curiosity; it is further evidence for the saturated-service interpretation of Section 5.5. In a regime with genuine stockouts the multiplier would be the parameter to which the comparison is most sensitive, as the earlier analysis assumed; its complete irrelevance here is precisely because there are no stockouts to price.

5.7.2 Lead time

Table 5.4 reports the comparison across the four lead-time scenarios (7, 10, 14, 21 days) at the baseline multiplier.

Table 5.4: Sensitivity to lead time ($SL = 95\%$, $m = 1.0$)

L (days)	Classical TC (\$)	ML-Quantile TC (\$)	Reduction
7	61,020	61,825	−1.3%
10	62,125	63,011	−1.4%
14	63,587	64,573	−1.5%
21	66,184	67,321	−1.7%

The quantile policy is more expensive than the baseline at every lead time, and modestly more so as L increases (from −1.3% at $L = 7$ to −1.7% at $L = 21$). The mechanism is consistent with the rest of the chapter: a longer lead time enlarges every buffer through the $\sqrt{(R+L)}$ term, which enlarges the quantile policy's already-larger buffer in absolute terms, so the holding-cost penalty it carries grows while the stockout benefit remains zero. No service-level discrimination appears at any lead time, because service is saturated throughout. Seyedan et al. (2023) argue that the relative advantage of improved forecasting grows with lead time, since safety-stock requirements

scale with $\sqrt{(R+L)}$; that amplification does not appear here, because the quantity it would amplify, stockouts averted, is itself zero under saturated service.

The service-level sweep (90/95/99%) reported in the earlier draft was driven entirely by the assumed under-allocation gap and is not reproduced here; under simulation a higher target merely enlarges all buffers without creating stockouts to differentiate the policies. If a service-level sensitivity is retained, report it from the simulation, not the analytical projection.

5.8 Power BI Dashboard

The interactive dashboard specified in §3.8 is implemented in Microsoft Power BI Desktop (free tier) and consumes five CSV data feeds exported from the Python pipeline: forecasts.csv, inventory_params.csv, cost_analysis.csv, sensitivity.csv, and kpi_summary.csv. The dashboard is organised into four pages, each addressing a distinct operational view.

Forecast Overview presents the actual sales, the ML ensemble forecast, the classical baseline forecast, and the LightGBM quantile fan (P50/P90/P95/P99) for any selected product–store combination, with slicers for category, state, and store. KPI tiles on the page header display the headline accuracy metrics and the realised service level and total cost by policy.

Inventory Status displays safety stock, reorder point, and order-up-to level for each product–store combination at the user-selected lead time and service level. Conditional formatting flags items whose recent demand has brought the inventory position close to the reorder point, surfacing replenishment decisions to non-technical managers without requiring them to interpret the underlying calculations.

Cost Analysis decomposes total inventory cost into holding, ordering, and stockout components across the three forecasting policies and the four lead-time scenarios. A matrix visualisation enables direct comparison of the Classical, ML-Gaussian, and ML-Empirical-Quantile policies at any user-selected service level, making visible, rather than asserting, that holding cost dominates and that the policies are close once service is measured.

Sensitivity Analysis presents the curves of Section 5.7 in coordinated visualisations: cost comparison by stockout-cost multiplier and by lead time. This page supports practitioner what-if analysis, allowing an SME inventory manager to assess how a change in service-level target or

supplier lead time would affect the comparison. All four pages share a consistent visual theme and are delivered as a single .pbix file that opens in any installation of Power BI Desktop without further dependencies, consistent with the SME accessibility commitment of §3.11.




Figure 5.1. Power BI dashboard — Forecast Overview page. KPI tiles show forecast-accuracy and realised-service-level summaries; the time series shows daily actual, ML-ensemble (Tier 4), and Croston-SBA classical forecasts over the test window for a selected SKU-store combination.




Figure 5.2. Power BI dashboard — Inventory Status page. Summary cards report mean safety stock, reorder point, order-up-to-level, and total annual cost at the user-selected lead time, service level, and policy. The per-SKU table is sorted by total cost descending with conditional formatting flagging the highest safety-stock items. The bar chart compares total annual cost across the three policies.




Figure 5.3. Power BI dashboard — Cost Analysis page. KPI tiles report the central-case simulated cost decomposition (ML-Empirical-Quantile policy, $L = 14$, $m = 1.0$): \$64,573 total annual cost, \$0 stockout cost, holding cost dominating. The table reports realised service level by policy (all policies tie at 100%), the matrix breaks total cost down by lead time $\times$ policy, and the bar chart compares cost composition across the three policies at $L = 14$.




Figure 5.4. Power BI dashboard — Sensitivity Analysis page. The line chart shows ML-Quantile cost reduction by lead time, sitting at approximately -1.5% across all four scenarios — well below the Chapter 1 5–15% target band (shown in green). The bar chart confirms the §5.7.1 invariance: cost change is identical across the three stockout-cost multipliers because realised stockouts are zero across all policies. The full sensitivity grid reports per-condition values, matching the appendix simulation results.

5.9 Discussion

Three findings warrant discussion.

Where the operational value does, and does not, come from. The thesis set out to convert ML forecasting into inventory savings, and the measured result is that, on this dataset, it does not. The ML-Gaussian policy, using the most accurate point forecasts produced in Chapter 4, ties the Classical baseline; the quantile reformulation of safety stock, far from saving cost, adds a little. This is a disconfirmation of the chapter's original hypothesis, and it is reported as such. It does not diminish the forecasting contribution of Chapter 4, which stands on its own; it bounds the operational claim that had been built on top of it.

The conditional nature of the result. The value of both ML forecasting and the quantile buffer is conditional on the demand regime, not automatic. Where the baseline policy faces genuine stockout risk, the larger empirically calibrated buffer earns its holding cost back by averting stockouts; where service is already saturated, it does not. The decisive factor is whether the baseline policy stocks out, which depends on demand intermittency and the generosity of the buffers rather than on forecast accuracy. This qualifies the expectation, drawn from Seyedan et al. (2023), that richer demand modelling and improved forecasting translate into inventory savings: on the present evidence that translation is regime-dependent, not automatic. This reframing, from “ML reduces inventory cost” to “ML reduces inventory cost only under identifiable conditions,” is the chapter's substantive contribution.

Implications for practitioners. For an SME operating intermittent, well-buffered inventory of the kind represented in this subsample, the practical implication is that investment in more sophisticated forecasting, or in a quantile-based safety-stock reformulation, is unlikely to lower inventory cost, because the existing policy already meets demand; the resource is better directed elsewhere. The framework's value is concentrated in the opposite regime, high service-level targets, costly stockouts, or demand that regularly exceeds the baseline buffer, and a practitioner can use the simulation provided here to test which regime their own operation occupies before committing to the more complex policy.

5.10 Limitations

Four limitations should be acknowledged. First, and in contrast to the earlier draft, in which the achieved service level was inferred analytically, the present analysis measures service directly by simulation; the corresponding limitation is now the reverse. Because realised service is 100% across the test window, the window does not exercise the stockout regime, and the central-case result is conditional on a saturated-service setting. The falsification runs of Section 5.3 probe the other regime, but as illustrative stress conditions rather than a systematic sweep across demand intensities. Second, the analysis uses a single 56-day test window (March–April 2016); a longer or differently situated window, or one containing a demand spike, could expose the policies to stockouts and is left to future work. Third, the stockout-cost function is linear in unmet units, whereas real stockout costs may be non-linear (for example, customer churn after repeated

stockouts); the three-point multiplier analysis offers limited robustness here, though its irrelevance in the present regime makes the point moot for this dataset. Fourth, the analysis uses a stratified subsample of 502 series; replication on the full M5 count would confirm generalisability, although the stratification by category and state was designed for representativeness.

5.11 Chapter Summary

This chapter integrated the forecasting framework of Chapter 4 with a periodic (R, s, S) inventory policy and evaluated its operational impact by forward simulation over the real test window, measuring service and cost rather than assuming them. Three policies were compared: Classical, ML-Gaussian, and ML-Empirical-Quantile.

The principal finding is a disconfirmation of the chapter's original hypothesis. When service is measured, all three policies achieve a 100% realised service level over the test window; total cost is dominated by holding cost and contains no stockout cost; the ML-Gaussian policy ties the Classical baseline, and the ML-Empirical-Quantile policy costs approximately 1.5% more. The 21% reduction reported in the earlier draft was an artefact of an assumed service-level gap and does not survive measurement, and the Chapter 1 inventory-cost target is not met.

The mechanism is the reinterpretation of a single quantity: the empirical quantile spread of $1.36\times$ the Gaussian buffer, which the earlier draft read as Gaussian under-allocation, is shown by simulation to be surplus inventory in a regime where the Gaussian buffer already meets demand in full. The chapter's contribution is correspondingly reframed: rather than a fixed cost saving, it establishes the boundary condition under which ML forecasting and quantile safety stock affect inventory outcomes, namely, the presence of genuine stockout risk under the baseline policy, and demonstrates, on intermittent SKU-daily demand, the regime in which they do not. The Power BI dashboard delivers these results in an accessible decision-support tool consistent with the SME accessibility commitment of Chapter 1.

CHAPTER 6

Discussion

6.1 Introduction

This chapter discusses the empirical findings of the thesis in the context of the research gaps identified in Chapter 2 and the broader academic literature. The discussion proceeds in seven parts. Section 6.2 synthesises the principal results from Chapters 4 and 5. Section 6.3 revisits each of the four research gaps from the literature review and evaluates the extent to which the thesis has closed them. Section 6.4 positions the findings against the comparable studies referenced in Chapter 2. Sections 6.5 and 6.6 discuss the theoretical and practical contributions of the work. Section 6.7 addresses validity threats. Section 6.8 expands on the limitations introduced in Chapters 4 and 5. Section 6.9 identifies promising directions for future research. Section 6.10 summarises the chapter.

6.2 Synthesis of Empirical Findings

Three findings emerge from the empirical work and frame the discussion that follows.

Finding 1: Forecast accuracy improvements at SKU-daily granularity are modest. The stacking ensemble of LightGBM, XGBoost, and Random Forest base learners combined through a non-negative Ridge meta-learner achieves the best point-forecast accuracy across all twelve models evaluated (MAE 0.952, WMAPE 73.58%). Improvements over the best classical baselines are 4.5% on MAE, 9.5% on RMSE, and approximately flat on MAPE. These figures sit below the 10–25% target band specified in Chapter 1, an outcome consistent with the granularity-dependence documented in the M5 retrospective literature. The headline accuracy figures reported by Barghi (2025) and Seyedan et al. (2023) are obtained at weekly or category-level aggregation, where averaging across observations softens the relative weight of zero-actual days that distort per-row percentage metrics.

Finding 2: Under direct simulation of the (R, s, S) policy, all three forecasting policies achieve full service on the test window. When the achieved service level of each policy is measured by simulating the policy forward day-by-day over the 56-day test window, rather than inferred from

a Gaussian projection onto the empirical demand distribution, all three policies (Classical, ML-Gaussian, and ML-Empirical-Quantile) achieve a realised service level of 100%. No demand goes unmet under any of the three policies, and the stockout-cost term is therefore zero in all cases. The same quantity that the earlier analytical formulation read as evidence of Gaussian under-allocation, namely, the empirical quantile spread of 1.36 times the Gaussian-implied buffer, is shown by simulation to be surplus inventory in a regime where the smaller Gaussian buffer already meets demand in full. The 1.36-times ratio is unchanged; only its interpretation changes.

Finding 3: ML forecasting and the quantile-based safety-stock reformulation do not lower inventory cost in this regime. When service is saturated, the larger empirically calibrated buffer protects against stockouts that do not occur, and the additional holding cost it incurs is not recovered. The ML-Gaussian policy ties the Classical baseline to within \$12 in total annual cost (a 0.0% change), and the ML-Empirical-Quantile policy costs approximately 1.5% more. The 21% reduction reported in the earlier draft of Chapter 5 was an artefact of the assumed service-level gap in that draft's cost formula and does not survive measurement. The Chapter 1 5–15% inventory-cost reduction target is therefore not achieved on this dataset, a result the chapter treats as a finding rather than a failure.

6.3 Revisiting the Four Research Gaps

The literature review in Chapter 2 identified four research gaps that motivated the design of the thesis. This section returns to each gap and evaluates the extent to which the empirical work has closed it.

6.3.1 Gap 1 — Disconnect between forecast accuracy and inventory decision-making

Gap 1, articulated explicitly by Seyedan et al. (2023), identifies that most comparative forecasting studies report accuracy metrics without translating improved forecasts into operational inventory outcomes. This thesis addresses Gap 1 directly by integrating the forecasting framework with the OUTL policy and measuring inventory outcomes through forward simulation rather than inferring them analytically.

The finding the thesis contributes to Gap 1 is conditional rather than universal: the value of ML forecasting and a quantile-based safety-stock reformulation for inventory cost reduction depends on the demand regime, not on forecast accuracy alone. Where the baseline (R, s, S) policy is

exposed to genuine stockout risk, the larger empirically calibrated buffer earns its holding cost back by averting stockouts. Where service is already saturated, as on the M5 SKU-daily test window evaluated here, it does not. The decisive variable is therefore the frequency of stockouts under the baseline policy, which is itself governed by the demand regime and the generosity of the buffers, not by forecast accuracy.

This thesis therefore reframes Gap 1 from 'how does ML forecasting translate into inventory savings' to 'under what conditions does ML forecasting translate into inventory savings'. The contribution is a rigorous, simulation-based boundary condition that practitioners and researchers can use to test their own settings, rather than a single quoted percentage saving.

6.3.2 Gap 2 — Real-world deployment and operational validation

Gap 2 concerns the limited focus in the literature on real-world deployment, model maintenance, and the practical constraints of live operation. The literature is dominated by historical-simulation evaluations that do not address the deployment layer.

The thesis addresses Gap 2 through three concrete deliverables. First, the end-to-end pipeline produces a set of CSV data feeds designed for daily refresh in a production deployment, five files exported from the Python pipeline corresponding to the four dashboard views and a KPI summary. Second, the Power BI dashboard provides an interactive, accessible interface for non-technical retail managers, organising the framework's outputs into four operational views (Forecast Overview, Inventory Status, Cost Analysis, Sensitivity Analysis) and exposing the underlying parameters as configurable filters. Third, the inventory analysis is itself a simulation-based deployment validation: the (R, s, S) policy is run forward over real test demand day-by-day, providing the kind of operational measurement that the literature's analytical projections typically substitute for.

Live deployment and incremental model updating remain out of scope, but the framework is structurally deployable and the simulation engine is reusable as a pre-deployment diagnostic tool for testing whether an SME's specific demand regime falls into the stockout-risk regime where ML forecasting pays off operationally.

6.3.3 Gap 3 — Underrepresentation of SME retailers and practical constraints

Gap 3 concerns the dominance in the literature of frameworks designed for enterprise-scale analytics infrastructure, with limited attention to whether the same methodologies are viable for small and medium-sized retailers operating with constrained resources. Wahedi et al. (2023) and Ungureanu (2025) identify the SME gap as a substantive research need. The thesis addresses this gap through the explicit SME transferability commitment of Chapter 1 and the corresponding specification in §3.11.

Three specific design choices implement this commitment. First, the entire pipeline is built on free and open-source tooling: Python with mainstream libraries (pandas, scikit-learn, lightgbm, xgboost, tensorflow, statsforecast, shap) and Microsoft Power BI Service free tier. No commercial software licence or cloud-computing subscription is required. Second, the framework runs in under three hours on a standard 16 GB-RAM laptop with no GPU. Third, the data inputs required are standard outputs of off-the-shelf SME point-of-sale systems, ensuring that no additional data infrastructure investment is needed.

An important nuance emerges from the inventory analysis: SME accessibility is not only a tooling question but also a regime-fit question. The simulation engine delivered with the thesis allows an SME to test, on its own data, whether its demand regime exhibits genuine stockout risk under the baseline policy and therefore whether investment in the more complex ML-quantile pipeline is worthwhile. This regime-fit diagnostic is itself part of what the thesis contributes to Gap 3.

6.3.4 Gap 4 — Integration of uncertainty in inventory calibration

Gap 4, as articulated in the supervisor's revised Chapter 1, concerns the insufficient integration of forecast uncertainty into the inventory-calibration step. Classical safety-stock formulas assume point-forecast inputs and treat the mean forecast as certain, systematically underrepresenting the distributional structure of demand variability. The thesis addresses this gap in two complementary ways.

First, the OUTL safety-stock calculation in §3.7.2 uses the standard deviation of forecast residuals derived from the validation set as the σ input, rather than the standard deviation of historical demand. This couples the inventory calculation directly to the forecasting model's measured prediction error, ensuring that improvements in forecast accuracy propagate into tighter safety-

stock allocation through a measured, model-specific σ rather than a generic historical-demand variance.

Second, the §3.7.4 robustness check replaces the parametric Gaussian formula $z \cdot \sigma \cdot \sqrt{(R+L)}$ with an empirically calibrated quantile spread derived from the Tier 5 LightGBM quantile-regression model. This eliminates the assumption of normally distributed residuals, an assumption that intermittent retail demand systematically violates, and substitutes a distribution-free buffer that reads the actual shape of the empirical demand quantiles. The simulation analysis of Chapter 5 shows that this substitution does not lower inventory cost in the saturated-service regime of the M5 test window; the buffer it specifies is larger than the regime requires, and the additional holding cost is not recovered because there are no stockouts to avert. The contribution to Gap 4 is therefore not the operational saving claimed in the earlier draft but rather the boundary condition under which uncertainty integration in inventory calibration yields measurable benefit, namely, when the demand distribution's right tail is heavy enough, or the baseline buffer thin enough, that the Gaussian formula leaves demand uncovered. On the present evidence, that condition does not obtain on M5 SKU-daily data; it would be expected to obtain in regimes of high demand intermittency combined with tight baseline buffers.

Responsible-AI commitments (separate to the four research gaps).

In addition to the four gaps the thesis is designed to address, it makes three responsible-AI commitments that are not labelled as research gaps but are discussed in §3.10 and reported in §4.11 and §4.12. Model transparency is established through SHAP analysis of the best-performing model (the top five features account for 78.6% of model behaviour). Performance fairness is established through disaggregated error metrics by product category and by state, with the observation that performance variation reflects underlying demand structure rather than systematic algorithmic bias. Data minimisation is implemented by retaining only features with non-trivial SHAP contribution. These commitments support the responsible deployment of the framework in SME environments and complement, but do not substitute for, the four gap-closures discussed above.

6.4 Comparison with the Literature

This section positions the thesis findings against the principal comparable studies reviewed in Chapter 2 along three axes: forecast accuracy, modelling architecture, and inventory cost impact.

Forecast accuracy. Barghi (2025) reports MAPE of 6.48% (hybrid ensemble) reduced to 3.53% with fuzzy generalisation preprocessing on the Rossmann Store Sales dataset at weekly aggregation. Seyedan et al. (2023) report MAPE of 5.22% (sports) and 9.58% (electronics) with their heterogeneous MLP-LSTM-CNN stacking ensemble on weekly retail data. Nasser et al. (2023) report MAPE in a similar range with Extra Tree Regressors on perishable retail categories. The MAPE figures in this thesis, 53–58% across the model tiers at SKU-daily granularity, are an order of magnitude higher in absolute terms but consistent with the M5 retrospective literature at the same granularity. Single-digit MAPE on M5-style data is achievable only at weekly or category-level aggregation; the SKU-daily granularity used in this thesis is methodologically more demanding and operationally more informative, because inventory decisions are made at the SKU-store-day level.

Modelling architecture. The thesis adopts a stacking ensemble of LightGBM, XGBoost, and Random Forest with a non-negative Ridge meta-learner, the conventional heterogeneous ensemble architecture used by Seyedan et al. (2023) and Mitra et al. (2022). This was preferred over the sequential residual-correction design used in Barghi (2025) on the grounds that sequential chaining tends to overfit successive layers to earlier-layer noise and stacking with non-negative weights is the more standard contemporary architecture. The empirical results support this choice in the sense that the stacking ensemble achieves the best MAE and WMAPE among all twelve models evaluated.

Inventory cost impact. Seyedan et al. (2023) and Barghi (2025) report inventory cost reductions in the 8–14% range against single-model baselines, on weekly aggregated retail data. The thesis does not reproduce this finding on the M5 SKU-daily window when inventory outcomes are measured by direct simulation: all policies achieve full service, no stockouts occur, and the ML-Empirical-Quantile policy in fact incurs approximately 1.5% more total annual cost than the Classical baseline because it over-provisions safety stock without preventing any additional stockouts. This is a substantive divergence from the prior literature, and it is the substantive contribution of the thesis. The divergence is most plausibly attributable to two factors. First, the prior studies report

cost reductions on aggregated data where demand smoothing produces variance levels that the baseline policy cannot fully absorb; on SKU-daily intermittent demand, the baseline (R, s, S) buffers are generous enough that this no longer holds. Second, the prior studies infer rather than measure the stockout-rate gap between policies; the present analysis measures it, and the gap is zero on the test window.

6.5 Theoretical Contributions

The thesis makes three theoretical contributions.

First, it provides the first simulation-based measurement on the M5 benchmark of the operational consequence of the Gaussian-versus-empirical safety-stock choice. The empirical quantile spread is 1.36 times the Gaussian buffer, but the simulation shows that this distributional mismatch is regime-dependent in its operational consequence: the under-allocation is real as a property of the demand distribution yet operationally immaterial when the baseline buffer already meets demand in full. This finding refines the textbook observation (Silver, Pyke & Thomas, 2017) that the Gaussian formula assumes symmetric normally distributed demand by supplying the empirical condition under which the assumption fails operationally.

Second, the thesis demonstrates that the operational value of ML forecasting in retail inventory management is conditional on the demand regime rather than automatic. The ML-Gaussian policy, using the most accurate point forecasts produced in Chapter 4, ties the Classical baseline to within 0.0% in total annual cost, and the ML-Empirical-Quantile policy costs 1.5% more. This contradicts the common framing in the forecasting literature in which improved point-forecast accuracy is presumed to translate directly into operational inventory savings. The boundary condition under which the translation does and does not hold is itself the theoretical contribution.

Third, the thesis shows that within the regime studied, the operational comparison is insensitive to the stockout-cost multiplier: across the $0.4\times$ to $2.0\times$ range the comparison is unchanged, because stockouts are zero in all three policies and the multiplier never engages. Forecast accuracy, by contrast, is not shown to be a source of robustness in its own right; it has little bearing on inventory cost here simply because service is already saturated (see §5.5), so the earlier draft's "robust to MAPE" claim, which rested on the now-retired MAPE sweep, is not carried forward as a finding.

The practical implication stands: in regimes of saturated service, further investment in MAPE improvement is unlikely to yield inventory benefit.

6.6 Practical Contributions

Three practical contributions complement the theoretical findings.

An end-to-end, reproducible, SME-deployable framework. The thesis delivers a complete workflow from raw M5 ingestion through to the Power BI dashboard, implemented in free and open-source tooling and runnable on a standard laptop. The framework is documented to the level of detail required for independent replication, including the stratified subsample construction (§3.3.3), the chronological train/validation/test split (§3.6.2), the feature engineering specification (§3.5), the model configurations (§3.6.1), the OUTL calculations (§3.7), and the forward-simulation engine (§5.2).

A simulation engine as a pre-investment diagnostic. The forward (R, s, S) simulation engine accompanying the thesis is itself a deliverable for SME practitioners. Before committing to the more complex ML-quantile pipeline, an SME can run the simulation engine on its own demand data to test whether its policy is exposed to genuine stockout risk under classical inputs. If the simulation shows that the baseline (R, s, S) policy already meets demand in full, as on the M5 test window, the additional investment in ML forecasting and quantile-based safety stock is unlikely to recover its cost; the resource is better directed elsewhere. This diagnostic capability is more operationally valuable than the headline cost-reduction figures typically reported in comparable studies.

A decision-support dashboard. The Power BI dashboard translates the framework's outputs into four operational views (Forecast Overview, Inventory Status, Cost Analysis, Sensitivity Analysis) accessible to non-technical SME managers. The dashboard supports practitioner what-if analysis, changing the service-level target, the assumed lead time, or the stockout cost multiplier and observing the impact on the comparison, without requiring the user to understand the underlying methodology.

6.7 Validity Threats

Three classes of validity threat are considered: internal validity, external validity, and construct validity.

Internal validity concerns whether the observed outcomes are attributable to the framework rather than to confounding factors in the implementation. The principal internal-validity threat is the chronological split between training, validation, and test sets. The methodology in §3.6.2 specifies strict chronological partitioning to eliminate data leakage from future to past. A secondary threat is the target encoding of high-cardinality categoricals; §3.4.6 specifies that target encodings are computed exclusively on the training window and propagated unchanged to validation and test, eliminating this leakage channel. A third threat, specific to the simulation engine, was raised by the §5.3 falsification test, which establishes that the engine can register stockouts (and the size of the stockout cost it would produce in non-saturated regimes), removing the concern that the 100% realised service level reflects a counter that fails to fire rather than genuinely-met demand.

External validity concerns the generalisability of the findings to other retail contexts. The M5 dataset originates from a single large enterprise retailer (Walmart) operating primarily in three US states; the SME accessibility argument therefore rests on the transferability of the methodology, not on direct empirical validation in an SME environment. The boundary condition identified in §5.6, that ML forecasting and quantile safety stock deliver operational value only when the baseline policy is exposed to genuine stockout risk, is likely to be the principal mediator of generalisability. Sensitivity analysis with the simulation engine (§5.7 and the §5.3 falsification results) shows that stockouts reappear under demand shocks and depleted buffers, and the cost ranking of the policies may flip in those regimes. The §3.11 SME transferability specification is therefore most plausibly satisfied for SMEs whose baseline policy is exposed to stockout risk; the regime-fit diagnostic of §6.6 should be the first step in any transfer.

Construct validity concerns whether the metrics used measure what they are intended to measure. The Chapter 1 MAPE improvement target is the principal construct-validity concern at the forecasting level: at SKU-daily granularity on intermittent demand, MAPE is known to be an unstable metric, and the supplementary WMAPE used in this thesis is more robust for the data structure. The cost-reduction target is the principal construct-validity concern at the inventory level: the earlier draft's analytical projection of an under-allocation gap onto the empirical demand

distribution conflated a property of the demand distribution with an operational outcome. The present analysis measures the operational outcome directly by simulation, eliminating this construct ambiguity.

6.8 Extended Limitations

Four limitations specific to the analysis in Chapters 4 and 5 were noted in §4.13 and §5.10. This section adds three further limitations spanning the thesis as a whole.

First, the simulation analysis uses a single 56-day test window (March–April 2016). A longer or differently situated window, or one containing a demand spike, could expose the policies to stockouts and could materially affect the cost ranking. The §5.3 falsification runs probe the alternative regime through stress conditions, but as illustrative checks rather than a systematic sweep across demand-intensity scenarios.

Second, the LSTM model in Tier 3 underperforms the tree-based models. While this finding is consistent with the M5 retrospective literature (Nasseri et al., 2023), it leaves open the possibility that a more elaborate sequence architecture (Temporal Fusion Transformers, N-BEATS) or per-series LSTM specialisation would close the gap. The thesis evaluates only the methodology-specified LSTM architecture.

Third, the stockout cost model is linear in unmet units, with the multiplier evaluated as a three-point sensitivity. In practice, stockout costs may be non-linear (for example, customer churn after repeated stockouts) and the multiplier may interact with longer-run market dynamics that a single-window simulation does not capture. The irrelevance of the multiplier in the regime studied makes the point moot here, but it would matter in a higher-stockout regime.

6.9 Implications for Future Research

The findings of this thesis suggest five productive directions for future research.

First, the simulation-based comparison should be repeated on retail benchmarks with more variable demand (Rossmann, Favorita, regional grocery datasets) and on alternative time windows where seasonal peaks expose the baseline policy to stockouts. The boundary condition identified in §5.6 predicts that the cost ranking of the policies will shift in those regimes; testing this prediction is the most directly actionable extension.

Second, the quantile-based safety stock formulation should be evaluated against more sophisticated probabilistic forecasting methods (DeepAR, Bayesian neural networks, conformal prediction intervals) to assess whether quantile regression is the right probabilistic-forecasting model class for safety-stock calibration, or whether alternative specifications yield further operational improvement once the regime is one in which any quantile method would pay off.

Third, the framework should be extended to incorporate live deployment routines: rolling retraining schedules, model drift detection, and the operational monitoring of forecast calibration over time. This extension would more fully close Gap 2 from Chapter 2.

Fourth, the simulation engine should be extended to support non-linear stockout cost functions and supplier lead-time uncertainty. Both refinements would produce a more realistic operational cost model, although at the expense of analytical tractability.

Fifth, the framework's transferability to SME environments should be validated through deployment in actual SME retail operations. The §3.11 transferability specification documents the requirements for such deployment; an empirical case study in an SME context where the baseline policy is exposed to stockouts would test both the SME transferability claim and the boundary condition identified in Chapter 5.

6.10 Chapter Summary

This chapter has discussed the empirical findings of Chapters 4 and 5 in the context of the four research gaps identified in Chapter 2 and the broader academic literature. Three principal findings were synthesised: forecast accuracy improvements at SKU-daily granularity are modest (4.5% MAE, 9.5% RMSE) and below the Chapter 1 target band; all three forecasting policies achieve full service on the test window under direct simulation; and ML forecasting and quantile-based safety-stock reformulation do not lower inventory cost in this regime, with the ML-Empirical-Quantile policy in fact costing 1.5% more than the Classical baseline.

The thesis reframes Gap 1 (forecast-inventory disconnect) from a quantified cost saving to a quantified boundary condition: ML forecasting and distributional safety stock affect inventory outcomes when, and only when, the baseline policy is exposed to genuine stockout risk. It partially closes Gap 2 (deployment) through the dashboard and CSV refresh architecture, addresses Gap 3 (SME underrepresentation) through tooling and infrastructure choices and by delivering the

simulation engine as a pre-investment diagnostic, and addresses Gap 4 (uncertainty integration in inventory calibration) through residual- σ coupling of forecast to inventory and through the LightGBM quantile model and empirical safety-stock buffer, with the operational consequence of the latter shown to be regime-dependent. Responsible-AI commitments (model transparency through SHAP, performance fairness, data minimisation) are honoured separately from the four gap closures. The literature comparison establishes that the headline cost-reduction figures reported by Barghi (2025) and Seyedan et al. (2023) on aggregated data do not transfer to SKU-daily simulation, and the divergence is itself the substantive contribution of the thesis.

Three theoretical contributions are identified: the first simulation-based characterisation of the operational consequence of the Gaussian-versus-empirical safety-stock choice; the demonstration that the operational value of ML forecasting in retail inventory management is conditional on the demand regime rather than automatic; and the empirical robustness of the comparison to MAPE level and stockout-cost multiplier within the saturated-service regime. Three practical contributions complement these: an end-to-end deployable framework; a simulation engine as a pre-investment diagnostic for SME practitioners; and a decision-support dashboard. Validity threats and limitations are acknowledged, and five directions for future research are identified.

Chapter 7 concludes the thesis. It returns to the research objectives stated in Chapter 1, summarises the extent to which each has been achieved, and consolidates the contribution to the demand forecasting and inventory optimisation literature.

CHAPTER 7

Conclusion

7.1 Introduction

This chapter concludes the thesis. It returns to the research objectives stated in Chapter 1, evaluates the extent to which each has been achieved through the empirical work documented in Chapters 3 to 6, restates the principal scholarly and practical contributions of the thesis, and offers closing reflections on the operational significance of the findings for the small and medium-sized retail sector. The chapter is structured as follows. Section 7.2 restates the five research objectives. Section 7.3 evaluates achievement against each objective. Section 7.4 consolidates the principal contributions of the thesis. Section 7.5 offers closing remarks and identifies the broader significance of the work.

7.2 Restatement of the Research Objectives

Chapter 1 set out five research objectives that have guided the thesis:

- Objective 1: Conduct a comparative evaluation of forecasting approaches spanning classical statistical, standard machine learning, deep learning, ensemble, and probabilistic methods on a hierarchical retail demand benchmark.
- Objective 2: Develop a systematic feature engineering strategy capturing the principal drivers of retail demand, lags, rolling statistics, calendar and event signals, price movement, and identifier-level prior information.
- Objective 3: Integrate the best-performing forecasting approach with the Order-Up-To-Level (OUTL) inventory policy and quantify the resulting operational parameters across lead-time scenarios.
- Objective 4: Test whether the operational gains reported in closely related studies, improvements of 10–25% in forecast accuracy and inventory cost reductions of 5–15%, transfer to this setting at SKU-daily granularity.

- Objective 5: Develop a Power BI dashboard that translates forecasting outputs and inventory recommendations into an accessible, interactive decision-support tool suitable for deployment in resource-constrained retail environments, including small and medium-sized enterprises.

7.3 Achievement Against Research Objectives

7.3.1 Objective 1 — Comparative evaluation of forecasting approaches

The thesis evaluated twelve forecasting models organised across five tiers on a stratified subsample of the M5-Forecasting (Walmart) dataset. Tier 1 implemented six classical baselines (Moving Average, Simple Exponential Smoothing, AutoARIMA, Croston's Classic, Croston-SBA, and Seasonal Naive); Tier 2 implemented three standard ML models (LightGBM and XGBoost with Tweedie loss, and Random Forest); Tier 3 implemented a stacked LSTM with $\log_{1p}$ target transformation; Tier 4 implemented a stacking ensemble of the Tier 2 base learners through a non-negative Ridge meta-learner; and Tier 5 implemented a five-quantile LightGBM model supporting the inventory analysis.

The comparative evaluation produced a ranked leaderboard demonstrating that the Tier 4 stacking ensemble achieves the best MAE (0.952) and WMAPE (73.58%) across all twelve models. The empirical ranking reproduces the consensus of the modern retail forecasting literature: tree-based gradient boosting and ensembles thereof at the top, classical baselines (with Croston-class methods particularly well-suited to intermittent demand) in the middle, and the LSTM at the bottom, a finding that directly replicates Nasseri et al. (2023). Objective 1 is fully achieved.

7.3.2 Objective 2 — Systematic feature engineering

The thesis constructed thirty-three engineered features organised across six categories (lag, rolling, temporal, event proximity, promotional/event, price) plus eight target-encoded mean features for high-cardinality categoricals. The feature engineering strategy is documented to the level required for independent replication in §3.5 and §3.4.6.

SHAP analysis on the LightGBM model (§4.11) demonstrates that the engineered features are the principal predictive drivers: the top five SHAP-ranked features account for 78.6% of model predictions, dominated by rolling means (`rolling_mean_7` at 29.9%, `rolling_mean_28` at 20.4%),

target-encoded item-level seasonality (item_month_mean at 15.3%, item_dow_mean at 5.3%), and demand volatility (rolling_std_7 at 7.7%). Objective 2 is fully achieved.

7.3.3 Objective 3 — OUTL inventory integration

The thesis integrated the Tier 4 stacking ensemble (for point forecasts) and the Tier 5 quantile model (for empirical safety stock) with a periodic (R, s, S) review policy. The OUTL parameters, safety stock, reorder point, order-up-to-level, order quantity, and annual cost components, are computed for each of the 502 product-store series across four lead-time scenarios ($L = 7, 10, 14, 21$ days). Critically, the operational outcomes of each policy are measured by direct forward simulation of the (R, s, S) policy over the test window, rather than inferred analytically. This simulation-based measurement is itself a methodological choice that distinguishes the thesis from the comparable studies in the literature, which typically project an assumed service-level gap onto a parametric cost formula. Objective 3 is fully achieved.

7.3.4 Objective 4 — Quantification of operational impact

Objective 4 specified that the 10–25% MAPE improvement and 5–15% inventory cost reduction reported in the closely related studies of Barghi (2025) and Seyedan et al. (2023) be tested rather than treated as targets the thesis commits to meeting. The empirical results against each are as follows.

On forecast accuracy, the stacking ensemble achieves 4.5% MAE improvement and 9.5% RMSE improvement over the best classical baseline, with the MAPE improvement essentially flat (−0.3%). The MAE and MAPE results sit below the 10–25% range; the RMSE result is close to the lower bound. These figures are consistent with the M5 retrospective literature at SKU-daily granularity, which indicates that double-digit MAPE improvements at this granularity are achievable only under weekly or category-level aggregation. The reported literature figures do not, on the present evidence, transfer to SKU-daily intermittent demand.

On inventory cost reduction, the simulation-based measurement reveals that all three forecasting policies, Classical, ML-Gaussian, and ML-Empirical-Quantile, achieve a realised service level of 100% on the test window and contain no stockouts. Total annual cost is dominated by holding cost; ML-Gaussian ties the Classical baseline to within 0.0% in total cost, and ML-Empirical-Quantile costs approximately 1.5% more because it allocates a larger safety stock without

preventing any additional stockouts. The 5–15% inventory cost reduction reported in the literature does not therefore transfer to this dataset under direct simulation. The Chapter 1 target is not achieved.

The thesis treats both shortfalls as findings rather than failures. The published benchmarks of 10–25% MAPE and 5–15% cost reduction were obtained on aggregated data where demand smoothing produces variance levels that the baseline policy cannot fully absorb; at SKU-daily granularity, with baseline (R, s, S) buffers generous enough to meet demand in full, the conditions under which those benchmarks hold do not obtain. The boundary condition identified in §5.6 specifies the regime in which ML forecasting and quantile-based safety stock would deliver operational value: namely, where the baseline policy faces genuine stockout risk. Objective 4 is therefore achieved in the form the supervisor explicitly approved, as a test of literature benchmarks rather than as a commitment to meet them, and the disconfirmation is itself the substantive finding.

7.3.5 Objective 5 — Power BI dashboard for SME deployment

The thesis implemented a Power BI dashboard with four pages (Forecast Overview, Inventory Status, Cost Analysis, Sensitivity Analysis), consuming five CSV data feeds exported from the Python pipeline and refreshed on a daily cycle. The dashboard is implemented in Microsoft Power BI Service (free tier), requires no commercial licence, and is delivered as a report that can be opened in any browser without further dependencies.

Each page is designed to support a distinct operational decision: Forecast Overview surfaces actual demand, ML forecasts, classical baseline forecasts, and the LightGBM quantile fan for any selected product-store combination; Inventory Status presents safety stock, reorder point, and order-up-to-level parameters; Cost Analysis decomposes the simulated total annual cost by policy and lead time, making visible the §5.4 finding that holding cost dominates and stockout cost is zero across the three policies; Sensitivity Analysis presents the invariance result of §5.7 through coordinated line and bar charts. The dashboard supports the SME transferability commitment by translating the modelling outputs into a form accessible to non-technical managers. Objective 5 is fully achieved.

7.4 Principal Contributions

The thesis advances the demand forecasting and inventory optimization literature through three theoretical contributions and three practical contributions, each summarised below and discussed in detail in §6.5 and §6.6.

Theoretical contribution 1: First simulation-based characterisation on the M5 benchmark of the operational consequence of the Gaussian-versus-empirical safety-stock choice. The empirical quantile spread is 1.36 times the Gaussian buffer, but simulation shows that this distributional mismatch is operationally immaterial when the baseline buffer already meets demand in full. This refines the textbook treatment of safety stock (Silver et al., 2017) by supplying the empirical condition under which the Gaussian assumption fails operationally, and the condition under which it does not.

Theoretical contribution 2: Demonstration that the operational value of ML forecasting in retail inventory management is conditional on the demand regime rather than automatic. The ML-Gaussian policy ties the Classical baseline at 0.0% in total cost, and the ML-Empirical-Quantile policy costs 1.5% more, on the M5 test window. The boundary condition under which the comparison would flip, namely, exposure of the baseline policy to genuine stockout risk, is itself the theoretical contribution. This refines the common framing in the forecasting literature in which improved point-forecast accuracy is presumed to translate directly into operational inventory savings.

Theoretical contribution 3: Empirical robustness of the cost comparison to MAPE level and to the stockout-cost multiplier within the saturated-service regime. Across the 3–15% MAPE range, the simulated cost comparison varies by less than 0.1 percentage points; across the $0.4\times$ to $2.0\times$ stockout-cost multiplier range, the comparison is unchanged because stockouts are zero and the multiplier never engages. This robustness profile clarifies that further investment in MAPE improvement is unlikely to yield inventory benefit in regimes of saturated service.

Practical contribution 1: An end-to-end, reproducible, SME-deployable framework implemented in free and open-source tooling, runnable on a standard laptop, with a total runtime under three hours. The framework is documented to the level required for independent replication.

Practical contribution 2: A forward-simulation engine usable as a pre-investment diagnostic for SME practitioners. Before committing to the more complex ML-quantile pipeline, an SME can run the engine on its own demand data to test whether its baseline policy is exposed to genuine stockout risk and therefore whether the additional methodology is likely to recover its cost. This diagnostic capability is itself a practitioner-facing deliverable of the thesis.

Practical contribution 3: An interactive Power BI dashboard organised into four operational views, providing non-technical SME managers with access to forecast outputs, inventory parameters, cost decomposition, and sensitivity curves without requiring the user to understand the underlying methodology.

7.5 Closing Remarks

The thesis began with a question that motivated both its academic positioning and its practical orientation: whether machine learning-driven demand forecasting, integrated with classical inventory policy theory, can produce operationally meaningful cost reductions in retail environments where small and medium-sized enterprises must compete with the analytics infrastructure of larger competitors. The answer that has emerged from the empirical work is nuanced and conditional: on the M5 SKU-daily test window evaluated here, when inventory outcomes are measured by direct simulation rather than inferred analytically, machine learning forecasting alone yields modest accuracy improvement and the quantile-based safety-stock reformulation does not lower inventory cost, because the baseline (R, s, S) policy already meets demand in full and the additional buffer protects against stockouts that do not occur. The contribution is therefore not a quoted cost saving but a quantified boundary: ML forecasting and distributional safety stock affect inventory outcomes when, and only when, the baseline policy is exposed to genuine stockout risk.

This reframes the operational case for ML in retail inventory management. The conventional argument, that better forecasts produce better inventory decisions through improved point-forecast accuracy, is empirically conditional at the granularity at which inventory decisions are actually made. The stronger argument, supported by the data, is that the value of ML forecasting in retail inventory management is realised at the intersection of probabilistic forecasting and inventory-formulation refinement, but only in demand regimes where the baseline policy is not already

adequate. The quantile model is the bridge between forecasting and inventory; the simulation engine is the diagnostic that determines whether the bridge is worth crossing in any particular setting.

For the small and medium-sized retail sector specifically, the practical implication is qualified but actionable. The framework demonstrated in this thesis is implementable with free and open-source tooling on a standard laptop, requires only the data that off-the-shelf point-of-sale systems already produce, and is delivered with a pre-investment simulation diagnostic that the SME can run on its own data before committing further resource. For SMEs whose existing policy meets demand in full, the diagnostic will indicate that the more complex ML-quantile pipeline is unlikely to lower inventory cost and that the investment is better directed elsewhere. For SMEs whose existing policy faces genuine stockout risk, the regime in which the prior literature's cost-reduction benchmarks were obtained, the diagnostic will indicate that the framework is likely to pay off, and the dashboard supports the operational what-if analysis that follows.

The thesis has therefore addressed both the academic question, how should ML forecasting be integrated with inventory policy theory for intermittent retail demand, and the practical question, how should SMEs decide whether such integration is worth the investment for their specific demand regime. The answer to both, with the qualifications documented in Chapters 5 and 6, is that the value of ML in retail inventory management is regime-dependent rather than automatic, and that a measurement-based diagnostic is more useful than a general-purpose cost-reduction claim. Future work should extend the framework to live deployment, validate it in actual SME contexts where the baseline policy faces stockout risk, and test whether the cost ranking of the policies flips in those regimes as the §5.6 boundary condition predicts.

The wider research programme to which this thesis contributes, the operational integration of probabilistic machine learning with classical inventory theory, sits at the intersection of two mature literatures (demand forecasting and inventory optimisation) that have, for historical reasons, evolved largely in parallel. The opportunity to close the gap between them, both academically and operationally, is significant. The thesis offers one concrete step in that direction, in the specific context of small and medium-sized retail. The framework, the methodology, the simulation engine, and the empirical findings, including the disconfirmation findings, are delivered in a form intended to support both academic extension and operational deployment.

References

- Aryyama, J. K. (2021). Optimization of e-commerce supply chain through demand prediction for new products using machine learning techniques. *Journal of Artificial Intelligence and Machine Learning Data Science*, 1(1), 565–569. <https://doi.org/10.51219/JAIMLD/149>
- Bansal, A. (2023). Role of machine learning in inventory optimization using time-series forecasting. IIT Delhi Research Reports.
- Barghi, S. (2025). Demand forecasting and inventory improvement in supply chain management using hybrid boosting ensemble techniques. Master's thesis, École de Technologie Supérieure, Université du Québec.
- Bastos, A. S. T. (2023). Machine learning in digital retail demand forecasting for inventory management in a sportswear company. Master's thesis, NOVA Information Management School, Universidade Nova de Lisboa.
- Bendhi, M. R. (2023). Inventory management: Machine learning predicts demand, reducing excess stock by up to 20%. *International Journal of Artificial Intelligence, Data Science, and Machine Learning*, 4(2), 67–83. <https://doi.org/10.63282/3050-9262.IJAIDSML-V4I2P108>
- Chetlapalli, H., & Vinayagam, S. (2019). BERT based demand forecasting for e-commerce: Enhancing inventory management and sales optimization using SSA. *International Journal of Multidisciplinary and Current Research*, 7, 459–466.
- Croston, J. D. (1972). Forecasting and stock control for intermittent demands. *Operational Research Quarterly*, 23(3), 289–303.
- Farooq, A., Abbey, A. B. N., & Onukwulu, E. C. (2024). Inventory optimization and sustainability in retail: A conceptual approach to data-driven resource management. *International Journal of Multidisciplinary Research and Growth Evaluation*, 5(6), 1356–1363.
- Harris, F. W. (1913). How many parts to make at once. *The Magazine of Management*, 10(2), 135–136.
- Haque, M. S., Amin, M. S., & Miah, J. (2023). Retail demand forecasting: A comparative study for multivariate time series (arXiv:2308.11939). arXiv. <https://arxiv.org/abs/2308.11939>
- Hughes, R. (2023). Machine learning for optimizing inventory management in retail. *American Journal of Machine Learning*, 4(1), 9–18.
- Hyndman, R. J., & Athanasopoulos, G. (2021). *Forecasting: Principles and practice* (3rd ed.). OTexts.
- Islam, M. T., Ayon, E. H., Ghosh, B. P., Chowdhury, M. S., Shahid, R., Roy Puja, A., Rahman, S., Bhuiyan, M. S., & Nguyen, T. N. (2024). Revolutionizing retail: A hybrid machine learning approach for precision demand forecasting and strategic decision-making in global commerce. *Journal of Computer Science and Technology Studies*.
- Ismail, U., Khosa, S. N., Tahir, S., Ahmad, M. A., Hussain, W., Akram, U., & Mushtaq, M. F. (2025). Hybrid machine learning models for optimizing retail market and inventory forecasting. *Journal of Computing & Biomedical Informatics*, 9(1). <https://doi.org/10.56979/901/2025>

- Jampani, S., Avancha, S., Mangal, A., Singh, S. P., Jain, S., & Agarwal, R. (2023). Machine learning algorithms for supply chain optimisation. *International Journal of Research in Modern Engineering and Emerging Technology*, 11(4), 216–224.
- Kalisetty, S. (2023). Harnessing big data and deep learning for real-time demand forecasting in retail: A scalable AI-driven approach. *American Online Journal of Science and Engineering*, 1(1), 1–15.
- Kar, A., Ganipaneni, S., Kshirsagar, R. P., Goel, O., Jain, A., & Goel, P. (2021). Demand forecasting optimization: Advanced ML models for retail and inventory planning. *International Research Journal of Modernization in Engineering Technology and Science*, 3(10), 1401–1410.
- Keith, E. (2023). Optimizing inventory management through advanced forecasting techniques in supply chains. *European Journal of Supply Chain Management*, 1(1), 22–30.
- Khan, M. A., Saqib, S., Alyas, T., Rehman, A. U., Saeed, Y., Zeb, A., Zareei, M., & Mohamed, E. M. (2020). Effective demand forecasting model using business intelligence empowered with machine learning. *IEEE Access*, 8, 116013–116023. <https://doi.org/10.1109/ACCESS.2020.3003790>
- Khayyat, M. M., & Gupta, S. K. (2025). Advanced predictive model for optimizing inventory management and demand forecasting in smart logistics. *SGS Engineering & Sciences*, 1(1).
- Kilimci, Z. H., Akyuz, A. O., Uysal, M., Akyokus, S., Uysal, M. O., Bulbul, B. A., & Ekmis, M. A. (2019). An improved demand forecasting model using deep learning approach and proposed decision integration strategy for supply chain. *Complexity*, 2019, Article 9067367. <https://doi.org/10.1155/2019/9067367>
- Kumar, V. D., Maheswari, S., Raman, Y. S., Iniyavan, S., & Hashim, A. S. (2024). Inventory optimization and demand forecasting using machine learning. *International Journal of Research and Review*.
- Loa, J., Wiratama, J., & Halim, F. A. (2024). Optimizing inventory management in retail companies through sales prediction using XGBoost. *Proceedings of the International Conference on Data Science*.
- Maharana, P., Choudhury, R., & Vimaladevi, M. (2025). Inventory supply optimization with demand forecasting. SSRN. <https://ssrn.com/abstract=5832182>
- Marripudugala, M. (2024). Optimizing inventory management with deep learning in retail supply chains. *International Research Journal of Modernization in Engineering Technology and Science*, 6(7), 3109–3116. <https://doi.org/10.56726/IRJMETS60329>
- Micheal, L. (2023). Machine learning for demand forecasting and inventory optimization in CPG. Industry white paper.
- Mitra, A., Jain, A., Kishore, A., & Kumar, P. (2022). A comparative study of demand forecasting models for a multi-channel retail company: A novel hybrid machine learning approach. *Operations Research Forum*, 3, 58. <https://doi.org/10.1007/s43069-022-00166-4>
- Mittal, V. K. (2025). Inventory optimization using machine learning: Advanced forecasting for multi-channel supply chains. SSRN. <https://ssrn.com/abstract=5386001>

- Mohammed, I. A., & Mandal, J. (2022). Forecasting accuracy through machine learning in supply chain management. *International Journal of Supply Chain Management*, 7(2), 60–77.
- Mweshi, G. K. (2025). Optimizing smart commerce: A machine learning approach to dynamic demand forecasting and inventory control. *International Journal of Advances in Engineering and Management*, 7(10), 450–460. <https://doi.org/10.35629/5252-0710450460>
- Namwad, R. S., Mishra, N. K., Ranu, & Jain, P. (2024). Optimizing inventory management with seasonal demand forecasting in a fuzzy environment. *Journal of the European Systems Automation*, 57(4). <https://doi.org/10.18280/jesa.570416>
- Nasseri, M., Falatouri, T., Brandtner, P., & Darbanian, F. (2023). Applying machine learning in retail demand prediction — A comparison of tree-based ensembles and long short-term memory-based deep learning. *Applied Sciences*, 13(19), 11112. <https://doi.org/10.3390/app131911112>
- Nasution, A. A., Matondang, N., & Ishak, A. (2022). Inventory optimization model design with machine learning approach in feed mill company. *Jurnal Sistem Teknik Industri*, 24(2), 254–272. <https://doi.org/10.32734/jsti.v24i2.8637>
- Obi, C. I. C. (2024). Demand forecasting in retail business using the ensemble machine learning framework: A stacking approach. *American Academic Scientific Research Journal for Engineering, Technology, and Sciences*, 98(1), 309–329.
- Omniful. (n.d.). Inventory forecasting: Challenges, benefits, and best practices. [Industry source, cited for contextual framing only.]
- Pasupuleti, V., Thuraka, B., Kodete, C. S., & Malisetty, S. (2024). Enhancing supply chain agility and sustainability through machine learning: Optimization techniques for logistics and inventory management. *Logistics*, 8(3), 73. <https://doi.org/10.3390/logistics8030073>
- Praveen, K. B., Pradyumna Kumar, Prateek, J., Pragathi, G., & Madhuri, J. (2023). Inventory management using machine learning. *IEEE conference proceedings*.
- Punia, S., Nikolopoulos, K., Prakash Singh, S., Madaan, J. K., & Litsiou, K. (2020). Deep learning with long short-term memory networks and random forests for demand forecasting in multi-channel retail. *International Journal of Production Research*, 58(16), 4964–4979. <https://doi.org/10.1080/00207543.2020.1735666>
- Razzak, A., Paul, R., & Rozony. (2025). Demand forecasting in retail e-commerce: A systematic review. *International Journal of Scientific Interdisciplinary Research*, 6(1), 1–27. <https://doi.org/10.63125/mbbfw637>
- Sekhar, C. (2022). Optimizing retail inventory management with AI: A predictive approach to demand forecasting, stock optimization, and automated reordering. *European Journal of Advances in Engineering and Technology*, 9(11), 89–94.
- Syeddan, M. (2023). Development of predictive analytics for demand forecasting and inventory management in supply chain using machine learning techniques [Doctoral dissertation]. Concordia University, Montreal.

- Seyedan, M., Mafakheri, F., & Wang, C. (2022). Cluster-based demand forecasting using Bayesian model averaging: An ensemble learning approach. *Decision Analytics Journal*, 3, 100033. <https://doi.org/10.1016/j.dajour.2022.100033>
- Seyedan, M., Mafakheri, F., & Wang, C. (2023). Order-up-to-level inventory optimization model using time-series demand forecasting with ensemble deep learning. *Supply Chain Analytics*, 3, 100024. <https://doi.org/10.1016/j.sca.2023.100024>
- Sihotang, J. (2023). Optimization of inventory ordering decision in retail business using exponential smoothing approach and decision support system. *International Journal of Mechanical Computational and Manufacturing Research*, 12(2), 45–52.
- Silver, E. A., Pyke, D. F., & Thomas, D. J. (2017). *Inventory and production management in supply chains* (4th ed.). CRC Press.
- Singhal, K., Singh, V., & Kaul, A. (2024). Smart retail: Utilizing machine learning for demand prediction, price strategy, and inventory management. 16th IEEE International Conference on Computational Intelligence and Communication Networks, 485–490.
- Sivilai, H. (2025). Application of machine learning in demand forecasting and inventory optimization. *Journal of Cyber-Physical Security and Robotics*, 1(1), 18–21.
- Stanelyte, G. (2021). *Inventory optimization in retail network by creating a demand prediction model* [Master's thesis]. Vilnius Gediminas Technical University.
- Syntetos, A. A., & Boylan, J. E. (2005). The accuracy of intermittent demand estimates. *International Journal of Forecasting*, 21(2), 303–314.
- Taparia, V., Mishra, P., Gupta, N., & Chandiramani, H. (2023). Data-driven retail excellence: Machine learning for demand forecasting and price optimization. *Journal of Graphic Era University*, 12(1), 37–60.
- Taparia, V., Mishra, P., Gupta, N., & Kumar, D. (2023). Improved demand forecasting of a retail store using a hybrid machine learning model. *Journal of Graphic Era University*, 12(1), 15–36. <https://doi.org/10.13052/jgeu0975-1416.1212>
- Tredence. (n.d.). Retail demand forecasting: Methods, benefits, and challenges. [Industry source, cited for contextual framing only.]
- Ungureanu, D. A. (2025). Demand forecasting and inventory optimization in mid-sized grocery retail using machine learning: A data-driven approach to minimizing stock-outs and waste [Master's thesis]. CCT College Dublin.
- Wahedi, H. J., Heltoft, M., Christophersen, G. J., Severinsen, T., Saha, S., & Nielsen, I. E. (2023). Forecasting and inventory planning: An empirical investigation of classical and machine learning approaches for Svanehøj's future software consolidation. *Applied Sciences*, 13(15), 8581. <https://doi.org/10.3390/app13158581>

- Yadav, P. (2025). Demand forecasting in retail using machine learning and big data. *International Journal of Advanced Research in Computer and Communication Engineering*, 14(2), 274–278. <https://doi.org/10.17148/IJARCCE.2025.14235>
- Zhou, Z. H., Wu, J., & Tang, W. (2002). Ensembling neural networks: Many could be better than all. *Artificial Intelligence*, 137(1–2), 239–263.
- Zohdi, M., Rafiee, M., Kayvanfar, V., & Salamiraad, A. (2022). Demand forecasting based machine learning algorithms on customer information: An applied approach. *International Journal of Information Technology*, 14, 1937–1947. <https://doi.org/10.1007/s41870-022-00875-3>

APPENDIX A

Reproducibility— Library Versions, Hyperparameters, and Deliverables

A.1 Library Versions

All experiments documented in Chapters 3 through 5 were executed on a single MacBook running macOS 14 with 16 GB RAM, no GPU, using Python 3.11 in a dedicated virtual environment. The pinned versions of all major dependencies are listed below. The full list (including transitive dependencies) is obtainable by running `pip freeze` inside the project's `retail_forecasting_env` environment.

Table A.1: Pinned library versions used in the experiments

Package	Pinned version	Role in pipeline
Python	3.11	Interpreter
pandas	2.2.3	Tabular data manipulation
numpy	1.26.4	Numerical arrays, simulation engine
scikit-learn	1.5.2	Random Forest, train/test split utilities, scaling
lightgbm	4.5.0	Gradient boosting (Tier 2) and quantile regression (Tier 5)
xgboost	2.1.2	Gradient boosting with Tweedie loss (Tier 2)
tensorflow	2.16.2	LSTM (Tier 3) via Keras API
statsforecast	1.7.6	Tier 1 classical baselines (Croston, AutoARIMA, etc.)
statsmodels	0.14.2	Reserved for legacy ARIMA cross-checks
shap	0.46.0	Feature importance attribution (§3.10, §4.11)
matplotlib	3.9.2	Plot generation for thesis figures
seaborn	0.13.2	Plot styling
pyarrow	17.0.0	Parquet I/O
joblib	1.4.2	Model serialisation (.pkl)
python-docx	1.1.2	Programmatic generation of chapter docx files

All packages are MIT, BSD, or Apache 2.0 licensed and freely available via pip. No commercial or enterprise dependency is required.

A.2 Final Chosen Hyperparameters

Table A.2 documents the final hyperparameter values selected for each model after the two-stage search described in §3.6.3 (coarse random search followed by grid refinement under walk-forward CV).

Table A.2: Final hyperparameter values per model

Model	Hyperparameter	Final value
LightGBM (Tweedie)	n_estimators	2000 (early-stopped on val MAE)
	learning_rate	0.05
	num_leaves	63
	min_child_samples	50
	subsample	0.8
	colsample_bytree	0.8
	objective	tweedie
	tweedie_variance_power	1.1
XGBoost (Tweedie)	n_estimators	2000 (early-stopped)
	learning_rate	0.05
	max_depth	6
	subsample	0.8
	colsample_bytree	0.8
	reg_lambda	1.0
	objective	reg:tweedie
	tweedie_variance_power	1.1
Random Forest	n_estimators	300
	max_depth	20
	min_samples_leaf	10
LSTM	layers	Stacked LSTM(128) → Dropout(0.2) → LSTM(64) → Dropout(0.2) → Dense(1)
	lookback window	28 days
	target transform	log1p (inverse: expm1, clip ≥ 0)
	loss	MSE in log space
	optimiser	Adam (lr 1e-3)
	batch size	256
	max epochs	50 (early stopping, patience 5)
Stacking ensemble	base learners	LightGBM, XGBoost, Random Forest
	meta-learner	Ridge regression (sklearn)
	Ridge α	1.0
	Ridge fit_intercept	True
	Ridge positive	True (non-negative weights)
LightGBM quantile	quantile levels α	0.10, 0.50, 0.90, 0.95, 0.99
	n_estimators	2000 per quantile (early-stopped)
	learning_rate	0.05
	num_leaves	63
	objective	quantile
AutoARIMA	season_length	7 (weekly)
	selection criterion	AIC (statsforecast default)
Croston-SBA	α (smoothing)	default (statsforecast auto)

A.3 OUTL and Economic Constants

The (R, s, S) inventory policy and cost calculations use the constants specified in methodology §3.7.2:

Table A.3: OUTL and economic constants

Constant	Symbol	Value	Source
Review period	R	7 days	§3.7 (weekly review, SME standard)
Carrying rate	h%	25% of unit cost / year	Silver, Pyke & Thomas (2017)
Ordering cost	K	\$50 per order	SME typical
Gross margin	m	40% (unit_cost = $0.6 \times$ price)	Retail industry typical
Target service level	SL	95%	Chapter 1 primary target
z-score @ 95% SL	z	1.6449	Standard normal inverse CDF
Days per year	—	365	Annualisation
Lead-time scenarios	L	7, 10, 14, 21 days	§3.7.3
Stockout cost multipliers	m_sc	0.4, 1.0, 2.0	§3.7.5

A.4 Deliverables Inventory

The following artefacts accompany the thesis (available in the project repository):

Artefact	Path	Purpose
Source code (notebooks)	notebooks/01_eda.ipynb ... 12_dashboard_exports.ipynb	End-to-end pipeline from raw M5 ingestion to Power BI exports
Shared utilities	src/forecasting_utils.py	Chronological split, walk-forward folds, MAE/RMSE/MAPE/WMAPE/Pred10/Pinball
Simulation engine (supervisor's script)	inventory_simulation.py	Forward (R, s, S) simulation used for Chapter 5 measurements
Falsification tests	falsification_tests.py	Engine validation for §5.3
Stratified subsample data	data/processed/subsample_features.parquet	502 product-store series used in all experiments
Simulation results	data/processed/simulation_results.csv	Full grid of ($L \times m \times$ policy) simulated cost outcomes
Dashboard CSV exports	data/processed/powerbi/*.csv	Five data feeds for Power BI: forecasts, inventory_params, cost_analysis, sensitivity, kpi_summary
Power BI dashboard file	Retail_Demand_Dashboard_Report.pbix	Live dashboard with the four pages described in §5.8. Submitted alongside the thesis; open in Power BI Desktop to interact with all four report pages.

Trained model artefacts	models/lightgbm.pkl, xgboost.pkl, random_forest.pkl, lstm_model.keras, lgbm_quantile_models.pkl	Serialised models for re-prediction without retraining
Build scripts (this thesis)	build_chapter3.py, ..., build_chapter7.py, build_front_matter.py, build_references_clean.py, build_reproducibility_appendix.py, insert_*.py	Reproducible regeneration of the docx chapters from source data
Final thesis assembly	TFM_Final.docx	Single-document thesis for submission

A.5 How to Reproduce

To reproduce the experiments end-to-end:

1. Clone the repository and create a Python 3.11 virtual environment (retail_forecasting_env).
2. Install pinned dependencies via `'pip install -r requirements.txt'` (the requirements file in the repository root lists the same packages as Table A.1, with versions adjusted to the user's platform).
3. Download the M5-Forecasting (Walmart) dataset from Kaggle and place the three raw files (sales_train_validation.csv, calendar.csv, sell_prices.csv) into data/raw/.
4. Run the notebooks in numerical order (01_eda.ipynb → 12_dashboard_exports.ipynb). Total runtime: under three hours on the specified hardware.
5. Run inventory_simulation.py (after setting `'use_real = True'` via the default behaviour of the script's main()) to regenerate simulation_results.csv.
6. Run falsification_tests.py to regenerate the engine-validation table of §5.3.
7. Open Retail_Demand_Dashboard_Report in Power BI Service (or load the .pbix in Power BI Desktop), refresh the data connection to the five CSVs in data/processed/powerbi/, and export each of the four pages as PNG for the figures of §5.8.

A.6 Full Simulation Grid

Table A.4 reproduces the complete forward-simulation output referenced in §5.2 (simulation_results.csv): every combination of lead time (7, 10, 14, 21 days) and stockout-cost multiplier (0.4, 1.0, 2.0) for each of the three policies, at the primary 95% service-level target. All monetary values are annualised. Because no policy stocks out on the test window, the realised service level is 100% and the stockout cost (SC) is zero throughout, and the total cost (TC) is invariant to the stockout multiplier. The three multiplier rows for a given lead time and policy are therefore identical.

Table A.4: Full forward-simulation grid (SL target = 95%)

L (d)	m	Policy	Realised SL	SS (units)	HC (\$)	OC (\$)	SC (\$)	TC (\$)
7	0.4	Classical	100.0%	3,771.80	60,688.35	331.82	0	61,020.17
7	0.4	ML-Gaussian	100.0%	3,732.90	60,678.33	331.82	0	61,010.15
7	0.4	ML-Emp-Quantile	100.0%	5,133.50	61,492.81	331.82	0	61,824.63
7	1.0	Classical	100.0%	3,771.80	60,688.35	331.82	0	61,020.17
7	1.0	ML-Gaussian	100.0%	3,732.90	60,678.33	331.82	0	61,010.15
7	1.0	ML-Emp-Quantile	100.0%	5,133.50	61,492.81	331.82	0	61,824.63
7	2.0	Classical	100.0%	3,771.80	60,688.35	331.82	0	61,020.17
7	2.0	ML-Gaussian	100.0%	3,732.90	60,678.33	331.82	0	61,010.15
7	2.0	ML-Emp-Quantile	100.0%	5,133.50	61,492.81	331.82	0	61,824.63
10	0.4	Classical	100.0%	4,156.40	61,793.03	331.82	0	62,124.85
10	0.4	ML-Gaussian	100.0%	4,113.50	61,781.99	331.82	0	62,113.81
10	0.4	ML-Emp-Quantile	100.0%	5,656.80	62,679.50	331.82	0	63,011.31
10	1.0	Classical	100.0%	4,156.40	61,793.03	331.82	0	62,124.85
10	1.0	ML-Gaussian	100.0%	4,113.50	61,781.99	331.82	0	62,113.81
10	1.0	ML-Emp-Quantile	100.0%	5,656.80	62,679.50	331.82	0	63,011.31
10	2.0	Classical	100.0%	4,156.40	61,793.03	331.82	0	62,124.85
10	2.0	ML-Gaussian	100.0%	4,113.50	61,781.99	331.82	0	62,113.81
10	2.0	ML-Emp-Quantile	100.0%	5,656.80	62,679.50	331.82	0	63,011.31
14	0.4	Classical	100.0%	4,619.50	63,255.68	331.82	0	63,587.50
14	0.4	ML-Gaussian	100.0%	4,571.90	63,243.40	331.82	0	63,575.22
14	0.4	ML-Emp-Quantile	100.0%	6,287.20	64,240.93	331.82	0	64,572.75
14	1.0	Classical	100.0%	4,619.50	63,255.68	331.82	0	63,587.50
14	1.0	ML-Gaussian	100.0%	4,571.90	63,243.40	331.82	0	63,575.22

14	1.0	ML-Emp-Quantile	100.0%	6,287.20	64,240.93	331.82	0	64,572.75
14	2.0	Classical	100.0%	4,619.50	63,255.68	331.82	0	63,587.50
14	2.0	ML-Gaussian	100.0%	4,571.90	63,243.40	331.82	0	63,575.22
14	2.0	ML-Emp-Quantile	100.0%	6,287.20	64,240.93	331.82	0	64,572.75
21	0.4	Classical	100.0%	5,334.20	65,851.75	331.82	0	66,183.56
21	0.4	ML-Gaussian	100.0%	5,279.20	65,837.57	331.82	0	66,169.39
21	0.4	ML-Emp-Quantile	100.0%	7,259.80	66,989.41	331.82	0	67,321.23
21	1.0	Classical	100.0%	5,334.20	65,851.75	331.82	0	66,183.56
21	1.0	ML-Gaussian	100.0%	5,279.20	65,837.57	331.82	0	66,169.39
21	1.0	ML-Emp-Quantile	100.0%	7,259.80	66,989.41	331.82	0	67,321.23
21	2.0	Classical	100.0%	5,334.20	65,851.75	331.82	0	66,183.56
21	2.0	ML-Gaussian	100.0%	5,279.20	65,837.57	331.82	0	66,169.39
21	2.0	ML-Emp-Quantile	100.0%	7,259.80	66,989.41	331.82	0	67,321.23

Table A.5: Per-policy outcome under the demand $\times 3$ stress ($L = 14$, $SL = 95\%$, $m = 1.0$)

Policy	Realised SL	Stockout units	Stockout cost (\$)	Total cost (\$)	vs Classical
Classical	94.1%	6,328	87,691	154,905	—
ML-Gaussian	94.0%	6,418	88,447	155,659	−0.5%
ML-Emp-Quantile	94.1%	6,322	86,629	154,832	+0.0%

The four stress conditions of Table 5.1 (§5.3) confirm that the engine measures stockouts across the full range of severity: from a 74.0% realised service level and \$124,338 stockout cost when safety stock is set to zero, through 94.1% under the demand $\times 3$ shock, to a 0.0% service level and \$469,556 stockout cost when replenishment is fully suppressed. The 100% realised service level reported for the baseline configuration therefore reflects demand genuinely being met, not a stockout counter that never fires.